



Multi Threading



- » Skan.ai - chief Architect
- » Ai.robotics - chief Architect
- » Genpact - solution Architect
- » Welldoc - chief Architect
- » Microsoft
- » Mercedes
- » Siemens
- » Honeywell



Mubarak

Agenda

- Concurrency
- Locks
- Best Practices
- Anti Patterns
- Debugging

Parallelism

Task Parallelism

- Thread pool
- Dedicated Thread
 - Background Thread
 - Foreground Thread

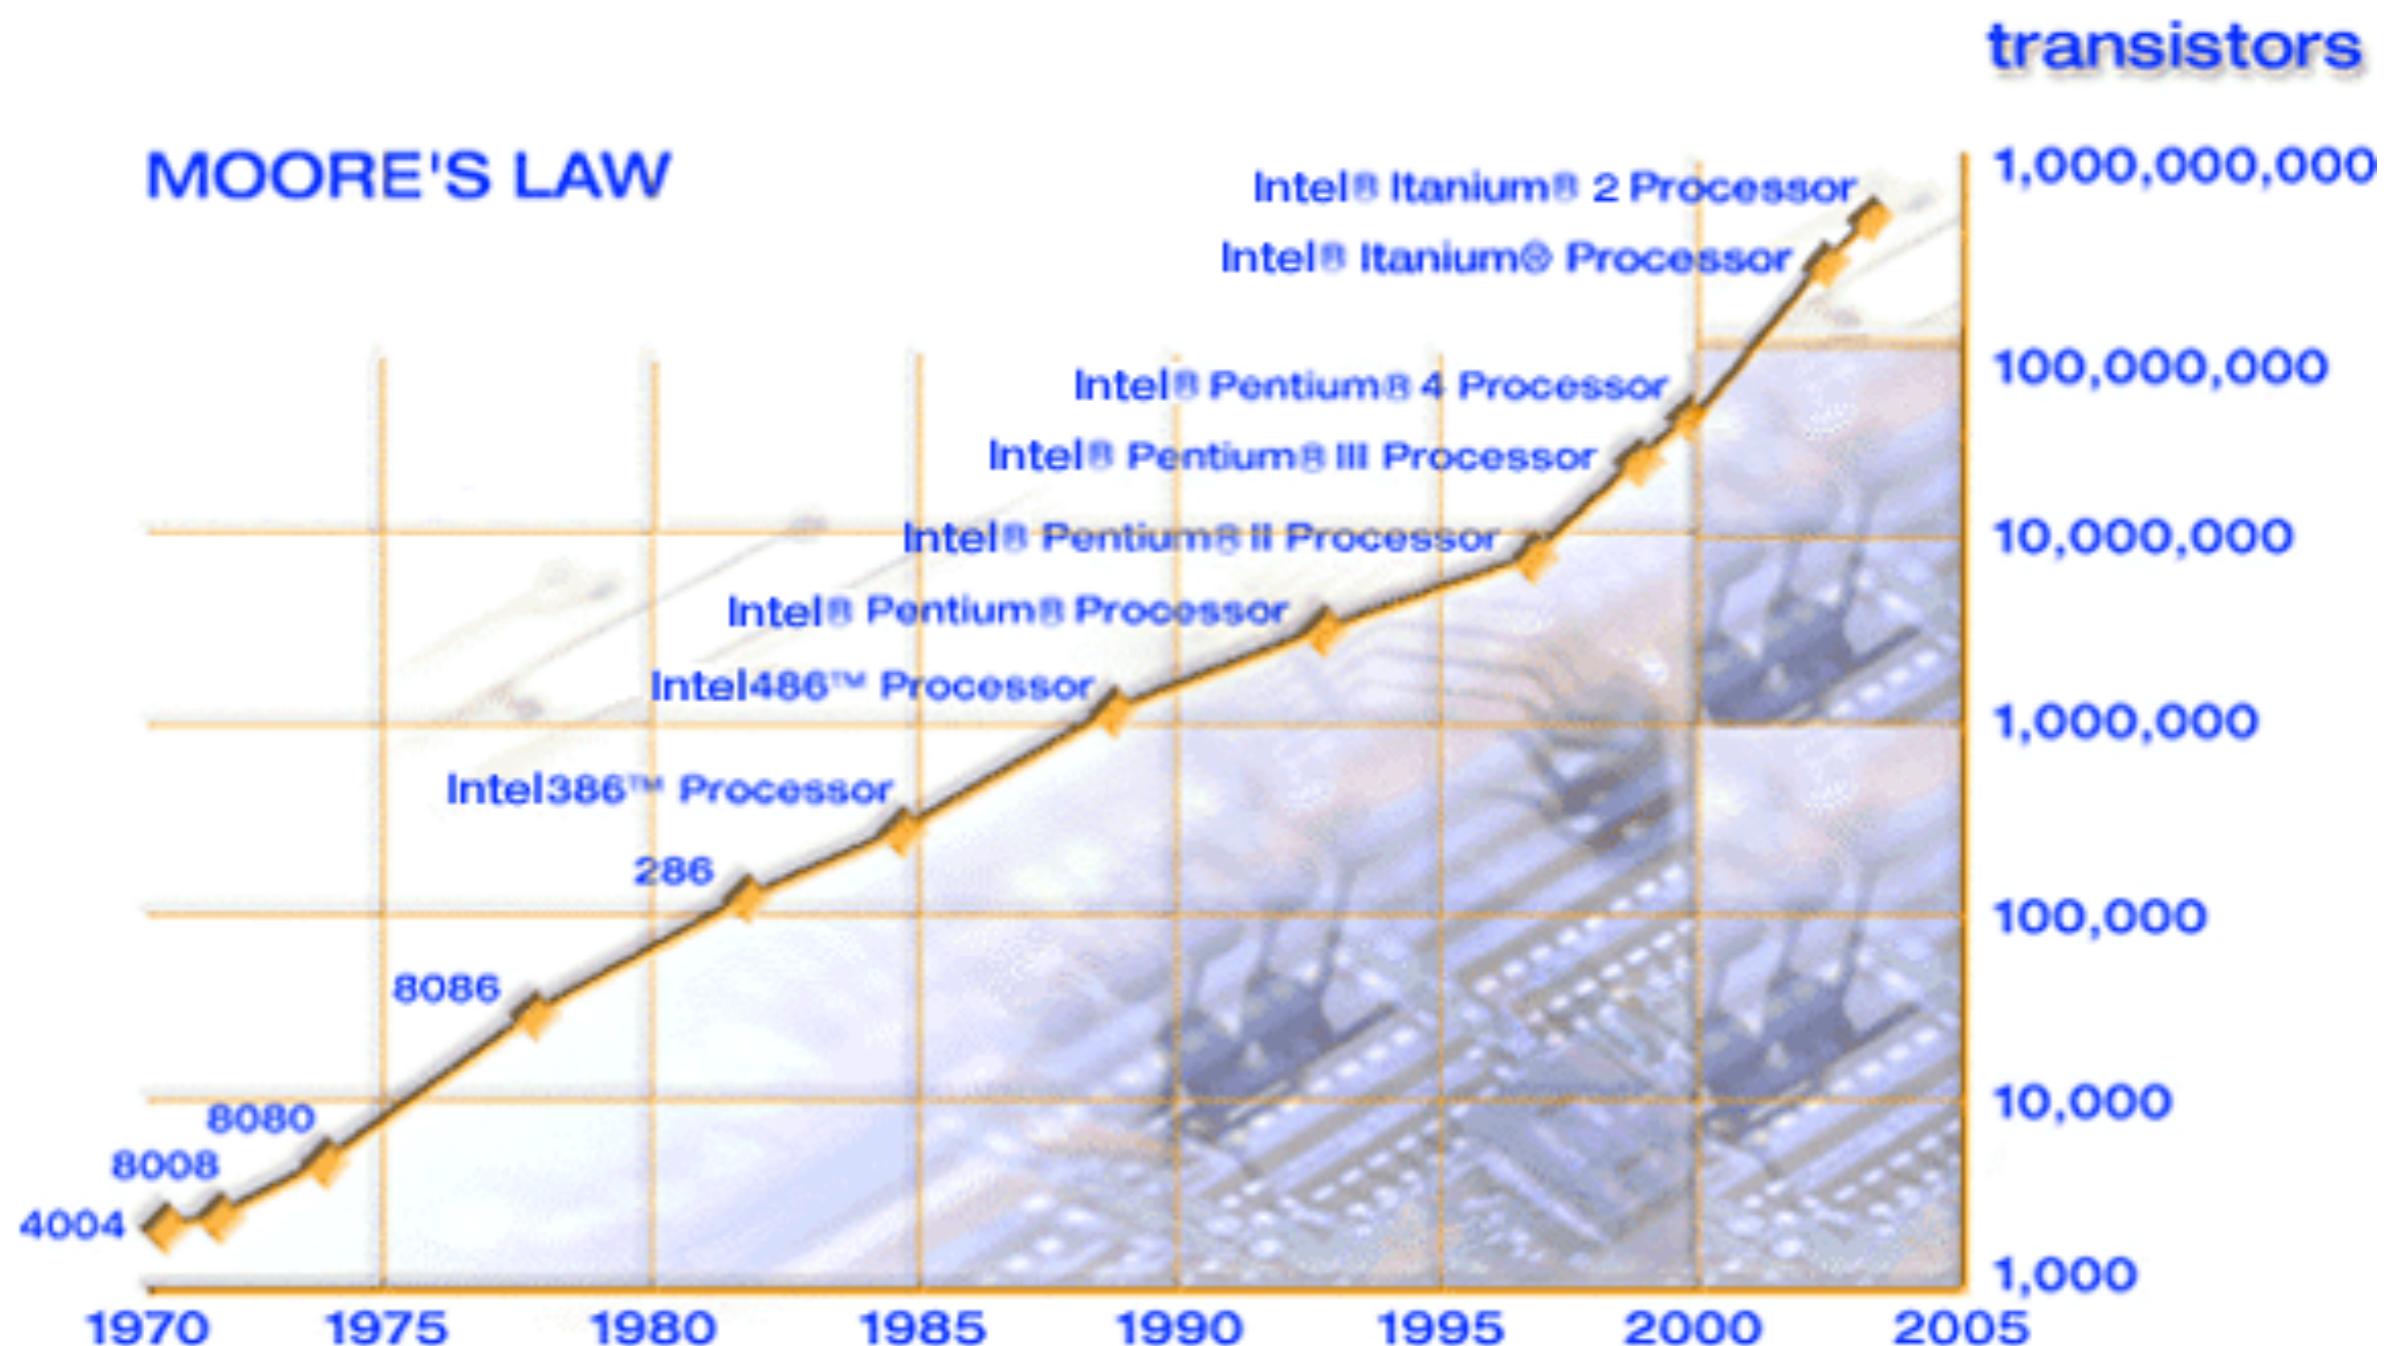
Data parallelism

- Declarative

IO parallelism

- Async IO
- Non Blocking IO
- Blocking IO

The Free Lunch is Over



Unfortunately adding transistors stopped translating into faster processors, primarily because of power consumption and the heat generated.

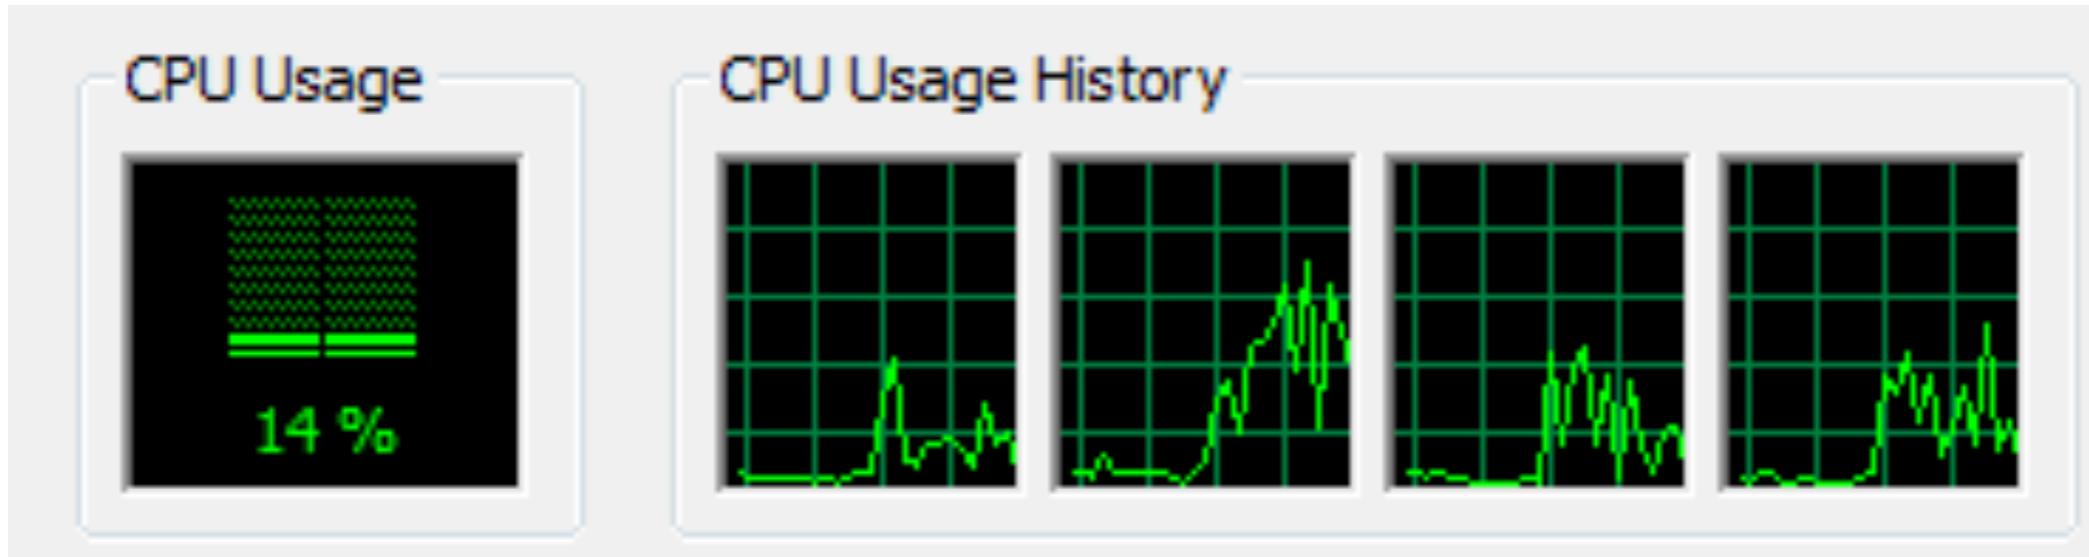
Example

```
private static void doWork(int instance)
{
    double result = Math.Acos(new Random().NextDouble()) *
Math.Atan2(new Random().NextDouble(), new Random().NextDouble());

    for (int i = 0; i < 20000; i++)
    {
        result += (Math.Cos(new Random().NextDouble()) *
            Math.Acos(new Random().NextDouble()));
    }
}
```

Single-Threaded Programming

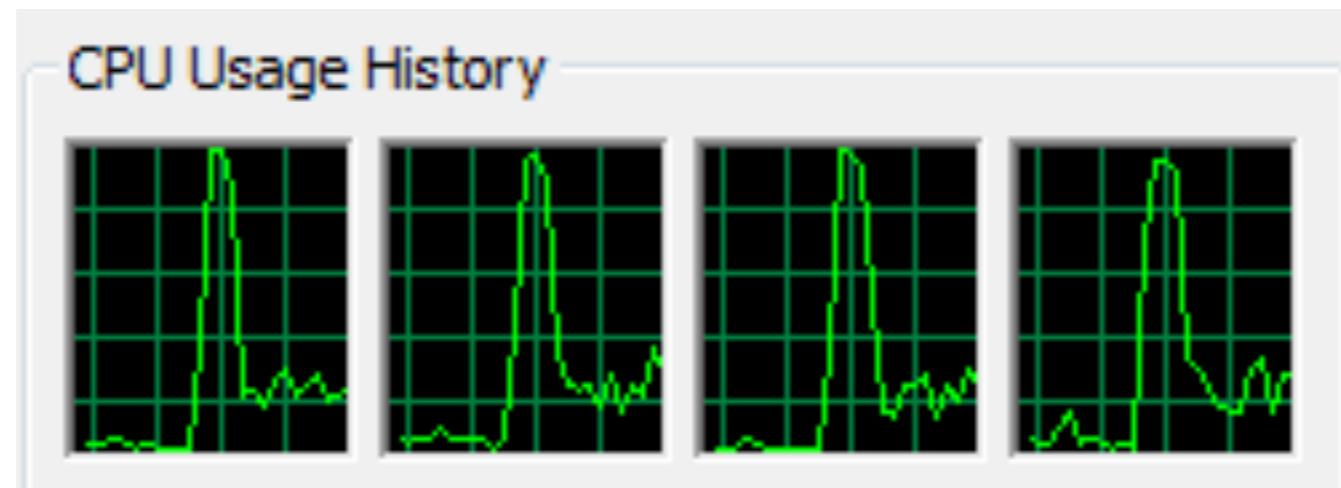
```
for (int i = 0; i < 100; i++)  
{  
    doWork(i);  
}
```



With four cores : over 12 seconds

Parallel Programming

```
//for (int i = 0; i < 100; i++)  
  
Parallel.For(0, 100,  
    delegate(int i){      doWork(i);          } );
```



The process finishes in 3.7 seconds instead of 12.

Parallel Programming

Individual tasks
may take a
little longer

```
0 took 145 milliseconds
1 took 124 milliseconds
2 took 121 milliseconds
3 took 125 milliseconds
4 took 122 milliseconds
5 took 124 milliseconds
6 took 122 milliseconds
7 took 123 milliseconds
8 took 121 milliseconds
9 took 121 milliseconds
10 took 122 milliseconds
11 took 124 milliseconds
12 took 123 milliseconds
13 took 122 milliseconds
14 took 122 milliseconds
```

Single-threaded

```
4 took 176 milliseconds
5 took 173 milliseconds
3 took 198 milliseconds
2 took 261 milliseconds
6 took 246 milliseconds
7 took 286 milliseconds
1 took 355 milliseconds
0 took 385 milliseconds
8 took 195 milliseconds
12 took 210 milliseconds
16 took 231 milliseconds
10 took 329 milliseconds
20 took 174 milliseconds
18 took 244 milliseconds
```

Parallel

some as high as
400-
milliseconds

Foreground Threads

vs

Background Threads



Try to avoid using foreground threads
as much as possible



When you really want to complete, like flushing data from a memory buffer out to disk.

Cost of a Memory

- x86, 700 bytes.
- x64, 1,240 bytes.
- IA64 CPUs, 2,500 bytes



Cost of a Thread

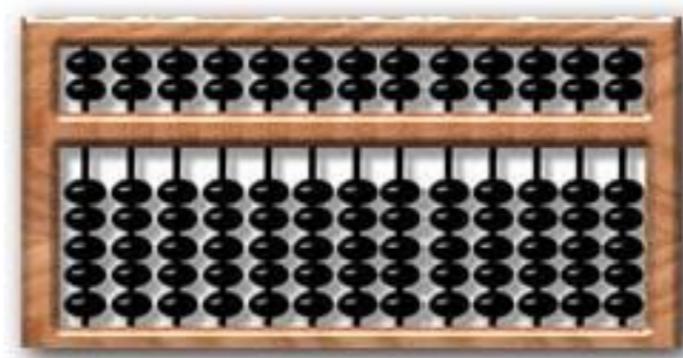
- 200,000 cycles to create a thread
- 100,000 to destroy one
- context switch costs 2,000-8,000 cycles



Don't create Threads

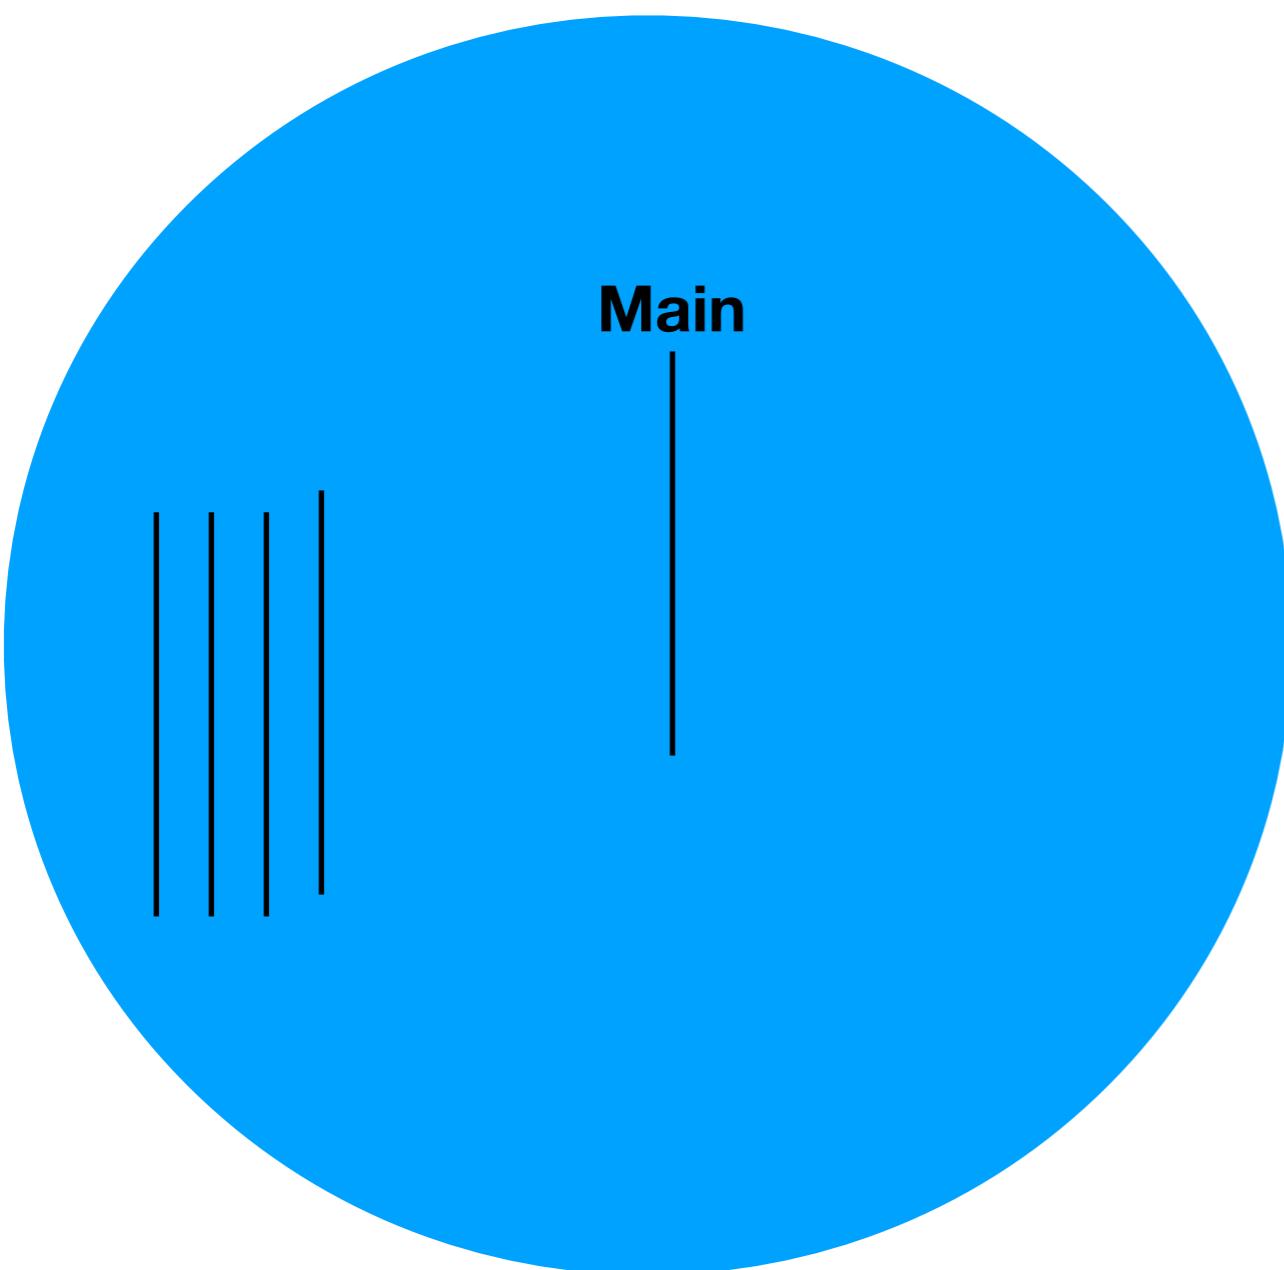
Total Threads

- A 32-bit process has 2 GB.
- After Win32 DLLs load, the CLR DLLs load, the native heap and the managed heap is allocated, there is approximately 1.5 GB left.
- Each thread requires more than 1 MB of memory for its user-mode stack and thread environment block (TEB)
- The most threads you can get in a 32-bit process is about 1,360.





What is the Ideal Number of Threads ?



Main





CARPOOLING

Because Gas is Just Too Expensive

When To Create Thread



Do you require a foreground thread

When To Create Thread



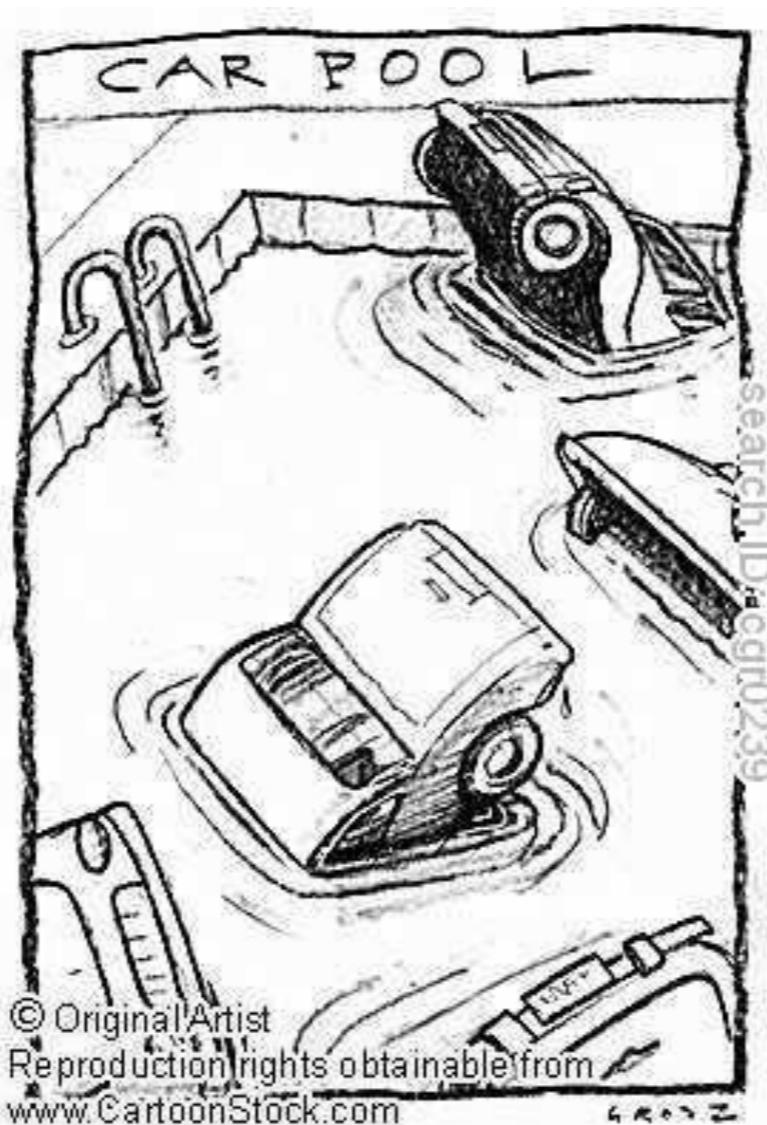
Change priority

When To Create Thread



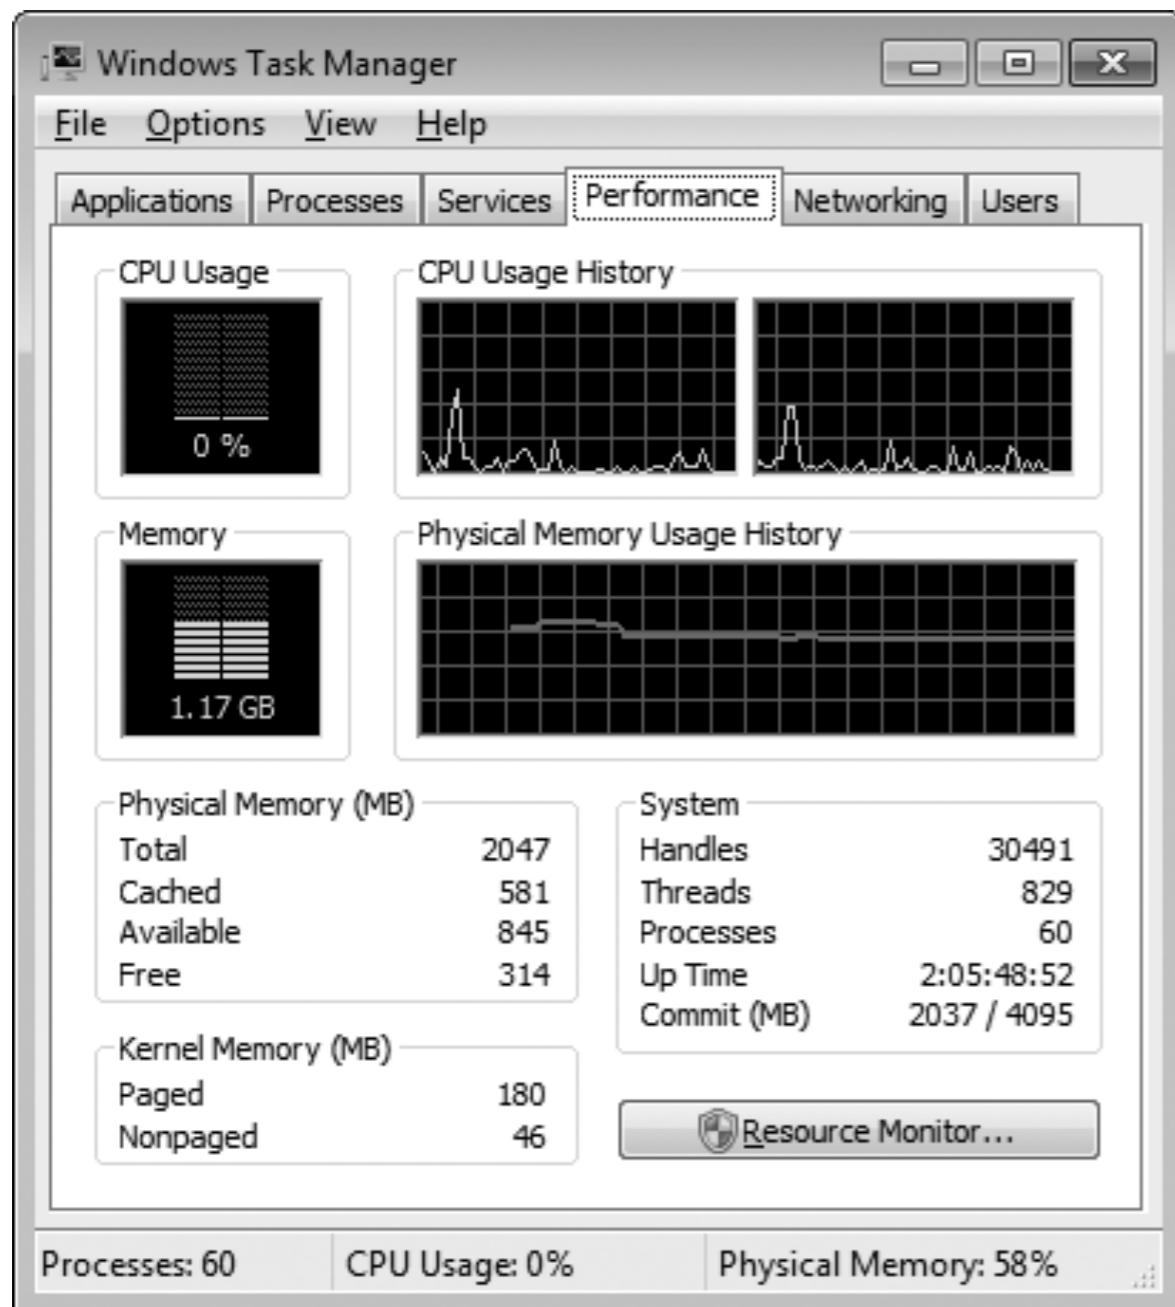
Long running

When To Create Thread



You need to have a stable identity associated with the thread

Do these applications need all these threads



Windows Task Manager

File Options View Help

Applications Processes Services Performance Networking Users

Image Name	CPU	Memory (K)	Threads	Description
OUTLOOK.EXE *32	00	86,392 K	38	Microsoft ...
devenv.exe *32	00	98,048 K	34	Microsoft ...
msnmsgr.exe *32	00	28,612 K	34	Windows ...
explorer.exe	00	24,336 K	32	Windows ...
iexplore.exe *32	00	44,388 K	21	Internet ...
WINWORD.EXE *32	00	34,888 K	13	Microsoft ...
MagicMouse.exe	00	7,944 K	13	
wlcomm.exe *32	00	6,208 K	11	Windows ...
csrss.exe	00	3,216 K	11	
iexplore.exe *32	00	3,716 K	10	Internet ...
msvsmon.exe	00	6,124 K	10	Visual Stu...
PushPin.exe	01	4,564 K	10	PushPin
MagicMouse.exe	00	3,948 K	10	
mspaint.exe	00	18,792 K	9	Paint

Show processes from all users End Process

Processes: 60 CPU Usage: 9% Physical Memory: 59%

This screenshot shows the Windows Task Manager's Processes tab. It displays a table of running processes with the following columns: Image Name, CPU, Memory (K), Threads, and Description. The table lists various system and application processes. At the bottom are buttons for 'Show processes from all users' and 'End Process'. Status bars at the bottom show 'Processes: 60', 'CPU Usage: 9%', and 'Physical Memory: 59%'.

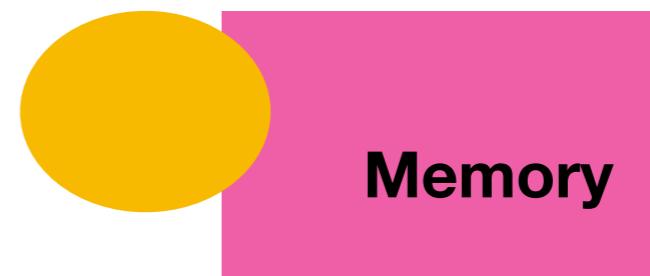


Playing with thread pool limits usually results in making an application perform worse, not better

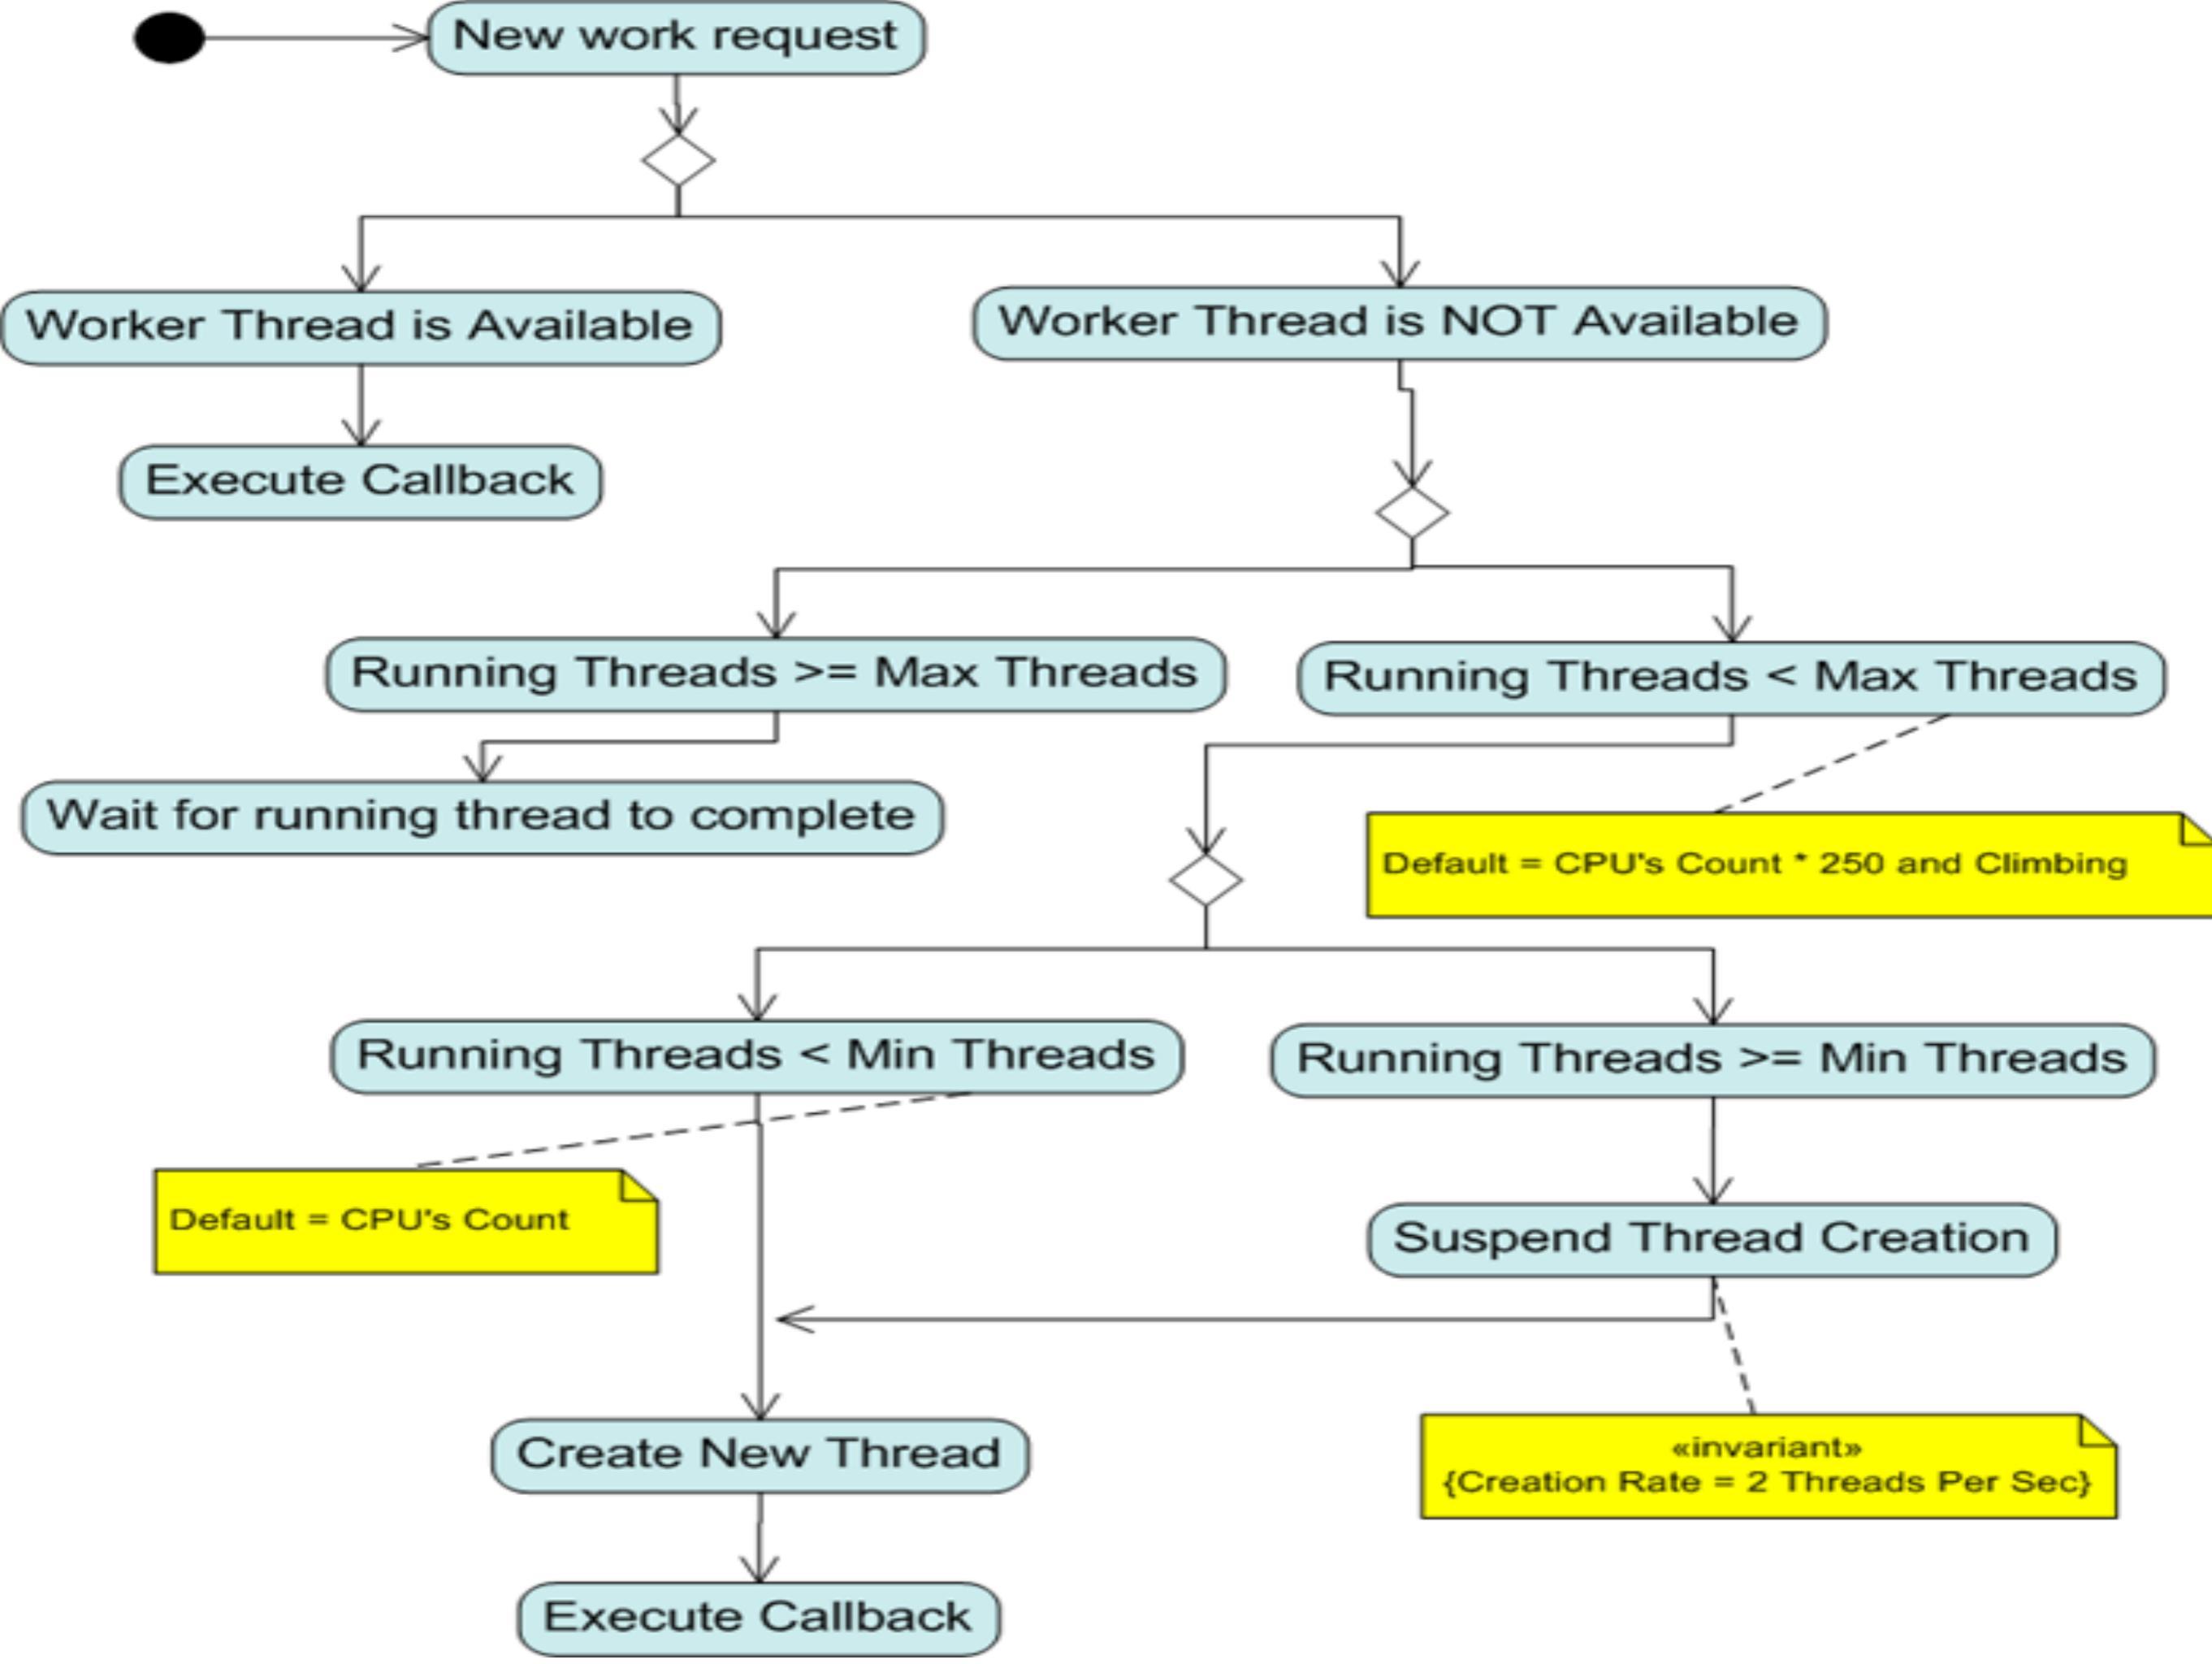
Max : 5 threads

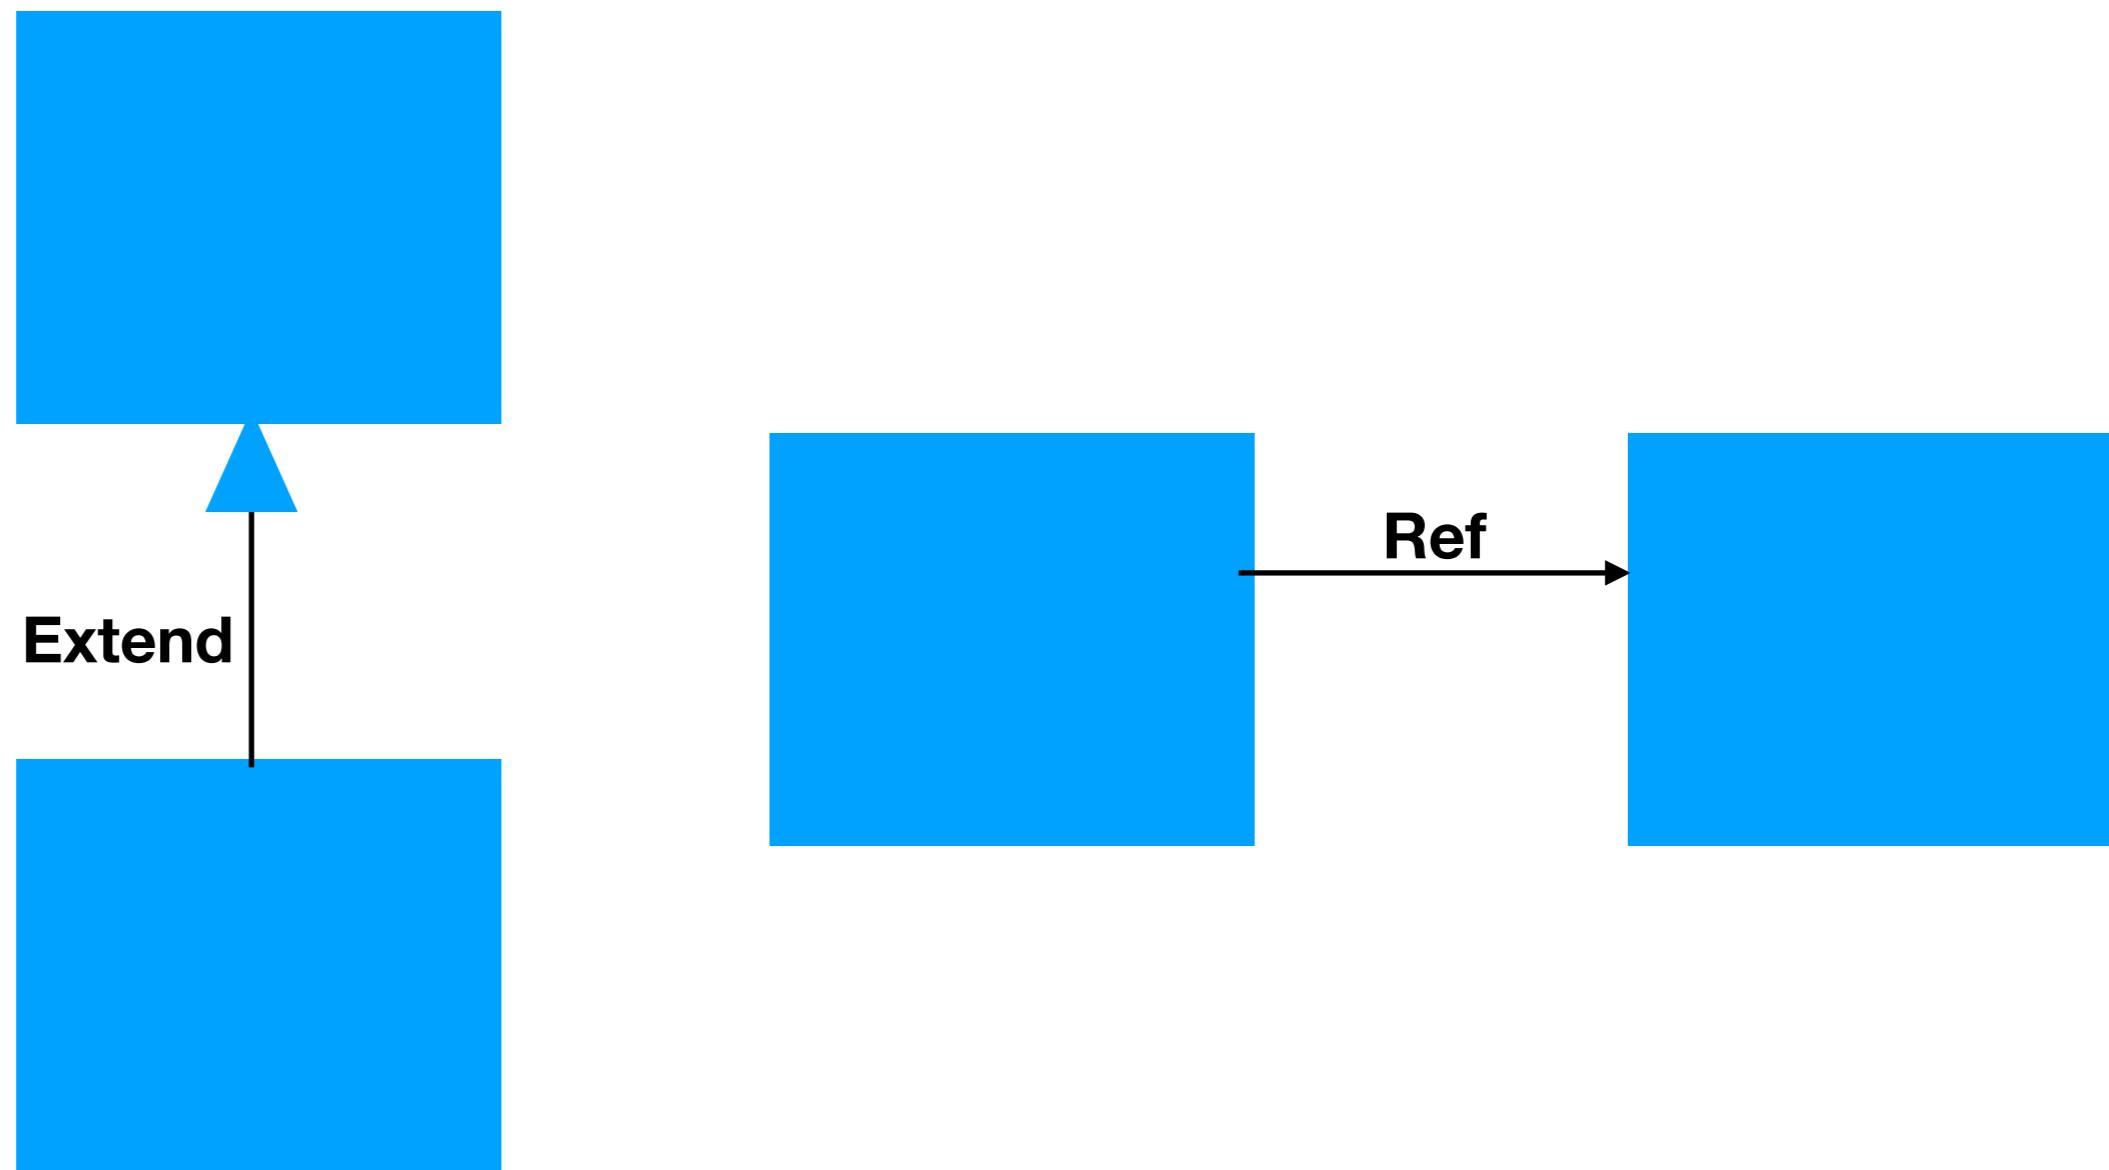
executor.execute()
Memory reader

executor.execute()
Memory writer



executor.execute()
executor.execute()
executor.execute()
executor.execute()
executor.execute()
executor.execute()





3 Approaches to Parallelism

1. Imperative data parallelism

- Applicable to common data oriented operations such as for and foreach loops.

2. Imperative task parallelism

- For performing parallel tasks via expressions and statements.

3. Declarative data parallelism

- Applicable while querying data.

job1(){ logic 1 }
job2(){ logic 2 }
job3(){ logic 3 }

data1, data2, data2
job(data){ logic }

System.Threading

Parallel.For

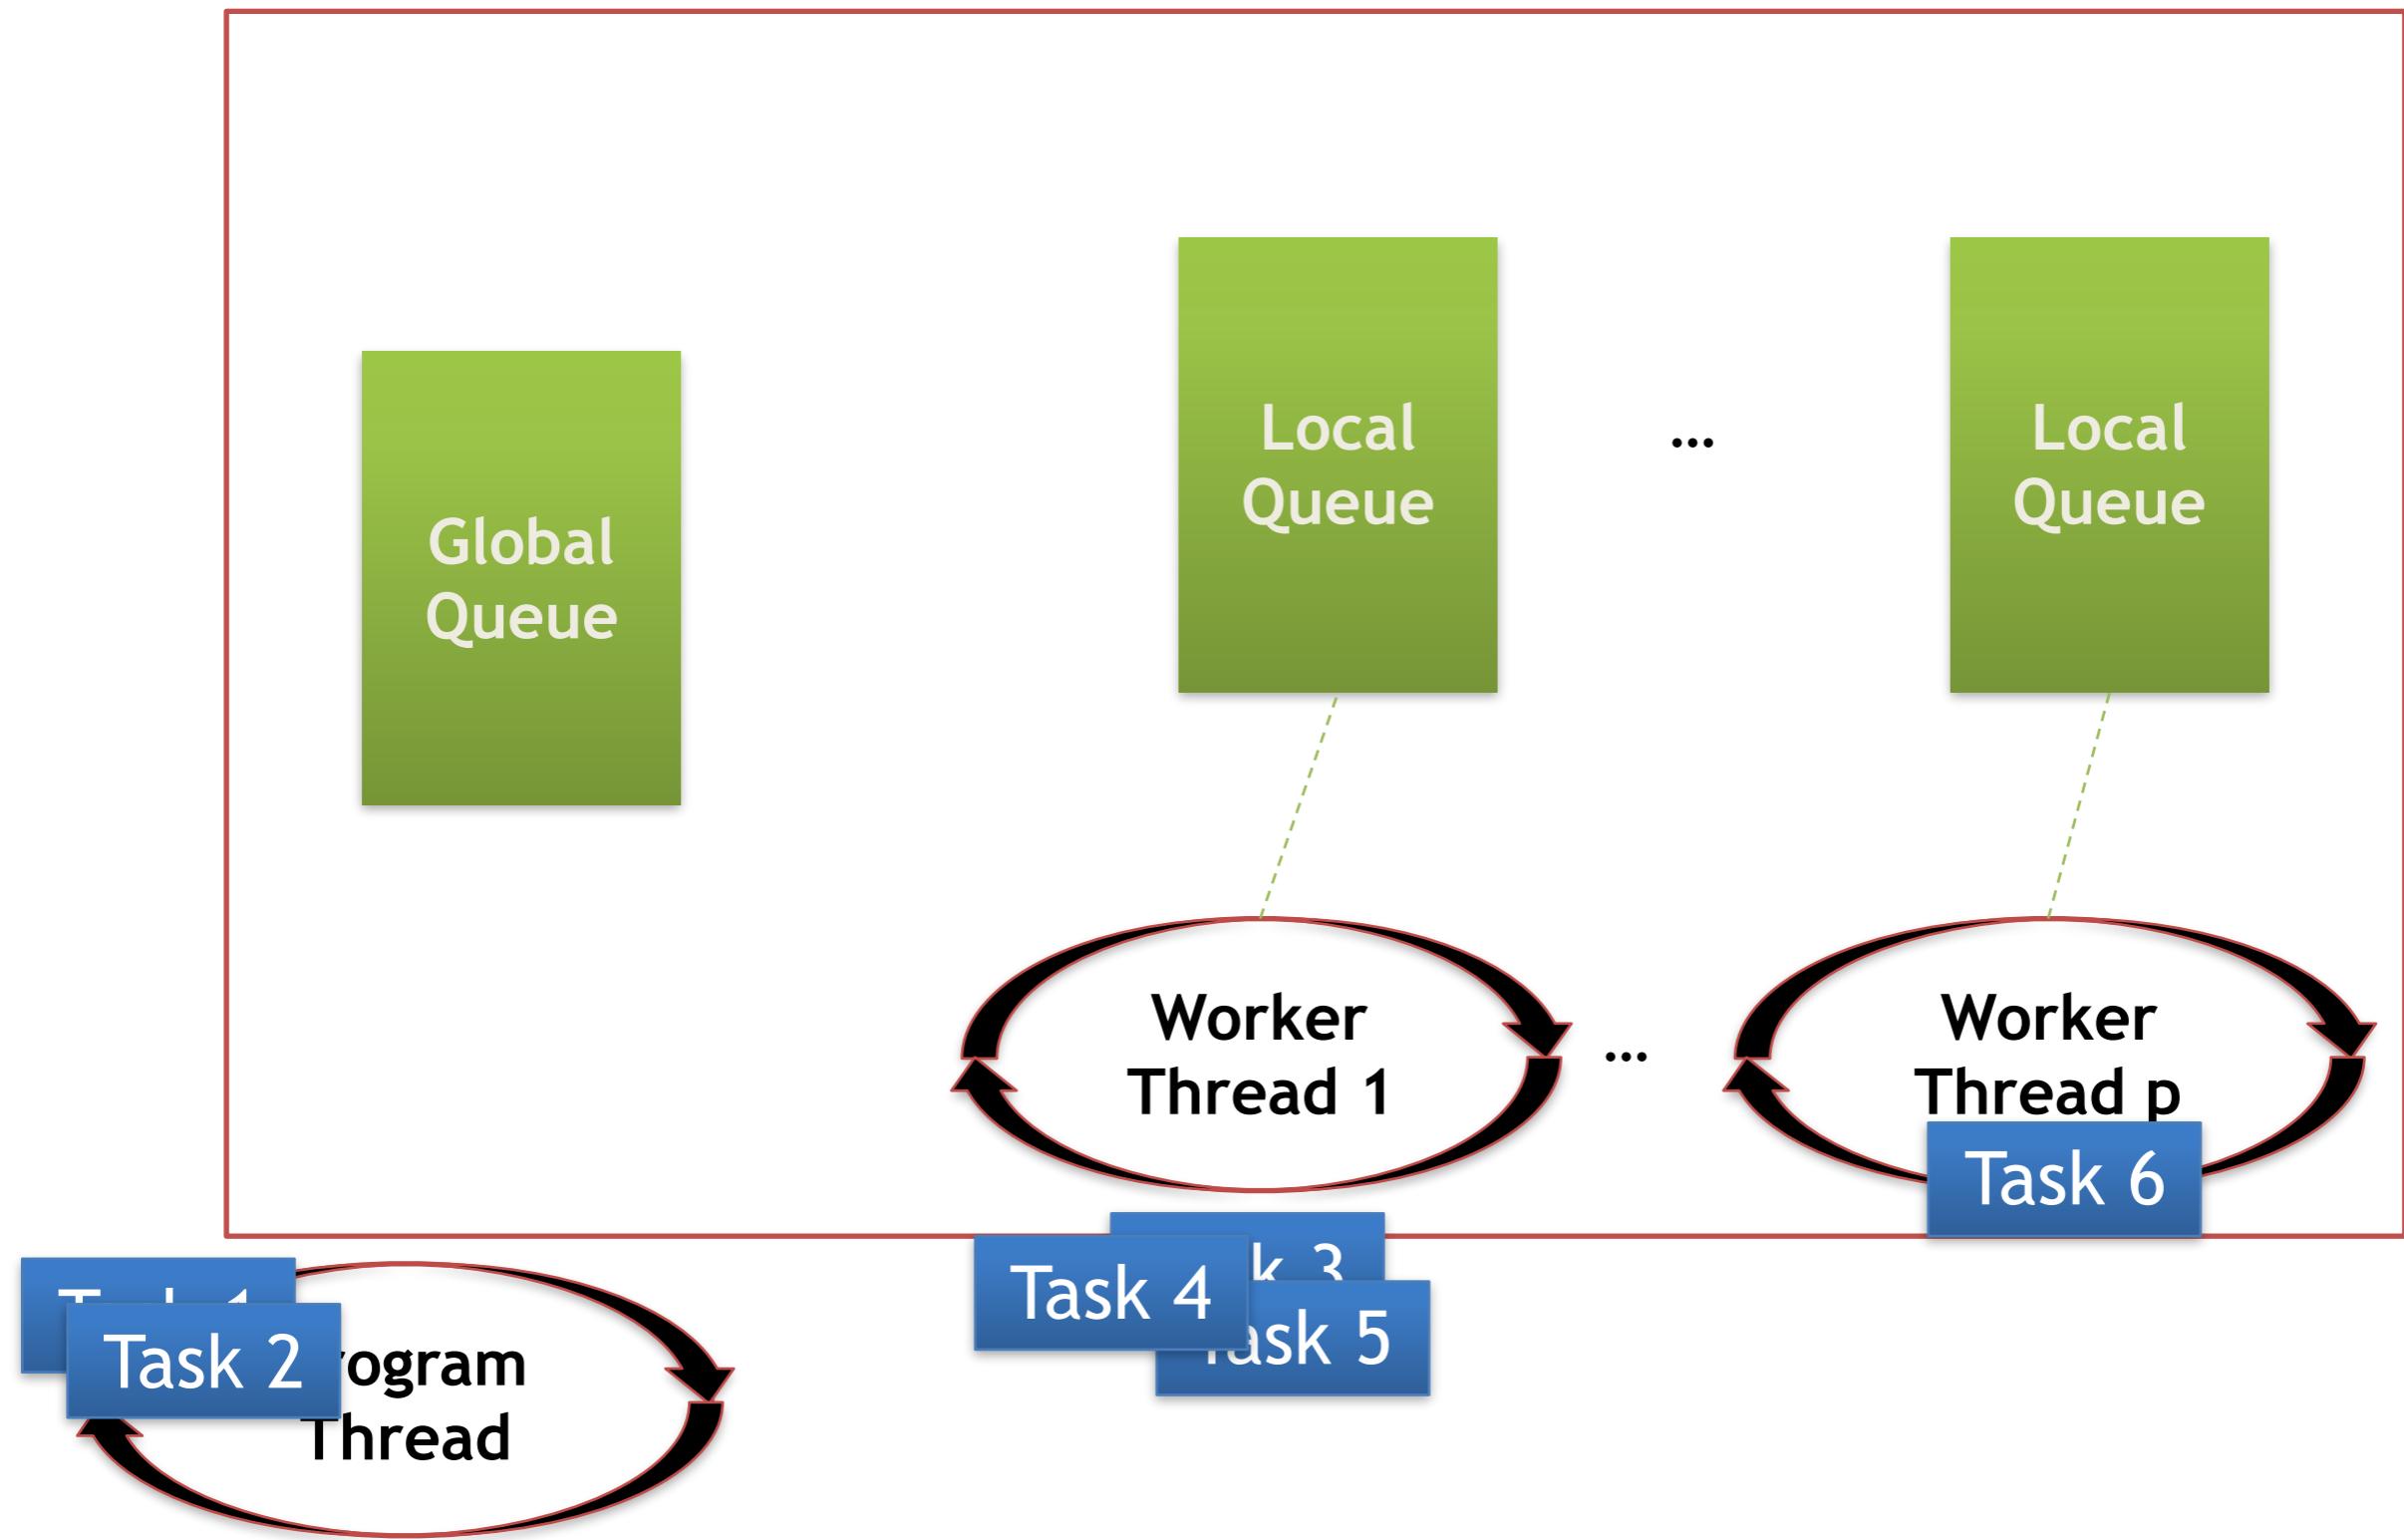
```
//for (int i = 0; i < 100; i++)
System.Threading.Parallel.For(0, 100,
    delegate(int i){ // parallel work here } );
```

Parallel.ForEach

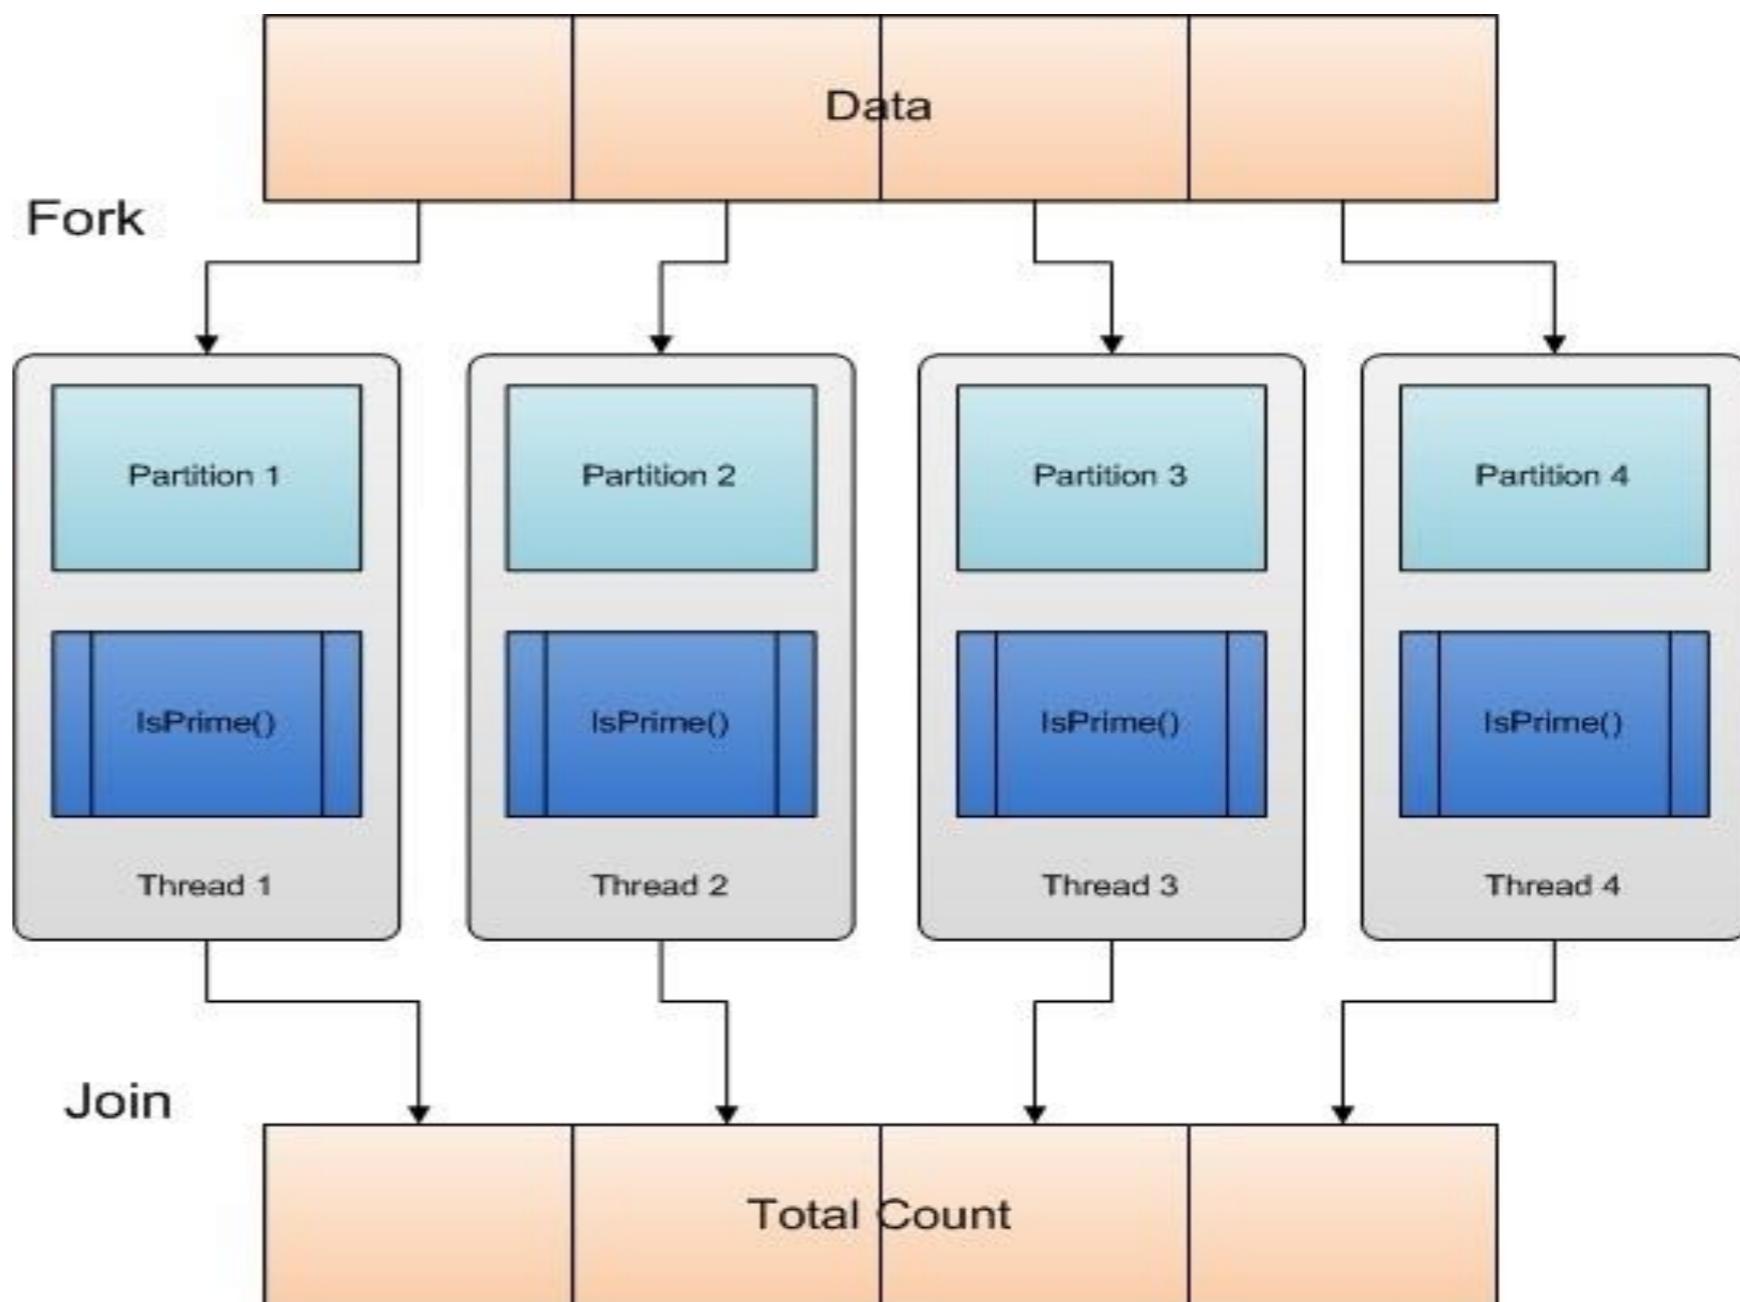
```
List<int> list = new List<int>();
Parallel.ForEach(list,
    item =>{ // parallel work here } );
```

Parallel.Invoke

```
Parallel.Invoke(A, B, C, D, E);
public static void A(){}
public static void B(){}
```



	Partitions work	Collates results
Parallel Stream	Yes	Yes
Thread class	No	No



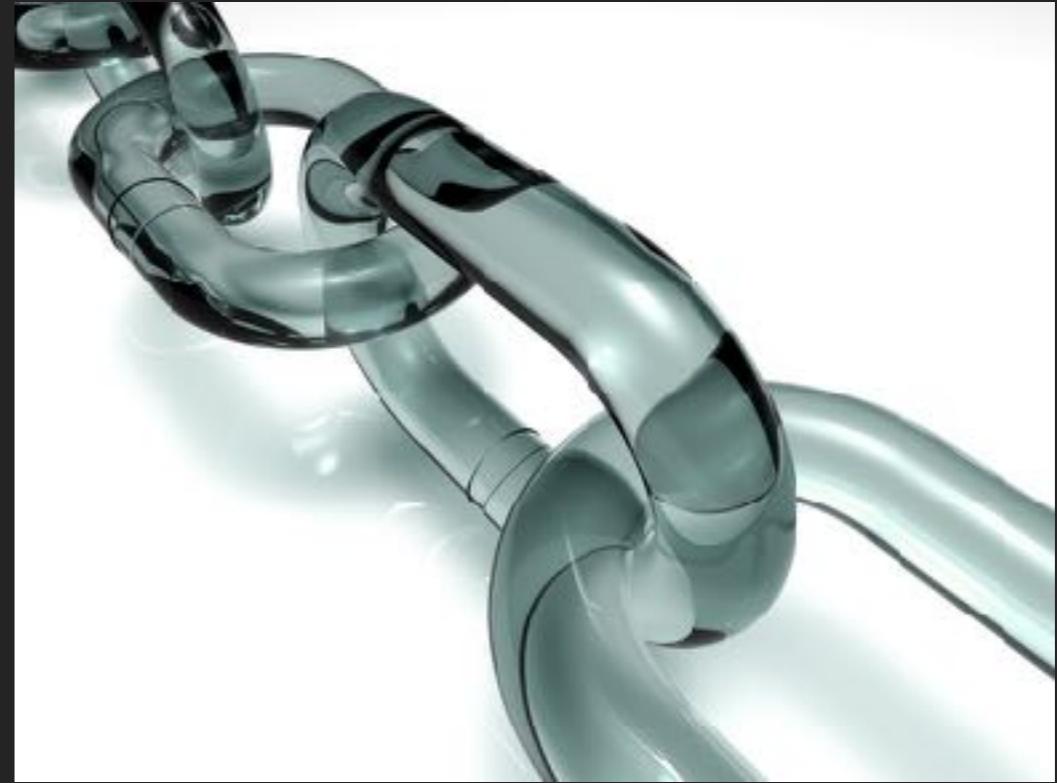
10000 employees

1 employee to each thread

100 employees to each thread

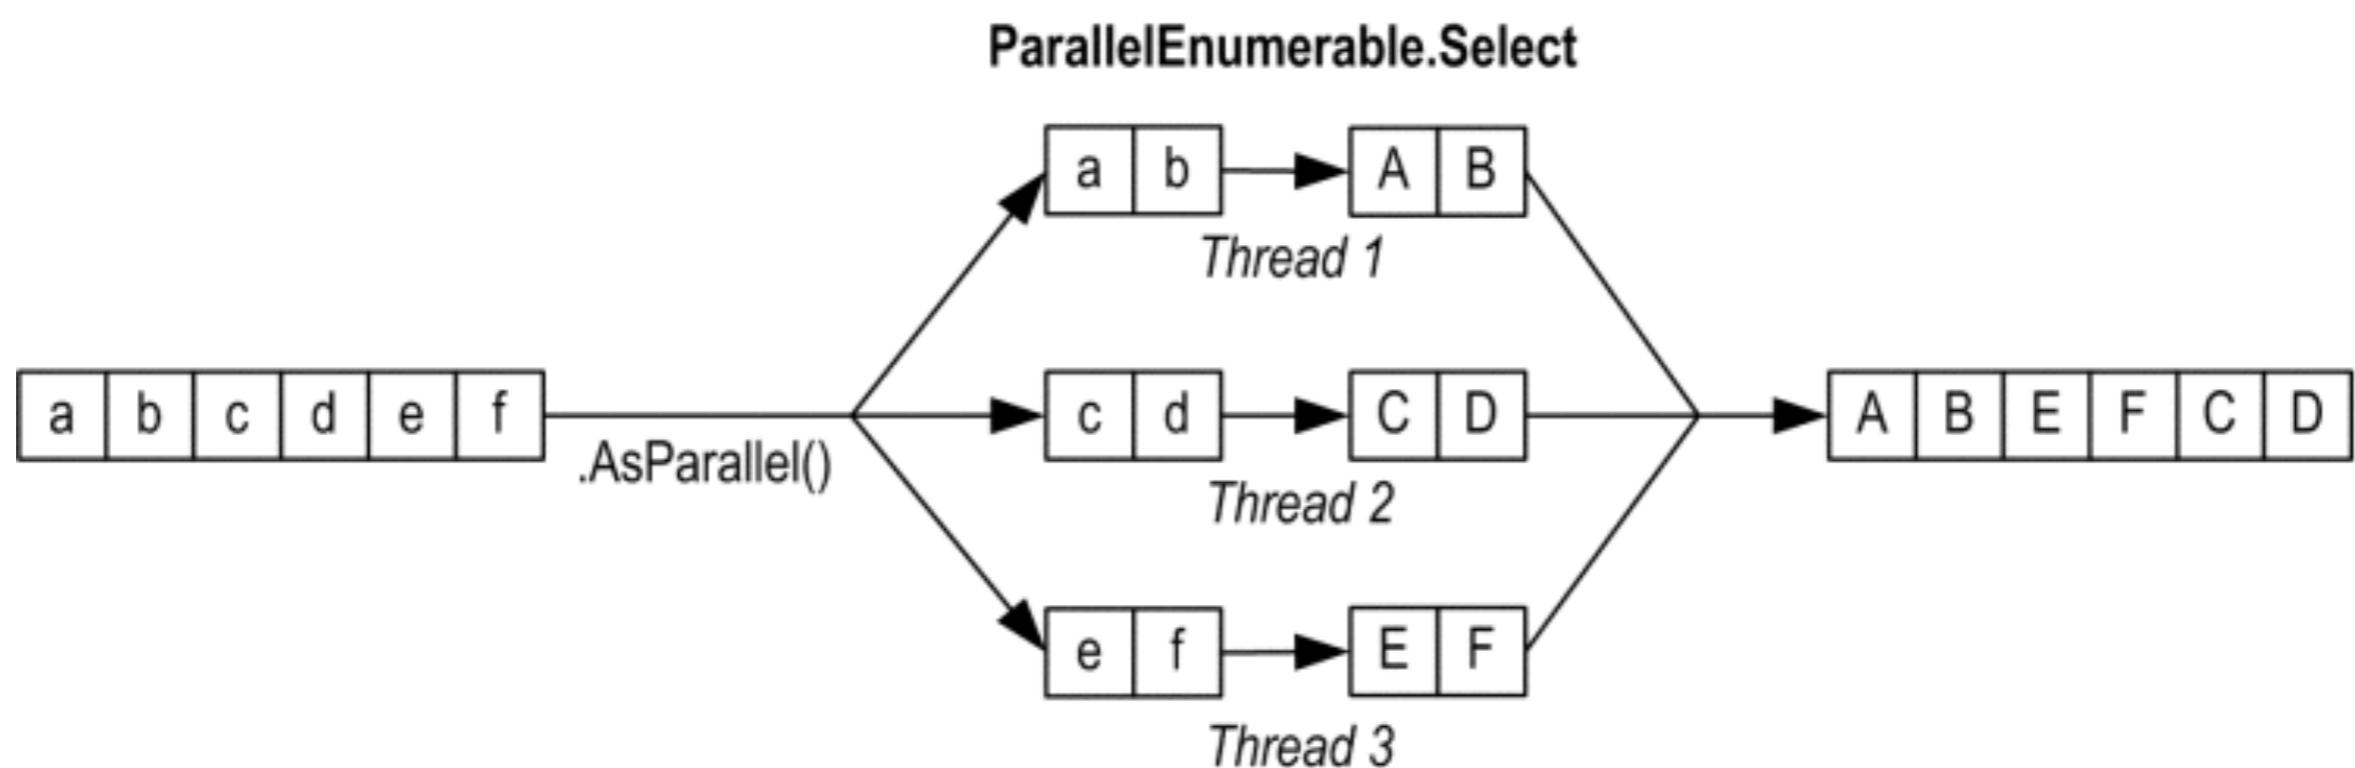
10000 employees in one thread

System.Threading



PLinq

```
var q = from p in people.AsParallel()
        where p.Name == queryInfo.Name &&
              p.State == queryInfo.State &&
              p.Year >= yearStart &&
              p.Year <= yearEnd
        orderby p.Year ascending
        select p;
```



```
"abcdef".AsParallel().Select (c => char.ToUpper(c)).ToArray()
```

Collection<10,000 employees>

```
#dedicated thread  
for(Employee e: employees){  
    Thread thread = new Thread(()-> save(e));  
    Thread.start();  
}  
10,000 threads  
Partition : 1 employee per thread
```

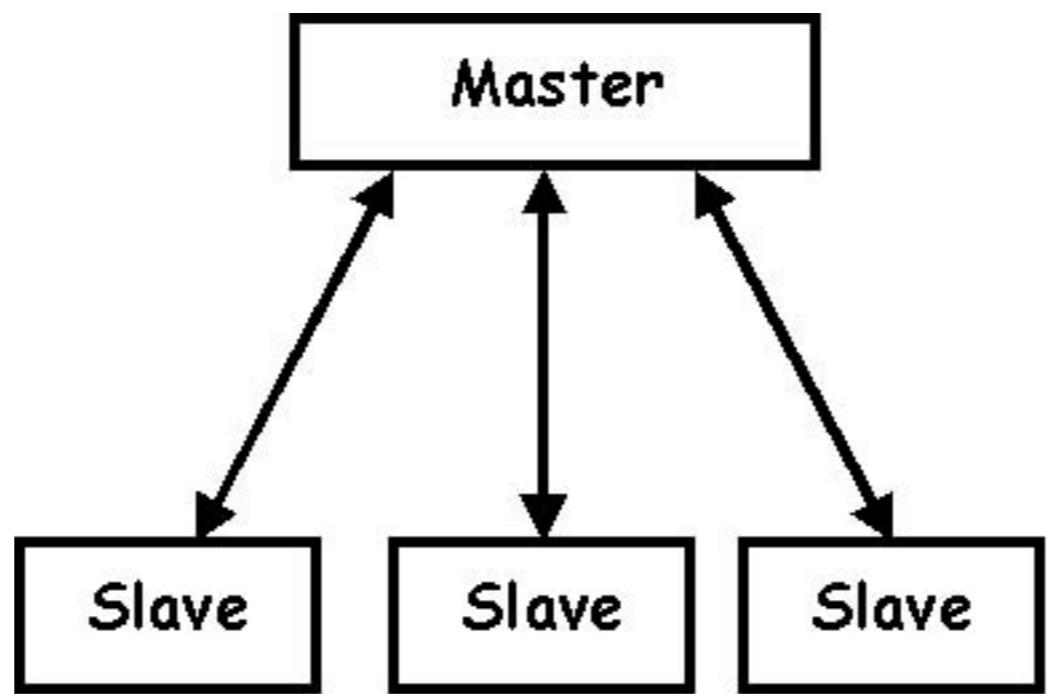
```
#threadpool
ExecutorService
executor=Executors.newFixedThreadPool(100);
for(Employee e: employees){
    executor.execute(()-> save(e));
}

```

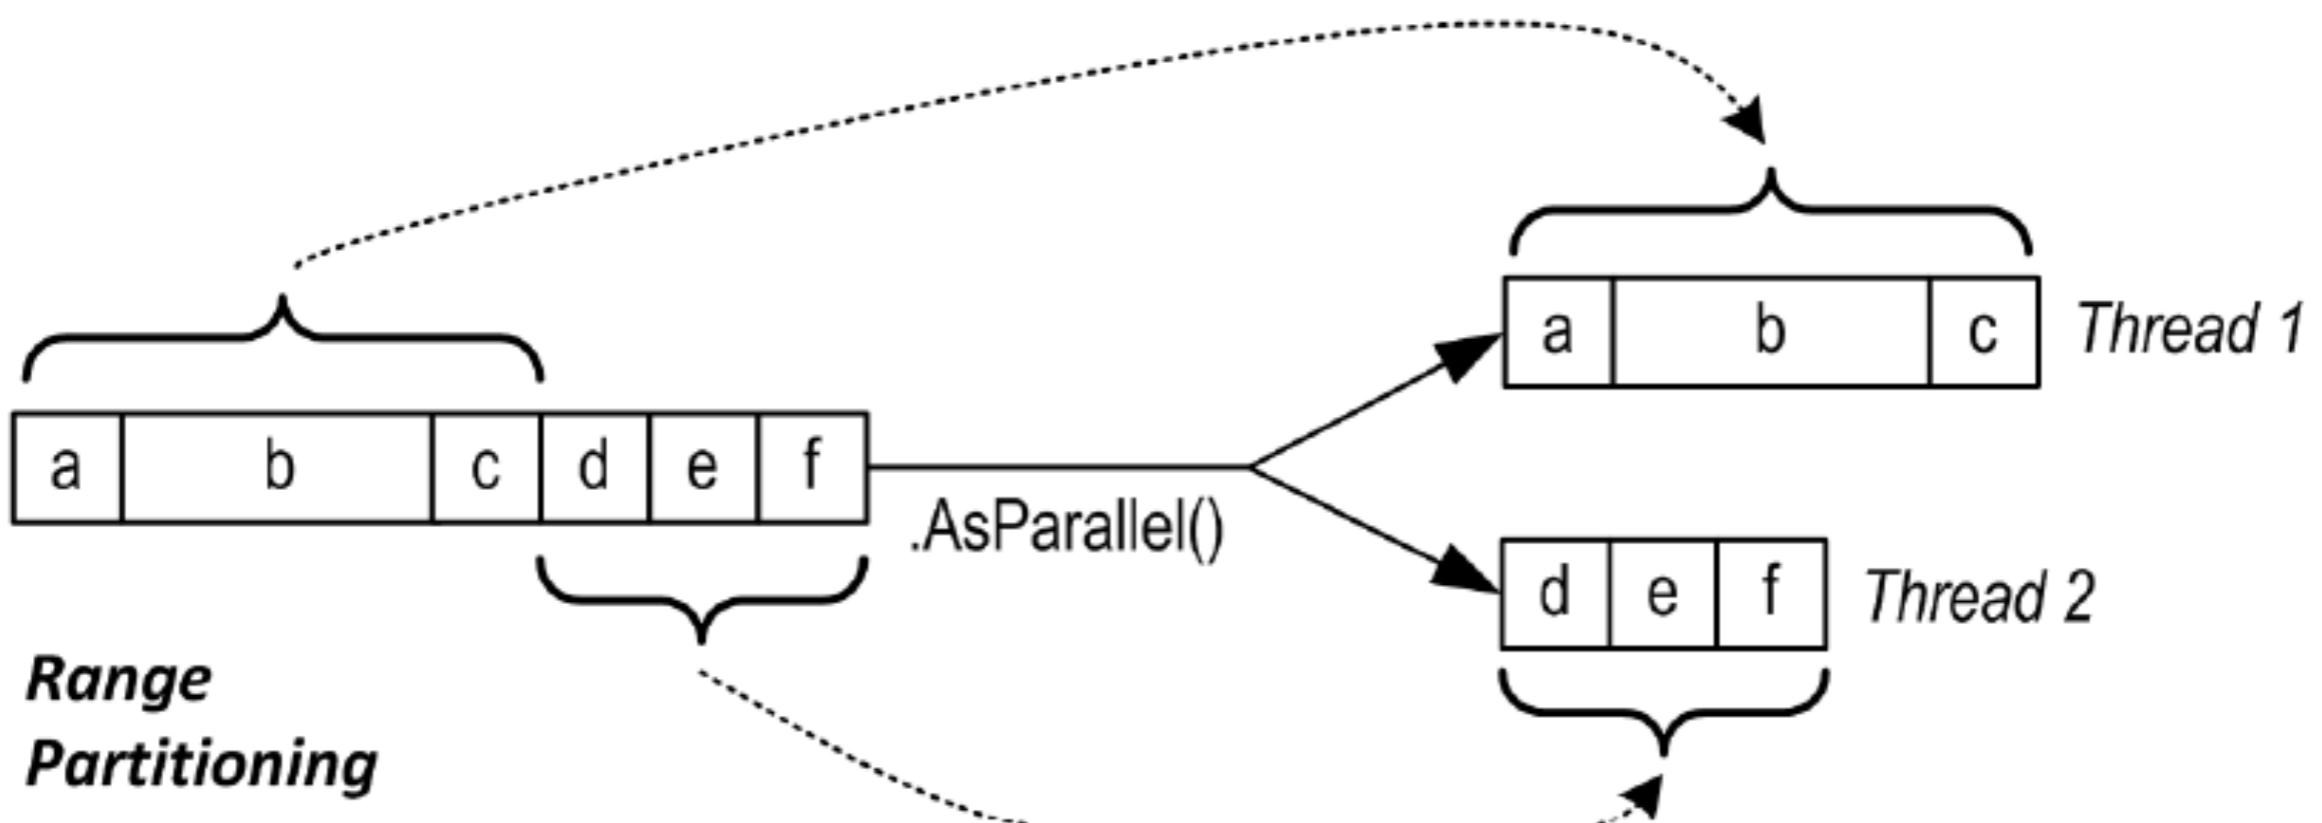
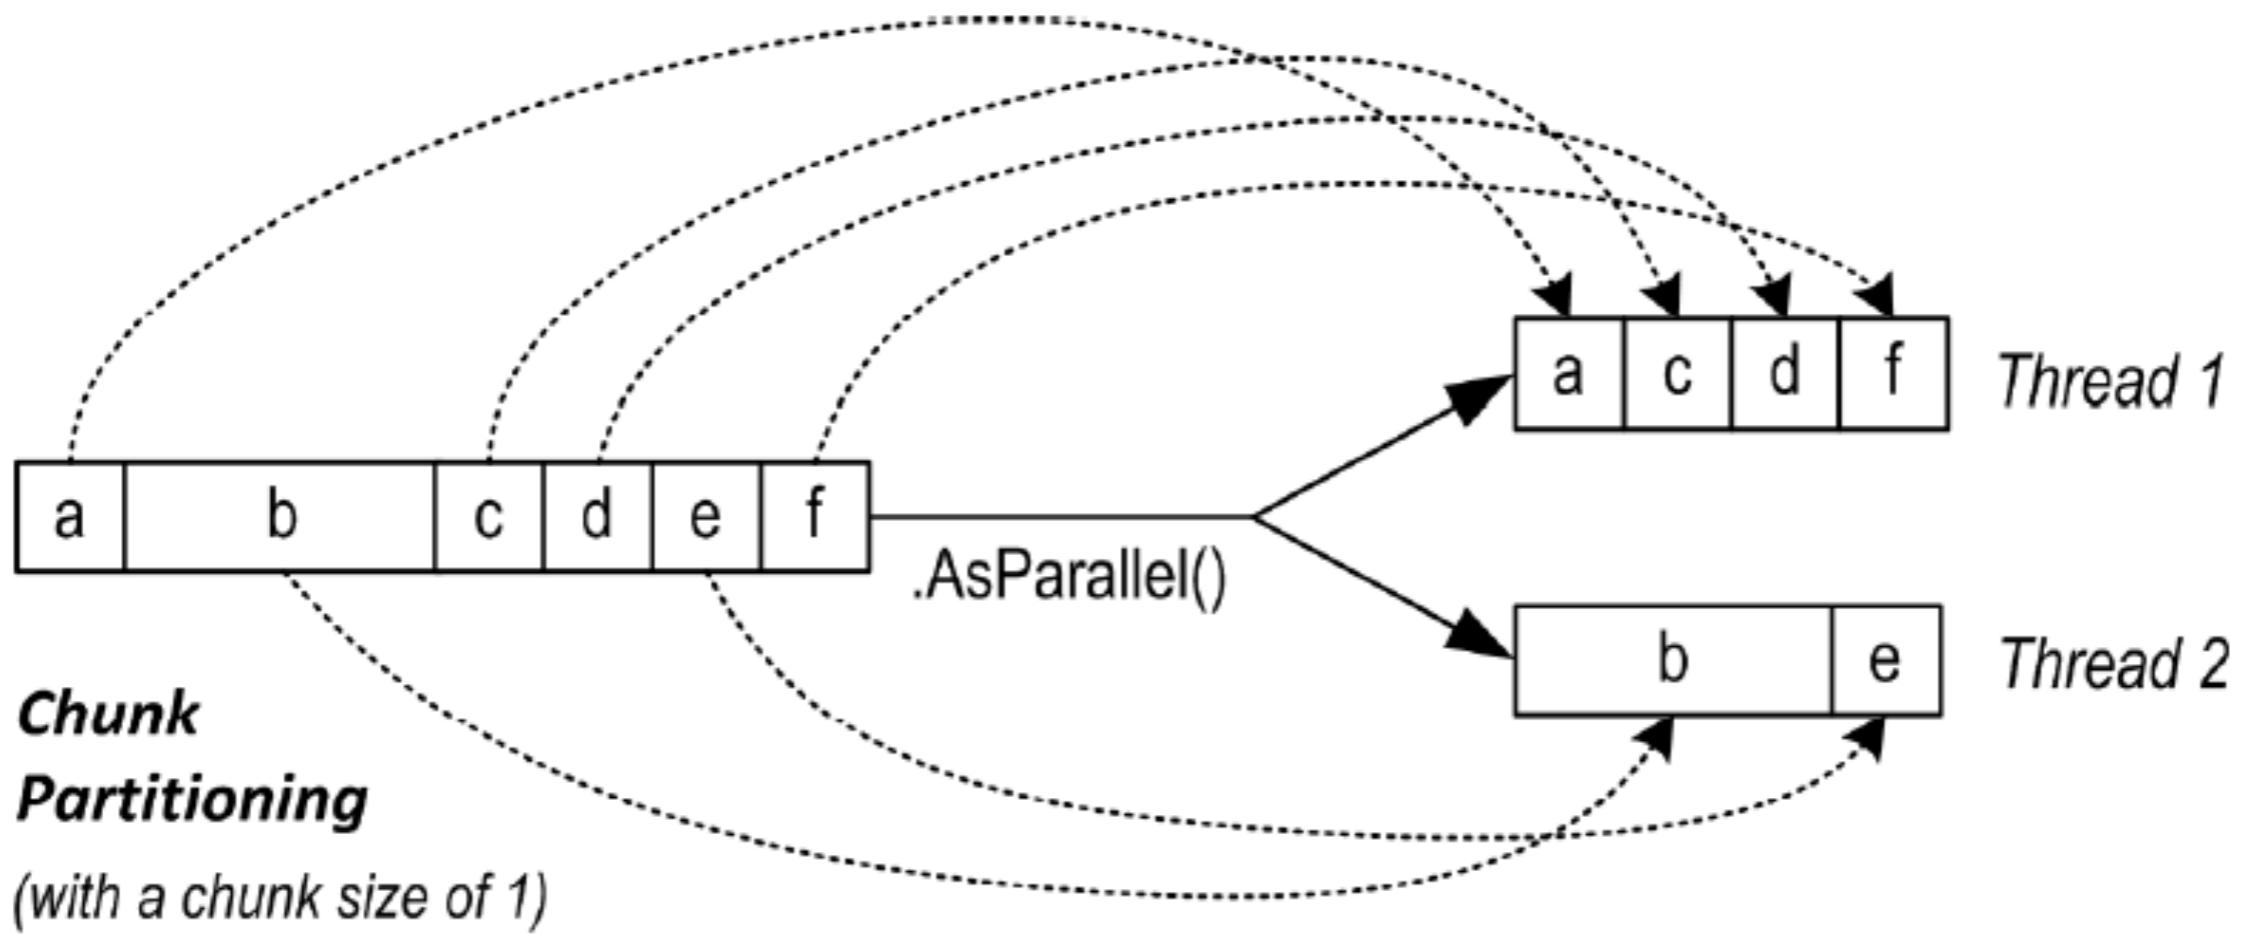
100 threads
partition: 1 employees per thread
Reused 10,000 times

```
#task  
employees.parallelStream().forEach(Main::save);
```

100 threads
partition: 10 employees per thread
Reused 1000 times



Why Isn't `.AsParallel()` the Default?



Partitioning Algorithms

- Several partitioning schemes built-in

- **Chunk**

- Works with any `IEnumerable<T>`
 - Single enumerator shared; chunks handed out on-demand



- **Range**

- Works only with `IList<T>`
 - Input divided into contiguous regions, one per partition



- **Stripe**

- Works only with `IList<T>`
 - Elements handed out round-robin to each partition



- **Hash**

- Works with any `IEnumerable<T>`
 - Elements assigned to partition based on hash code



- Custom partitioning available through `Partitioner<T>`

- `Partitioner.Create` available for tighter control over built-in partitioning schemes

Partitioning Strategy	Partitioning Strategy	Relative performance
<i>Chunk</i>	<i>Dynamic</i>	<i>Average</i>
Range	Static	Poor to excellent
<i>Hash</i>	<i>Static</i>	<i>Poor</i>
Stripe	Static	Poor to excellent
<i>Custom</i>	<i>Dynamic/static</i>	Poor to excellent

The cost of I/O

L1-cache	3 cycles
L2-cache	14 cycles
RAM	250 cycles
Disk	41 000 000 cycles
Network	240 000 000 cycles

Worker threads or I/O threads



Worker threads (Compute Bound/ CPU Bound)
I/O threads (Device Bound).

How Windows does Synchronous I/O

.NET

```
FileStream fs = new FileStream(...);  
Int32 bytesRead = fs.Read(...);
```

Win32
User-Mode

```
ReadFile(...);
```

Windows
Kernel-Mode

```
(Windows I/O Subsystem Dispatcher code)
```

①

⑨

③

⑧

④

⑦

Your thread blocks here!
Hardware does I/O;
No threads involved!

⑥

NTFS Driver

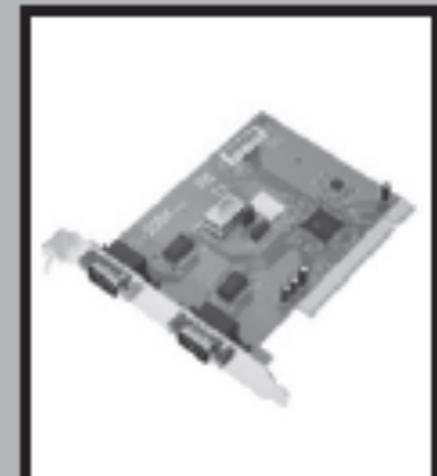
Network



DVD-ROM



RS-232



IRP
Queue



⑤

How Windows does Asynchronous I/O

.NET

```
FileStream fs = new FileStream(..., FileMode.Open);
fs.BeginRead(..., AsyncCallback, ...);
```

①

⑦

```
void AsyncCallback(...) { ... }
```

Win32
User-Mode

```
ReadFile(...);
```

③

IRP

⑥

Windows
Kernel-
Mode

```
(Windows I/O Subsystem Dispatcher code)
```

④

⑤

Your thread doesn't
block here; it
keeps running! ⑤

The CLR's Thread Pool

Threads can extract
completed IRP's
from here

NTFS Driver

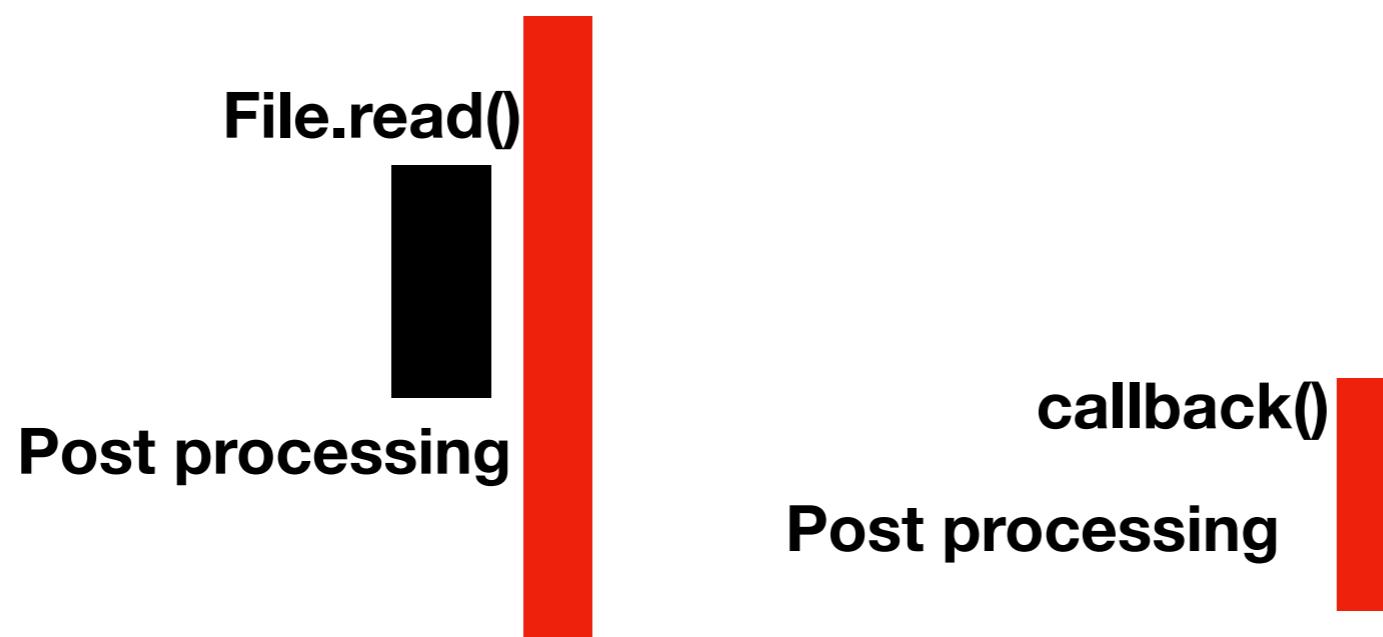
IRP
Queue



⑧

b

c



The diagram illustrates the execution flow of three client operations and a selector. On the left, four vertical bars represent the execution of **Client1.ConnectAsync()**, **Client2.ConnectAsync()**, **Client3.ConnectAsync()**, and **paysalary()**. These bars are colored blue, green, yellow, and red respectively. To the right of these bars, the word **Selector** is positioned above three vertical bars, each consisting of a black text label and a red rectangular bar. The first selector bar is labeled **callback()** and **Post processing**. The second selector bar is also labeled **callback()** and **Post processing**. The third selector bar is labeled **callback()** and **Post processing**.

Client1.ConnectAsync()

Client2.ConnectAsync()

Client3.ConnectAsync()

paysalary()

Selector

callback()

Post processing

callback()

Post processing

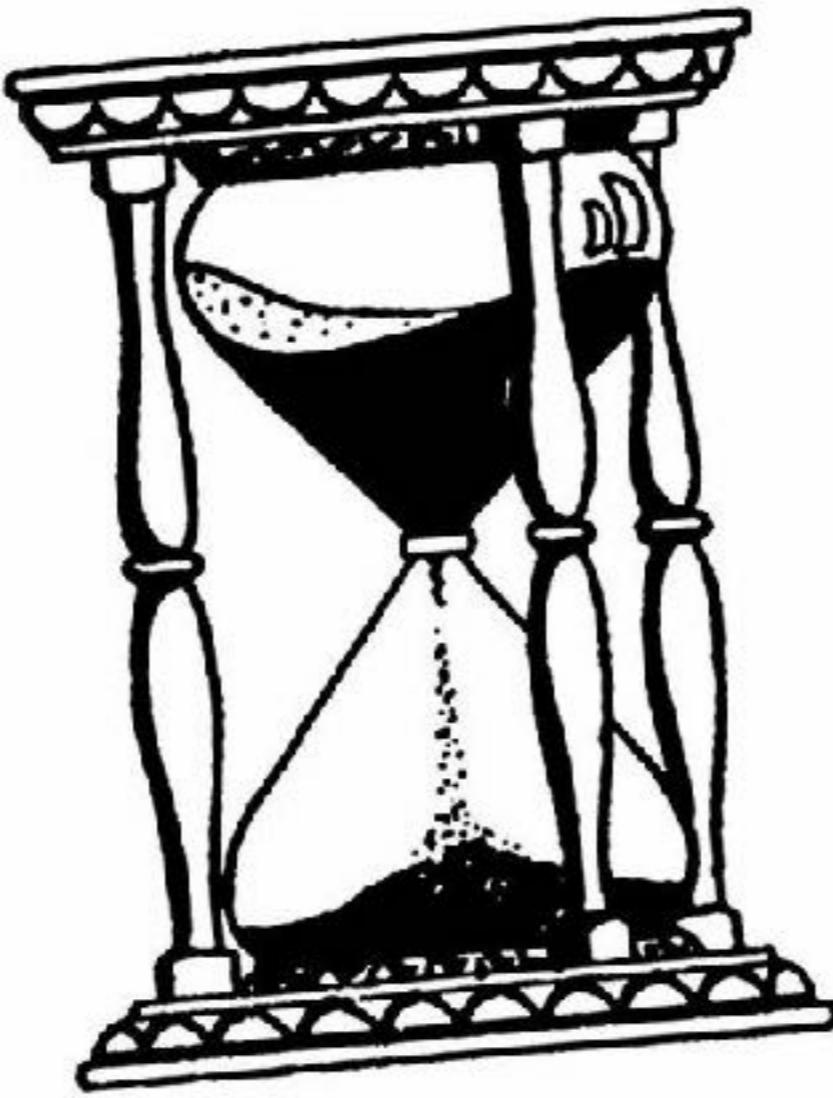
callback()

Post processing

- $a + b \rightarrow$ 3 cpu cycles
- Fun call \rightarrow 10 cpu cycles
- Create thread \rightarrow 200,000
- Destroy thread \rightarrow 100,000
- Switch context - > 4000
- Write file \rightarrow 10,00,000
-



When performing asynchronous I/O-bound operations, you have a device driver doing the work for you and no threads are required.



All synchronous I/O operations suspend the calling thread while the underlying hardware (disk-drive, network interface card, etc.) performs the work.

Sync IO

suspend the calling
thread while hardware
performs the work.

Non-blocking Async IO

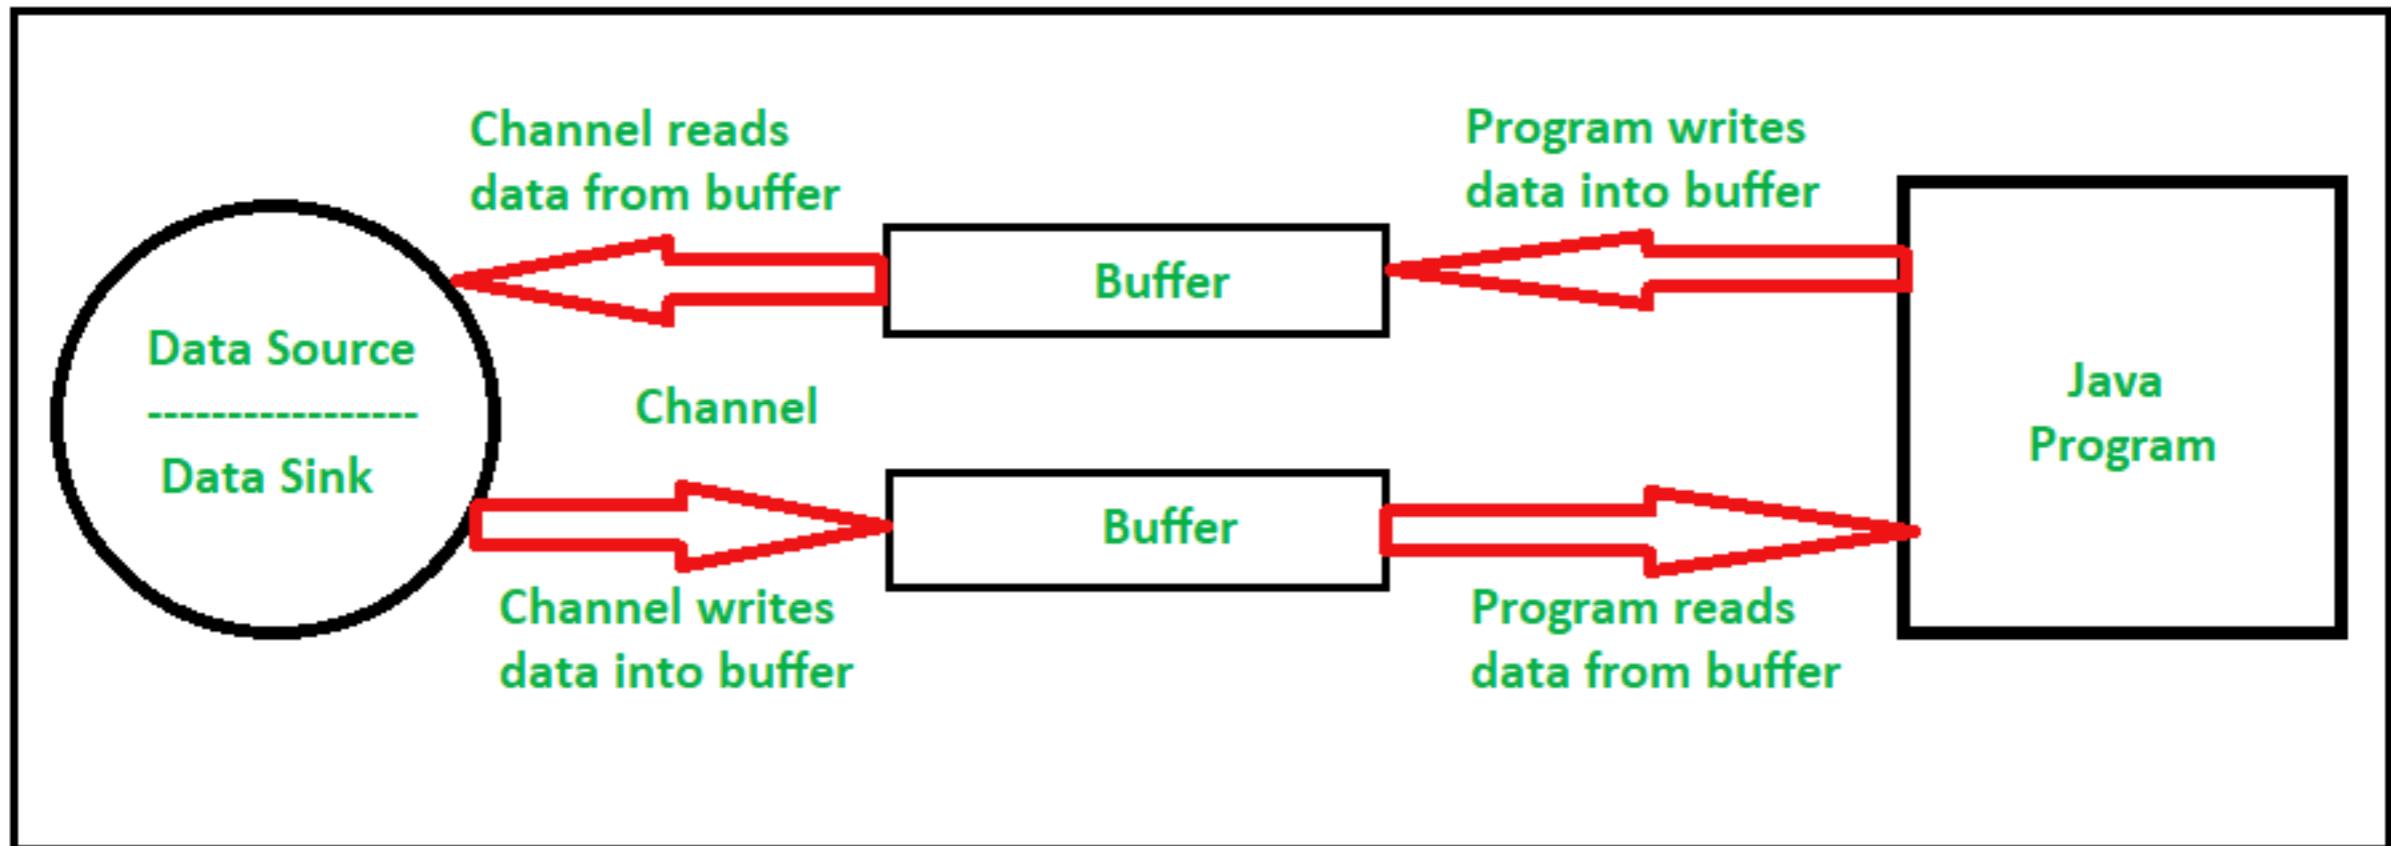
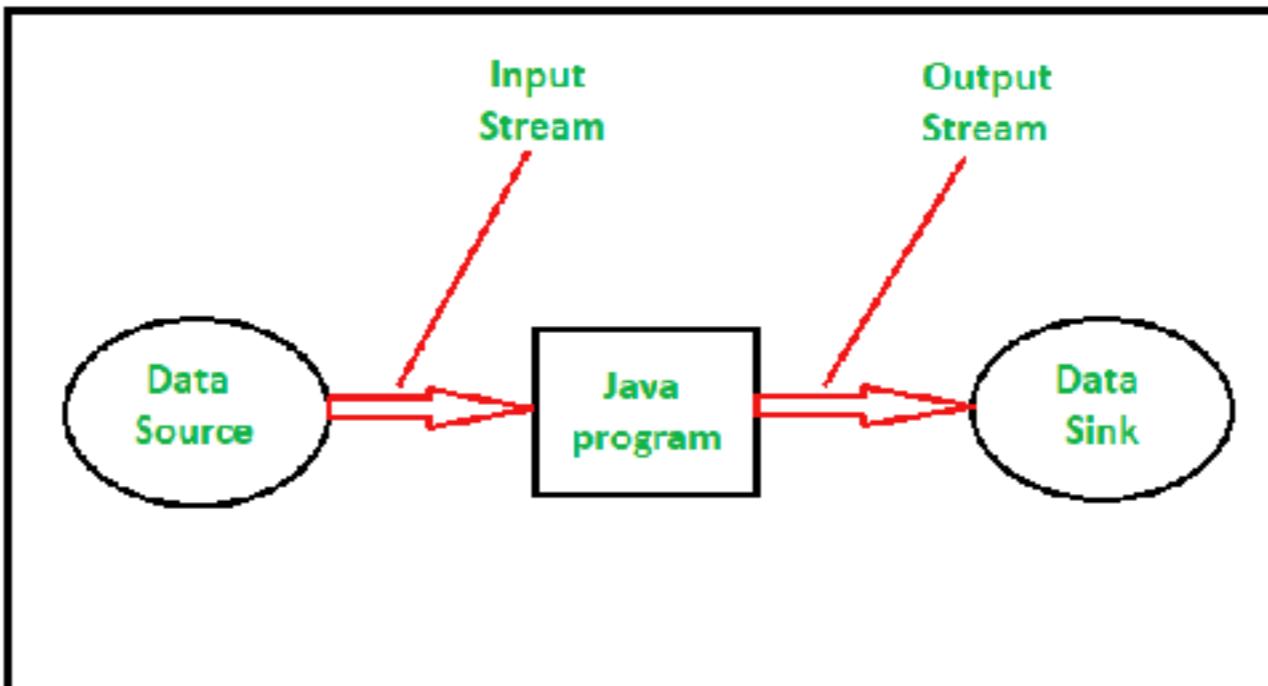
The kernel doesn't block your thread and returns from the IO call immediately. Your code then needs to check occasionally if the kernel has already done its job. As an example, Linux provides the `epoll` interface for the non-blocking IO.

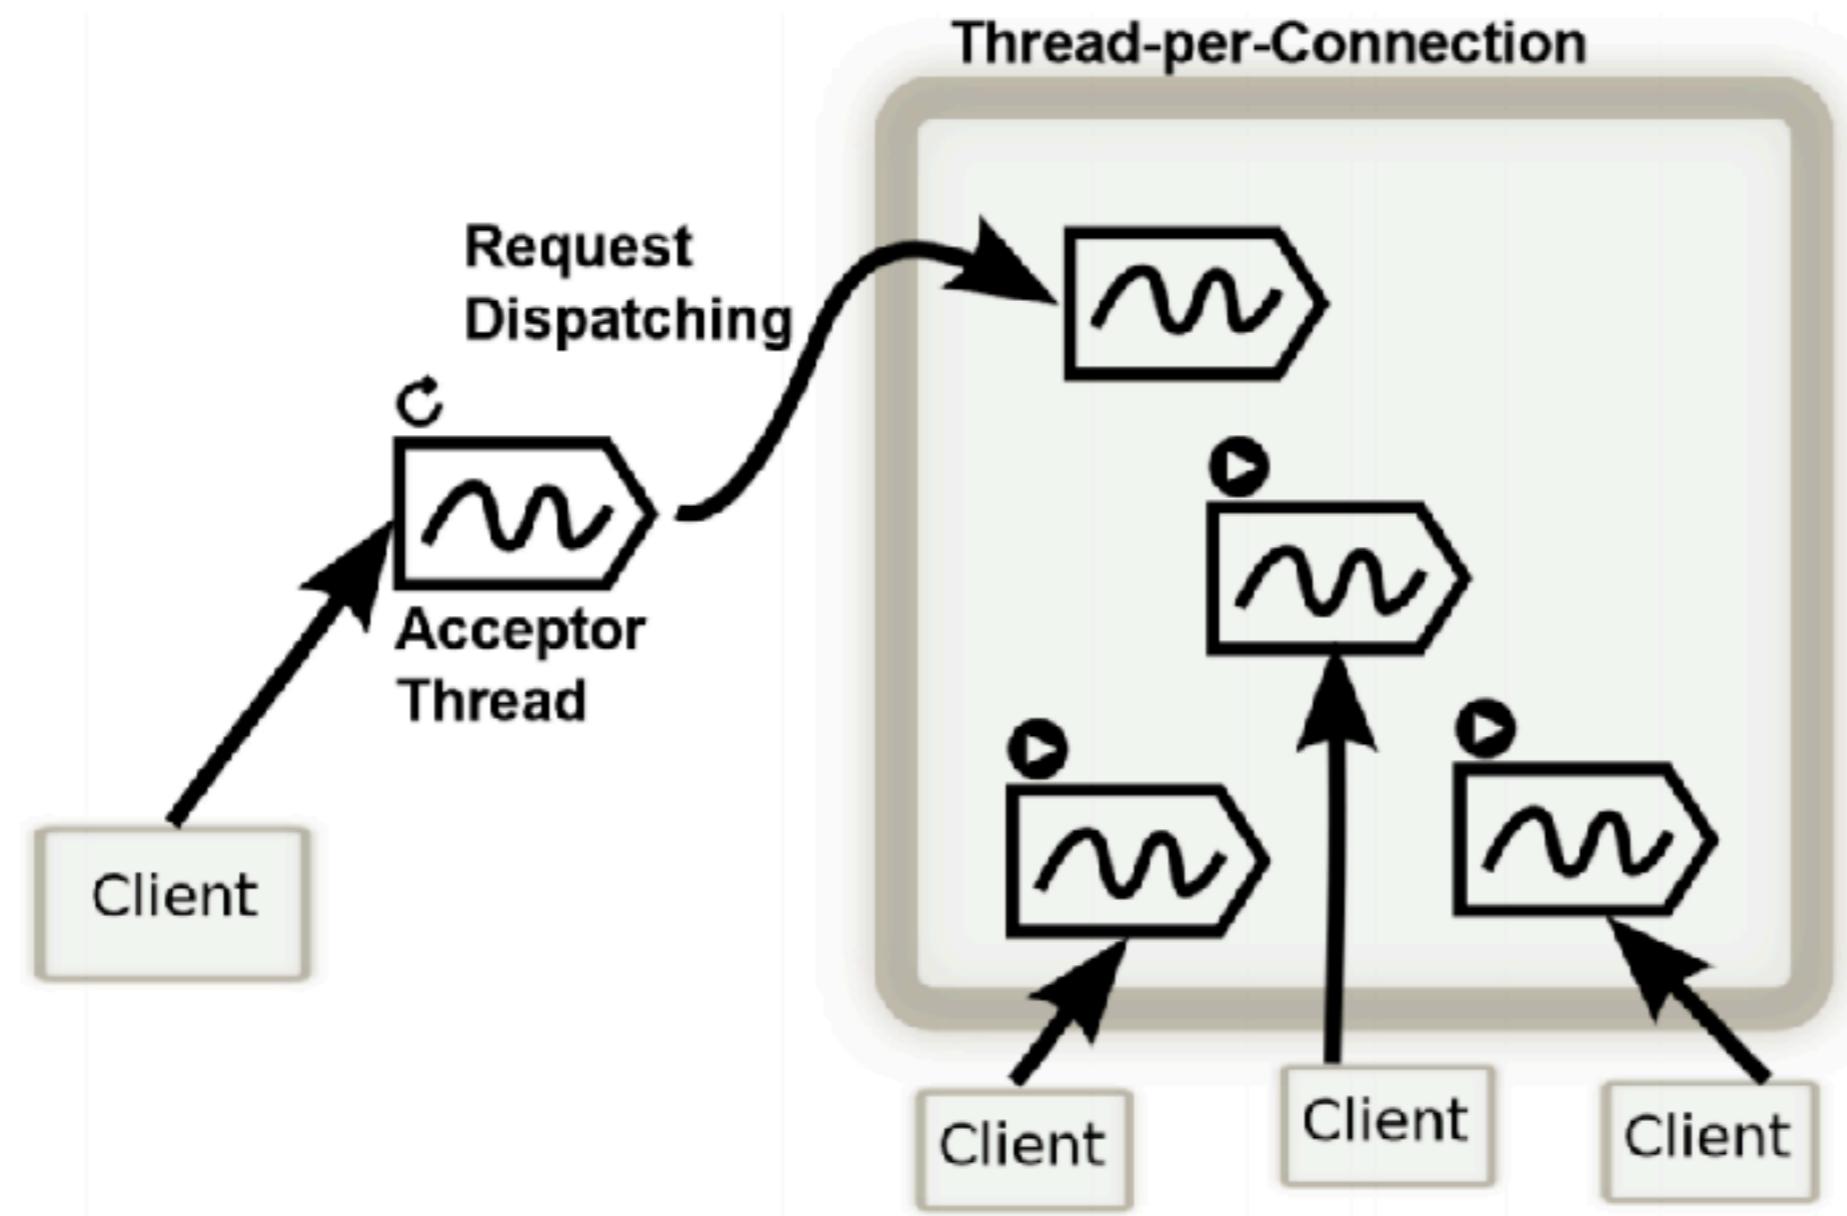
Blocking Async IO

There is still a thread lower down the stack which is blocked on the IO, but your thread isn't.
called wrapped inside a thread which will block until the IO is done and then delegate the handling of the result of the IO to your callback function.

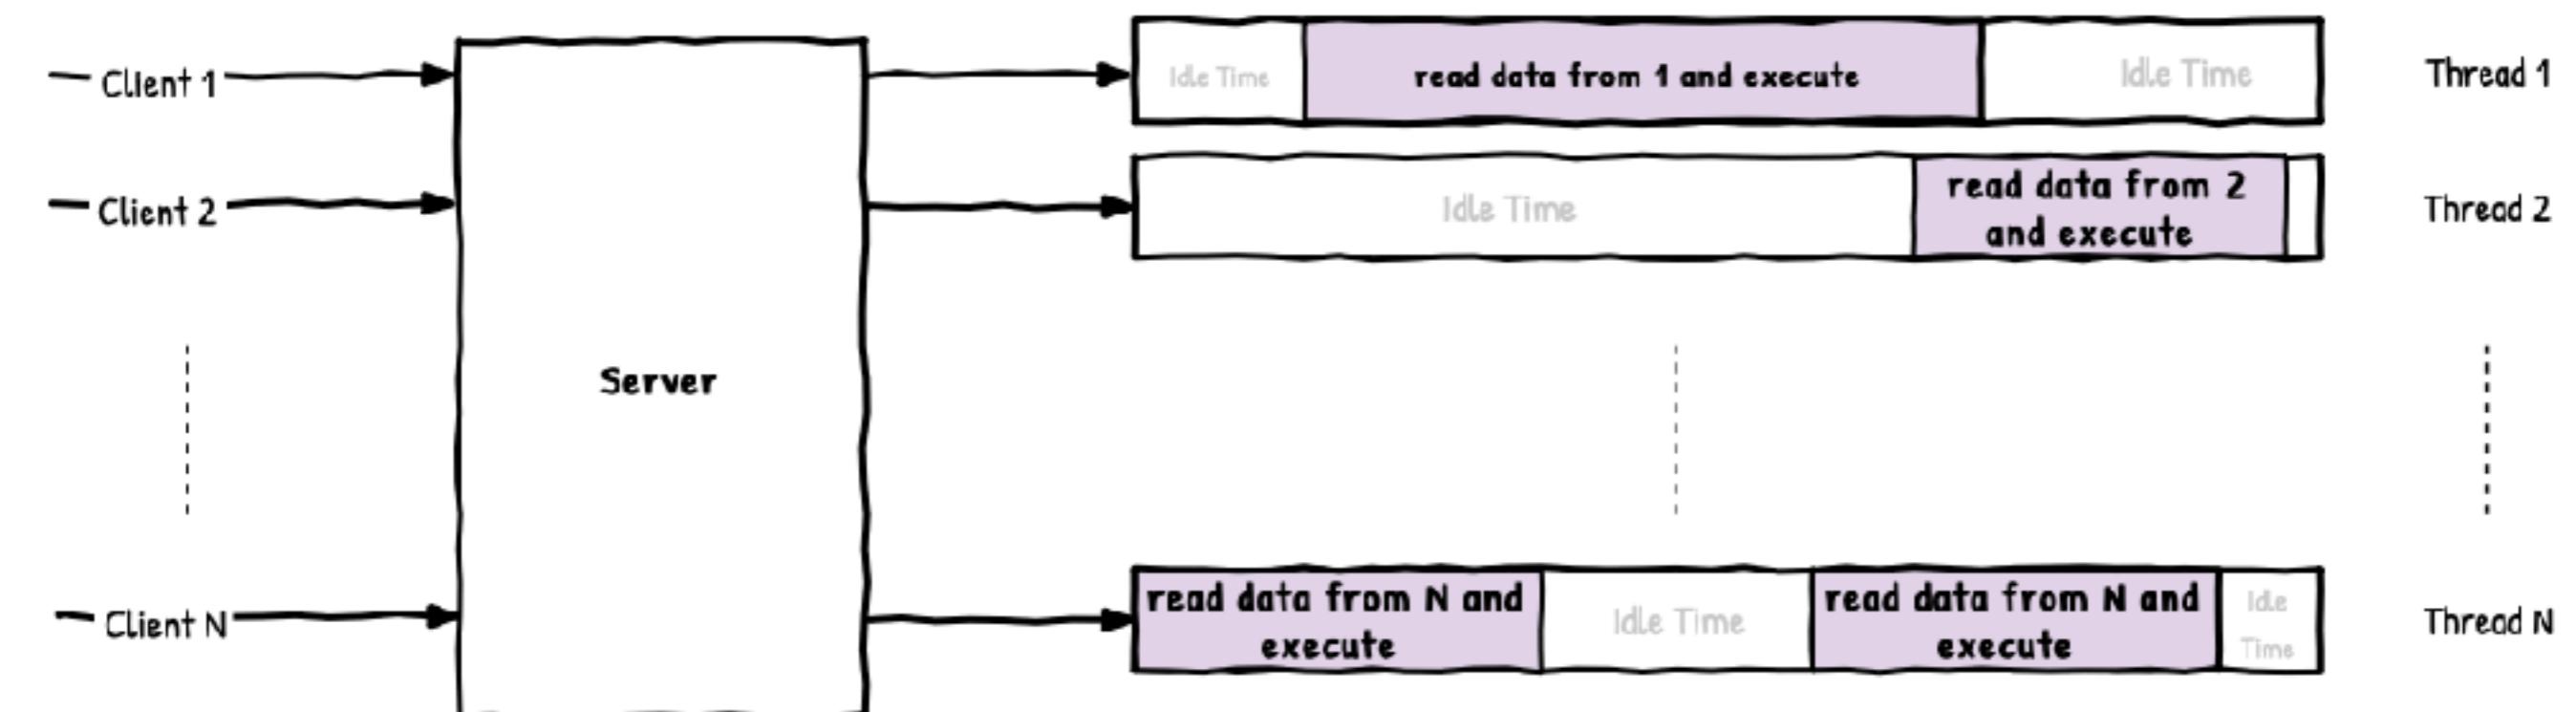
hardware (disk-drive, network interface card, etc.)

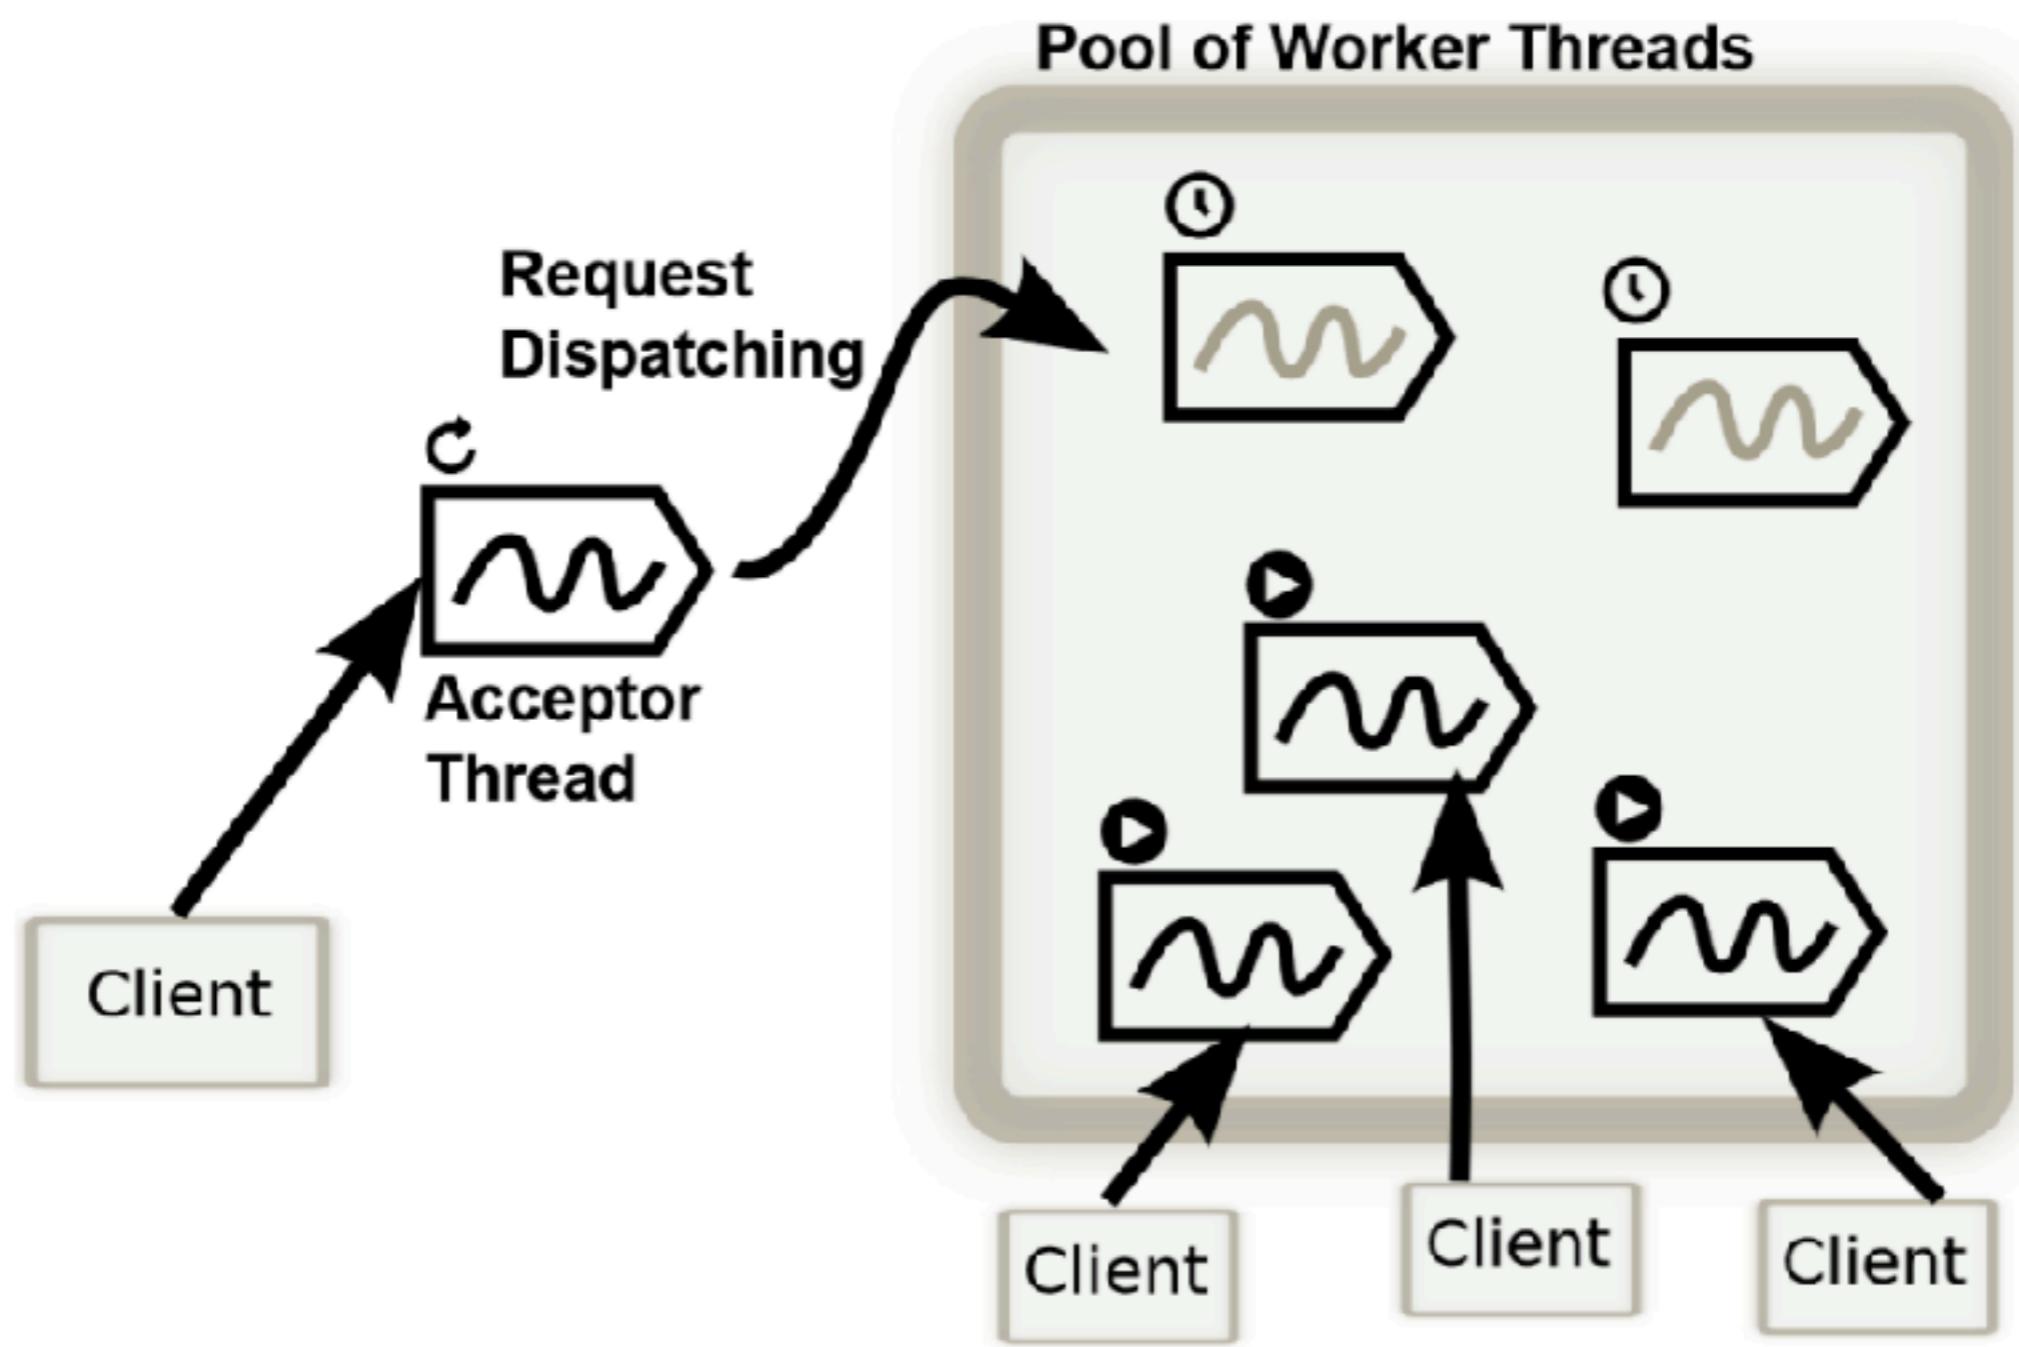
Blocking vs. Non-blocking IO

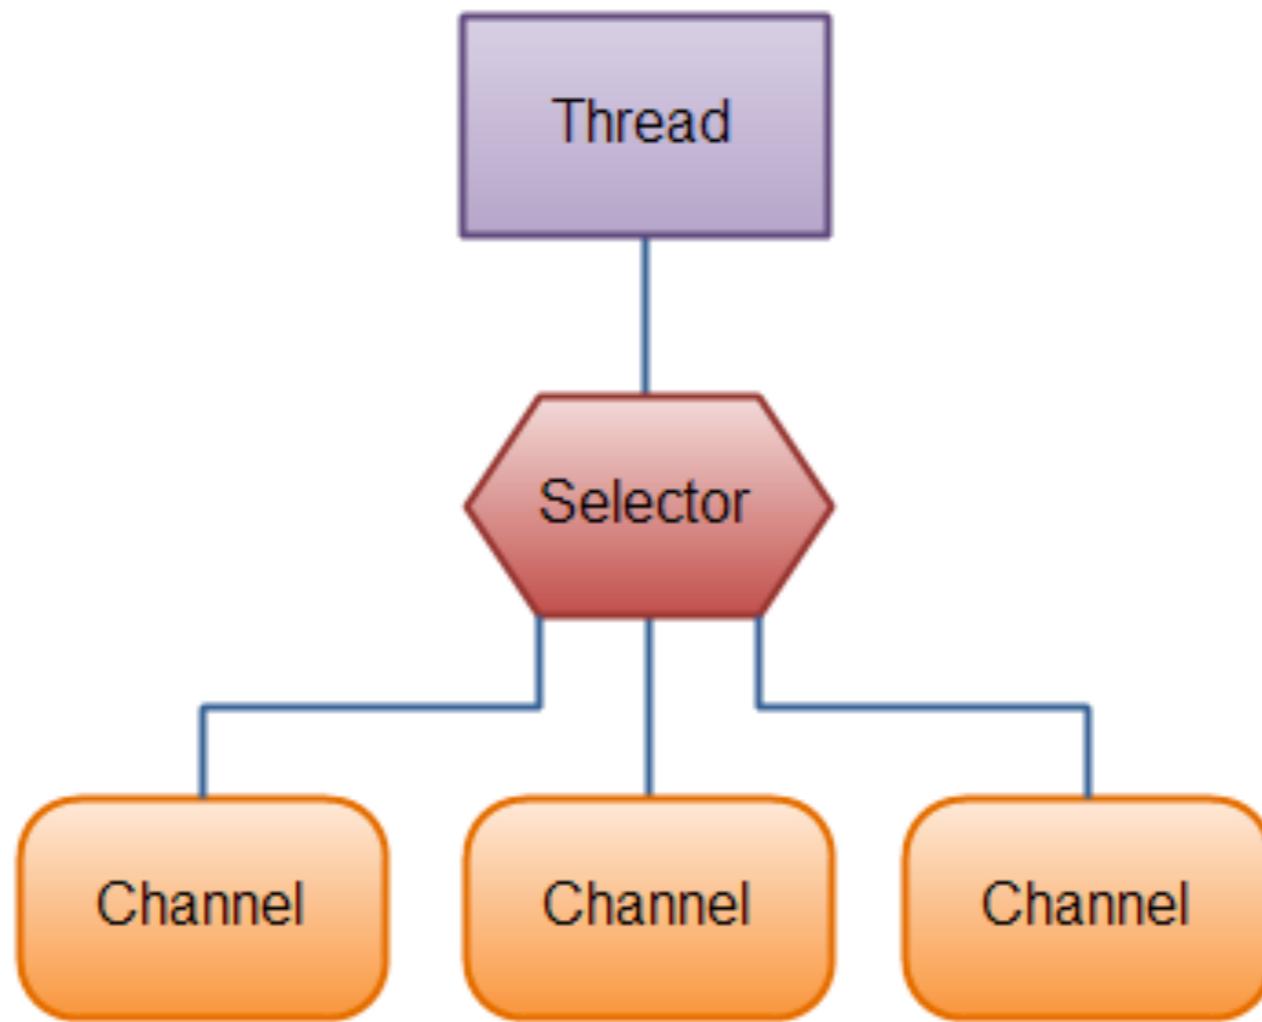




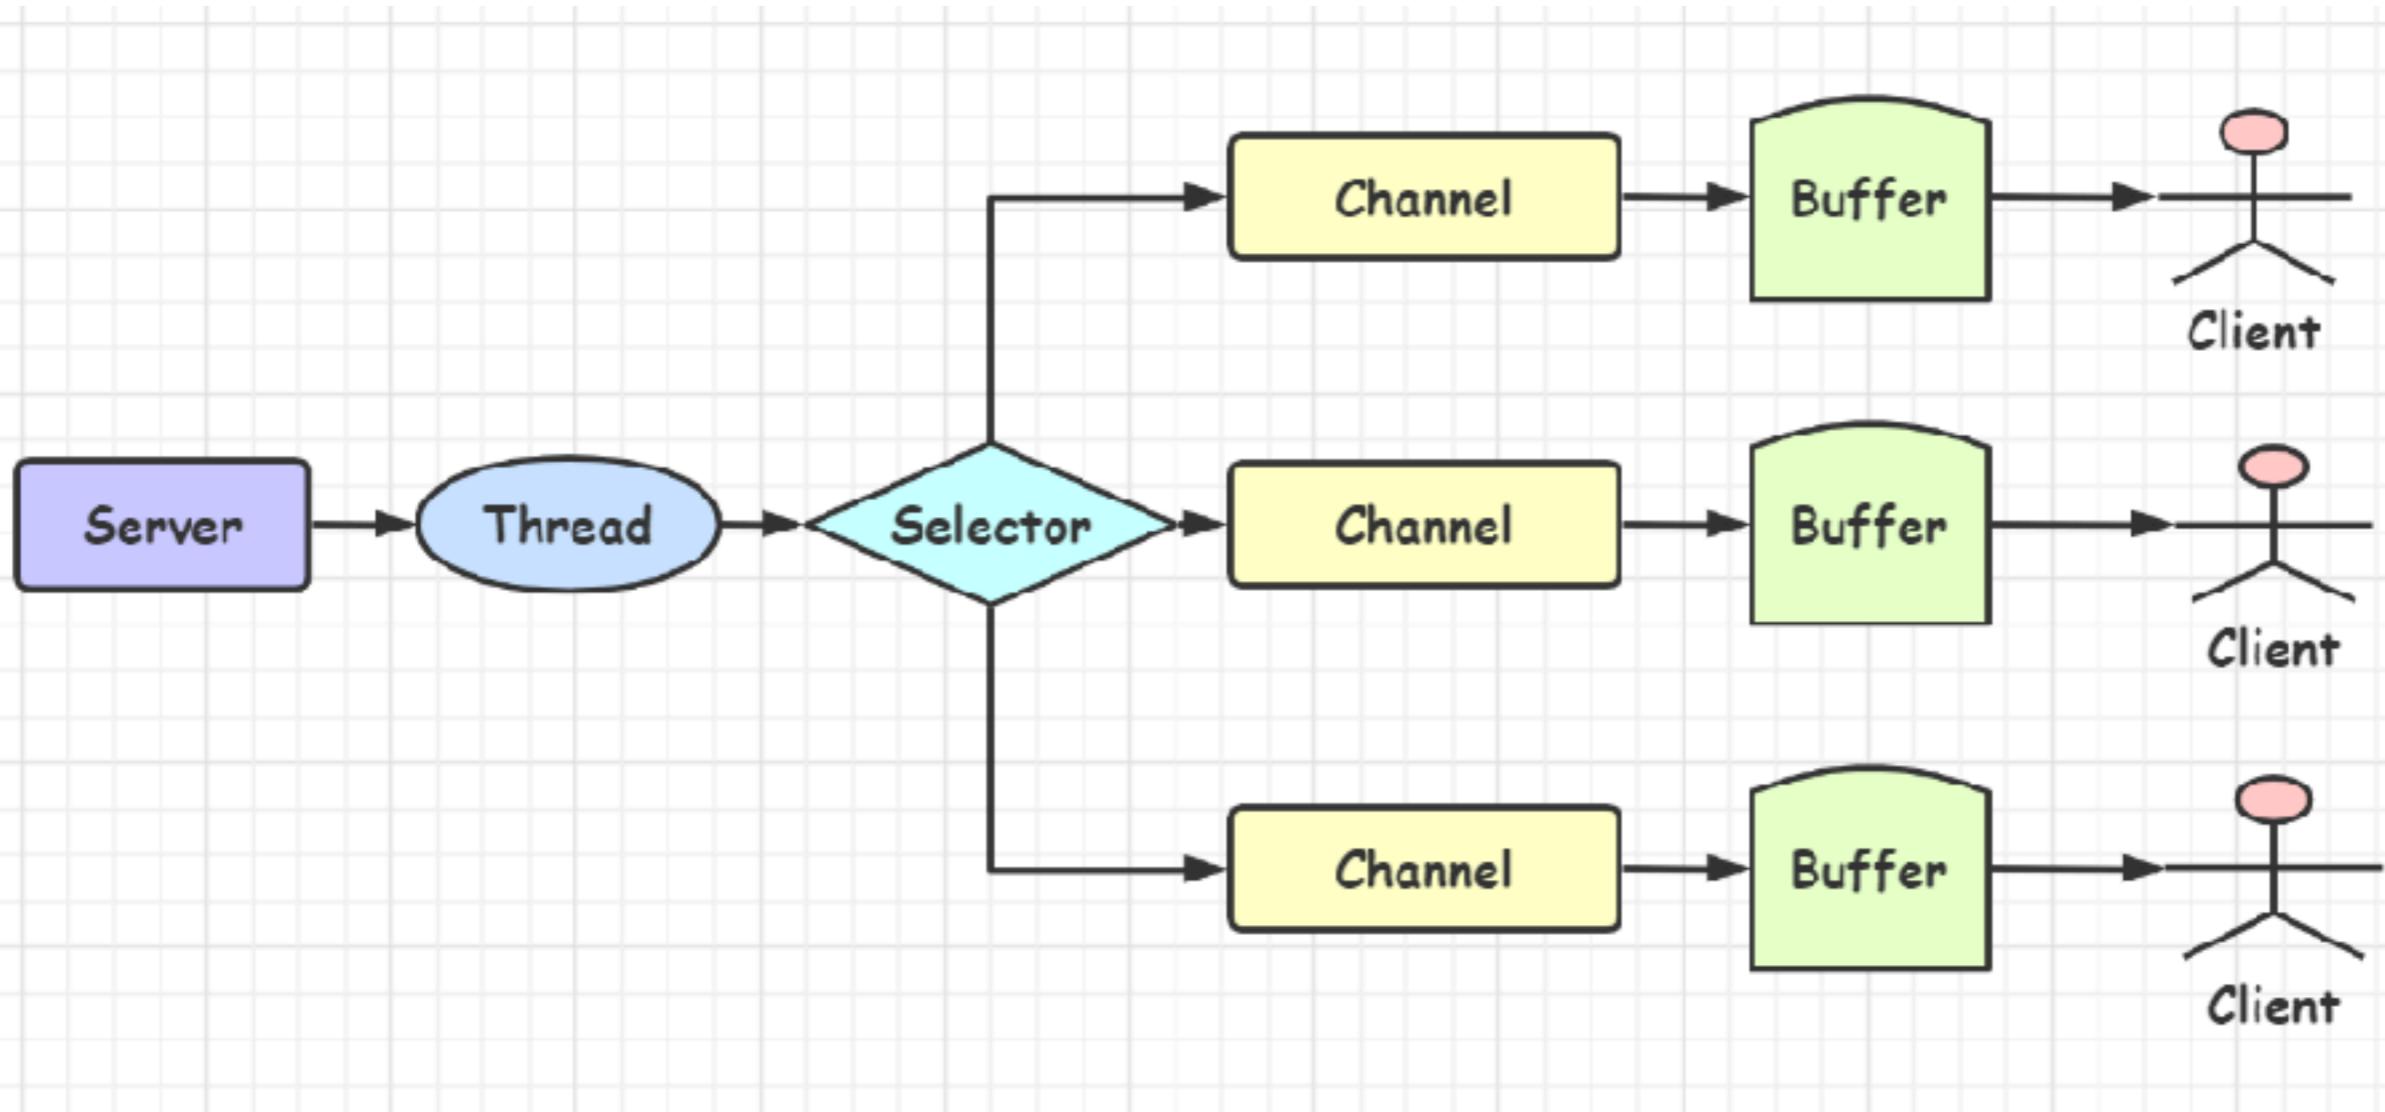
Server with a thread per connection (Blocking sockets)

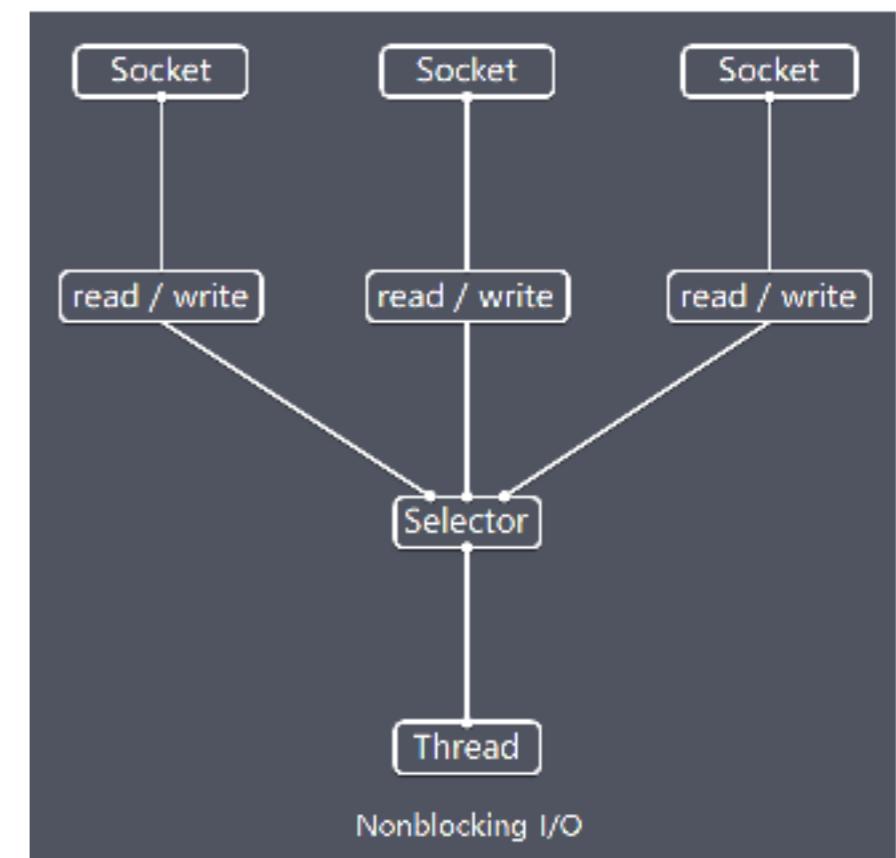
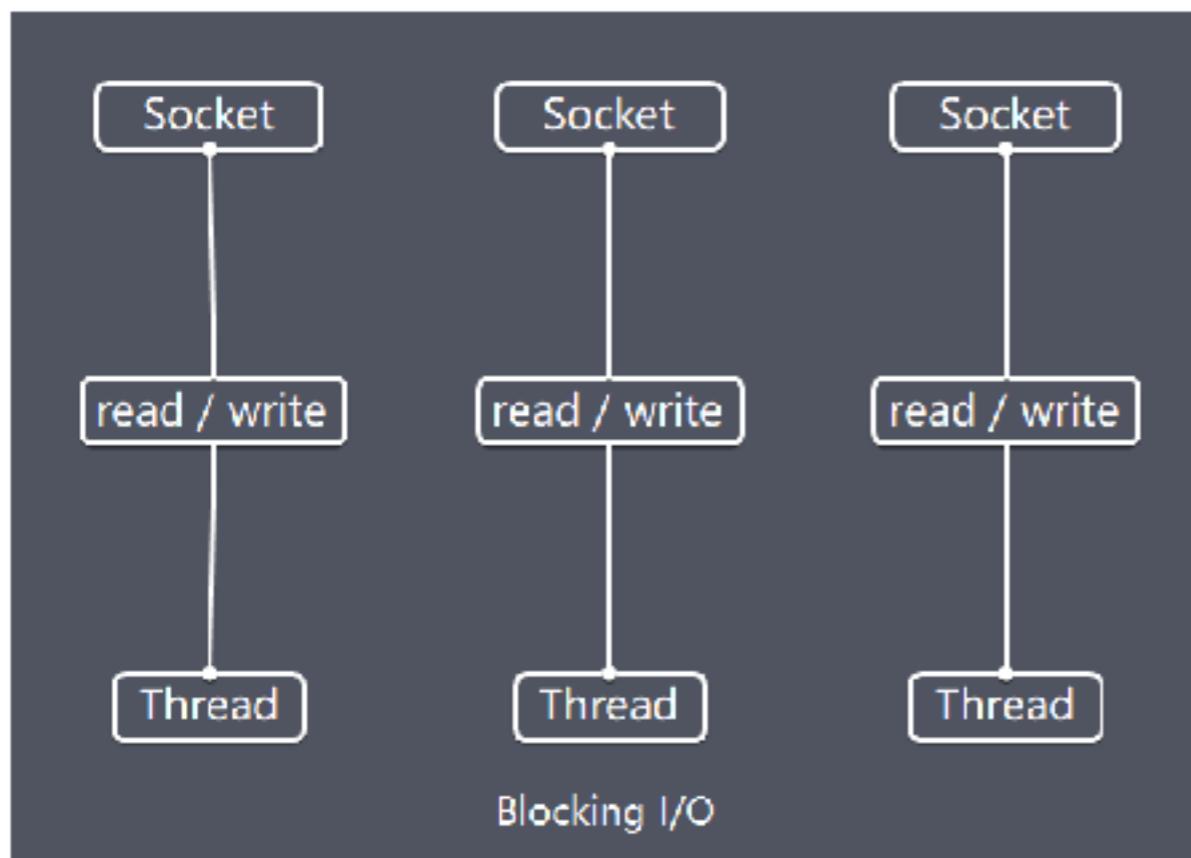






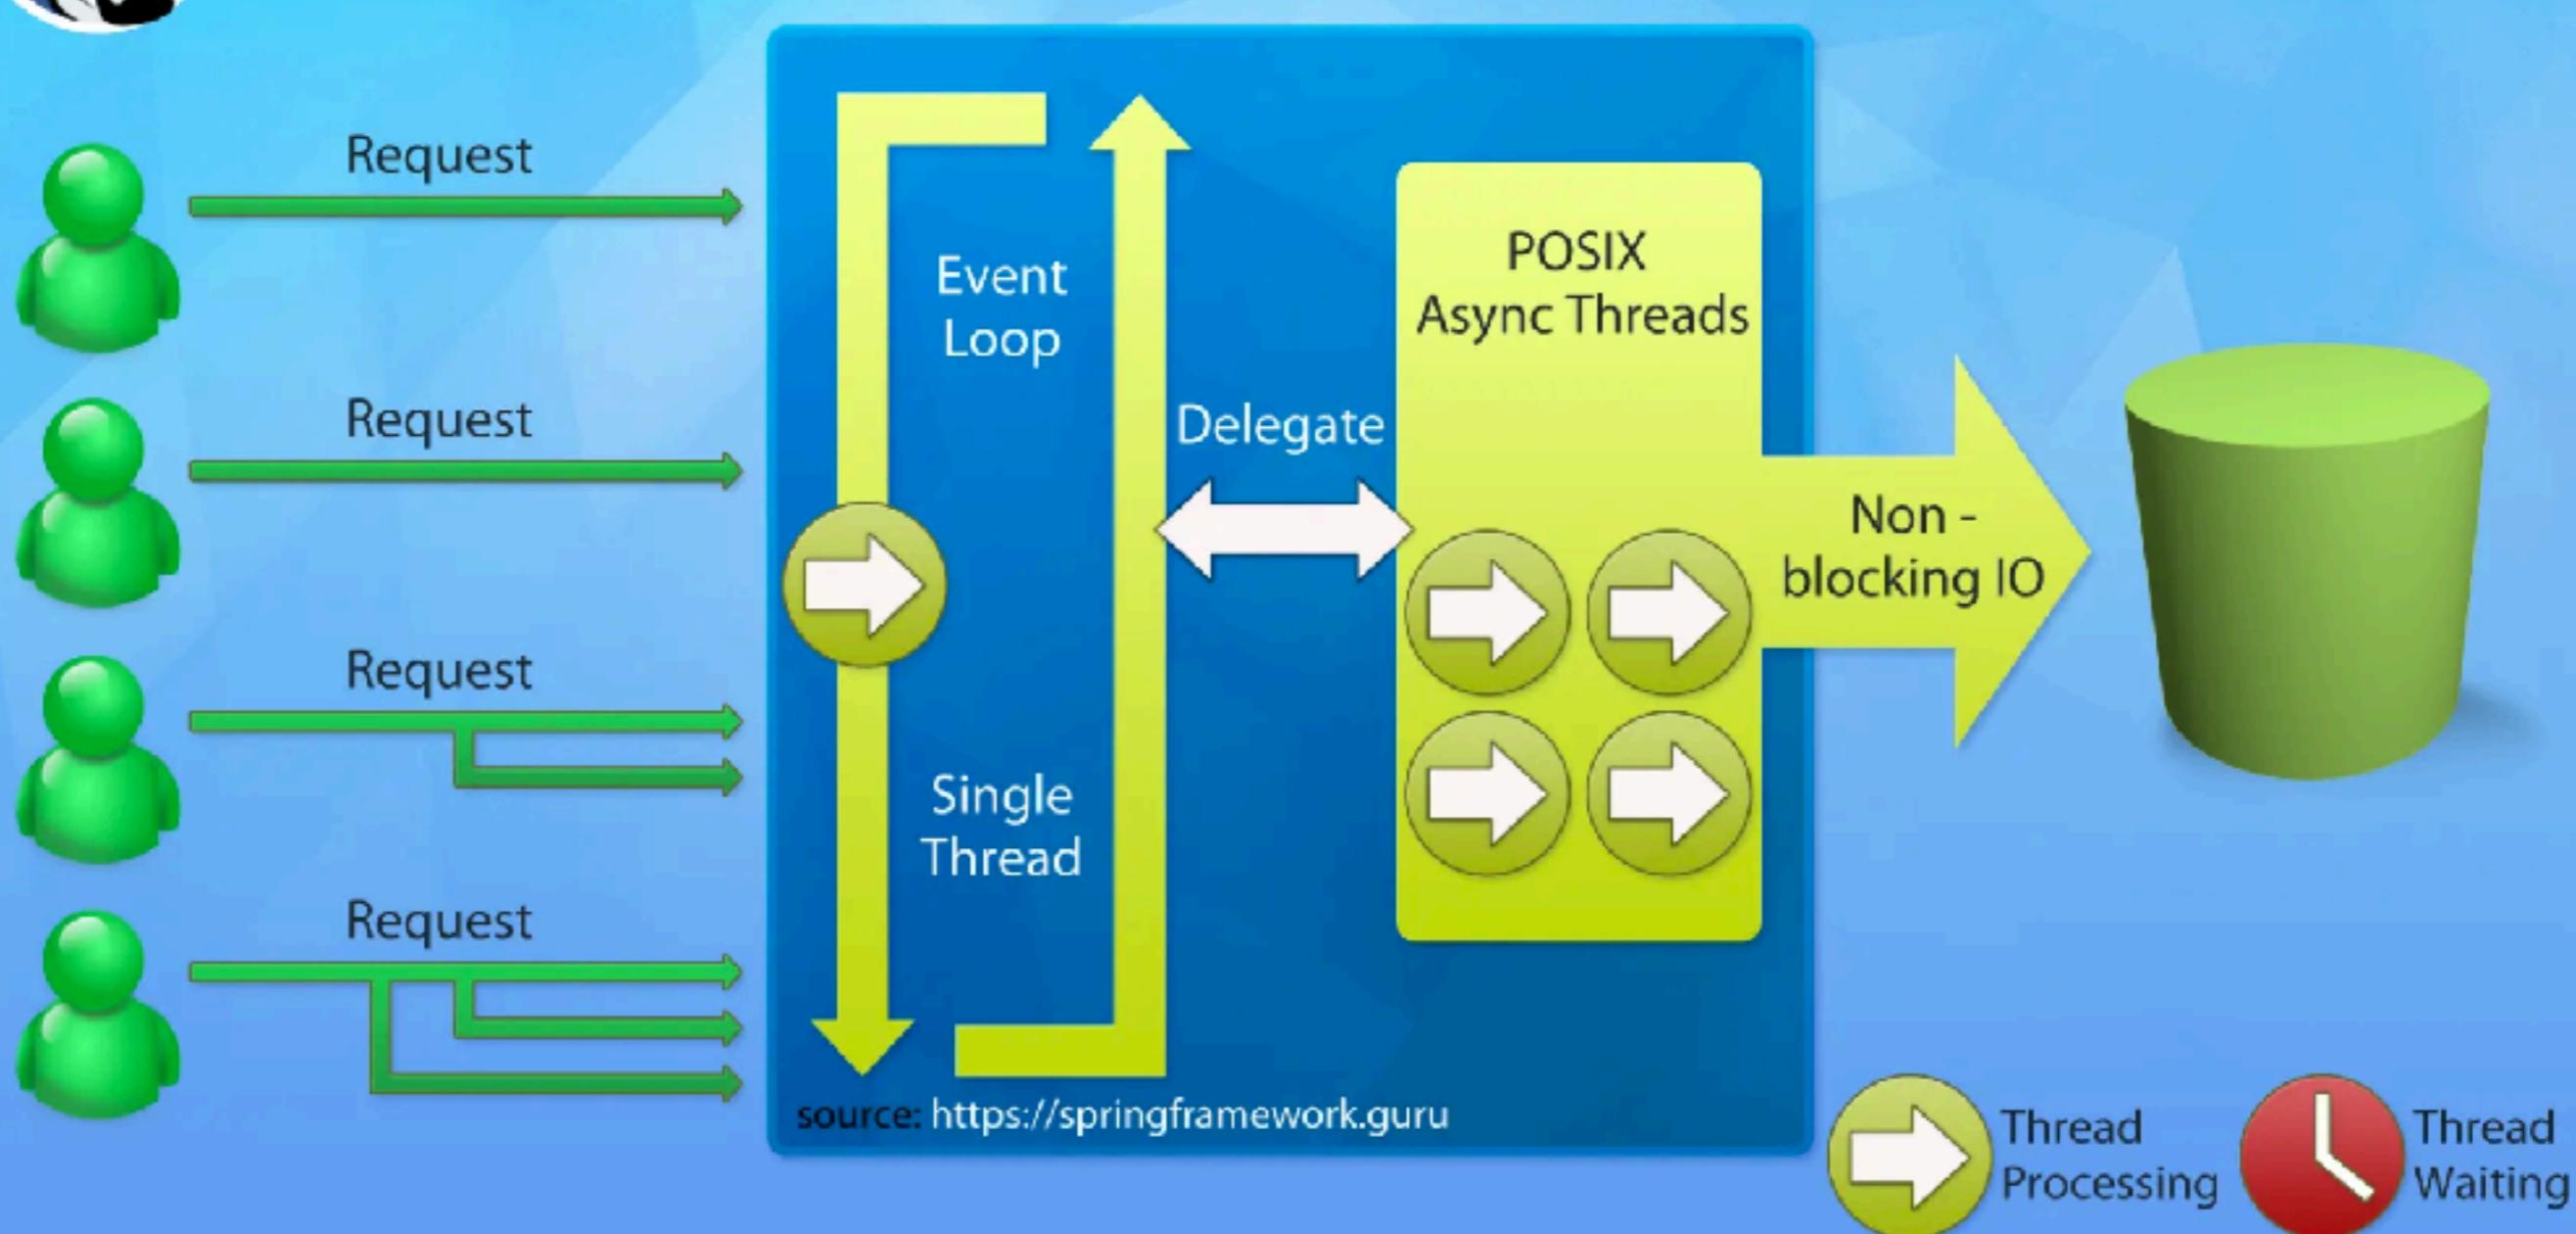
one thread to handle all of your channels. Switching between threads is expensive for an operating system, and each thread takes up some resources (memory) in the operating system too. Therefore, the less threads you use, the better.



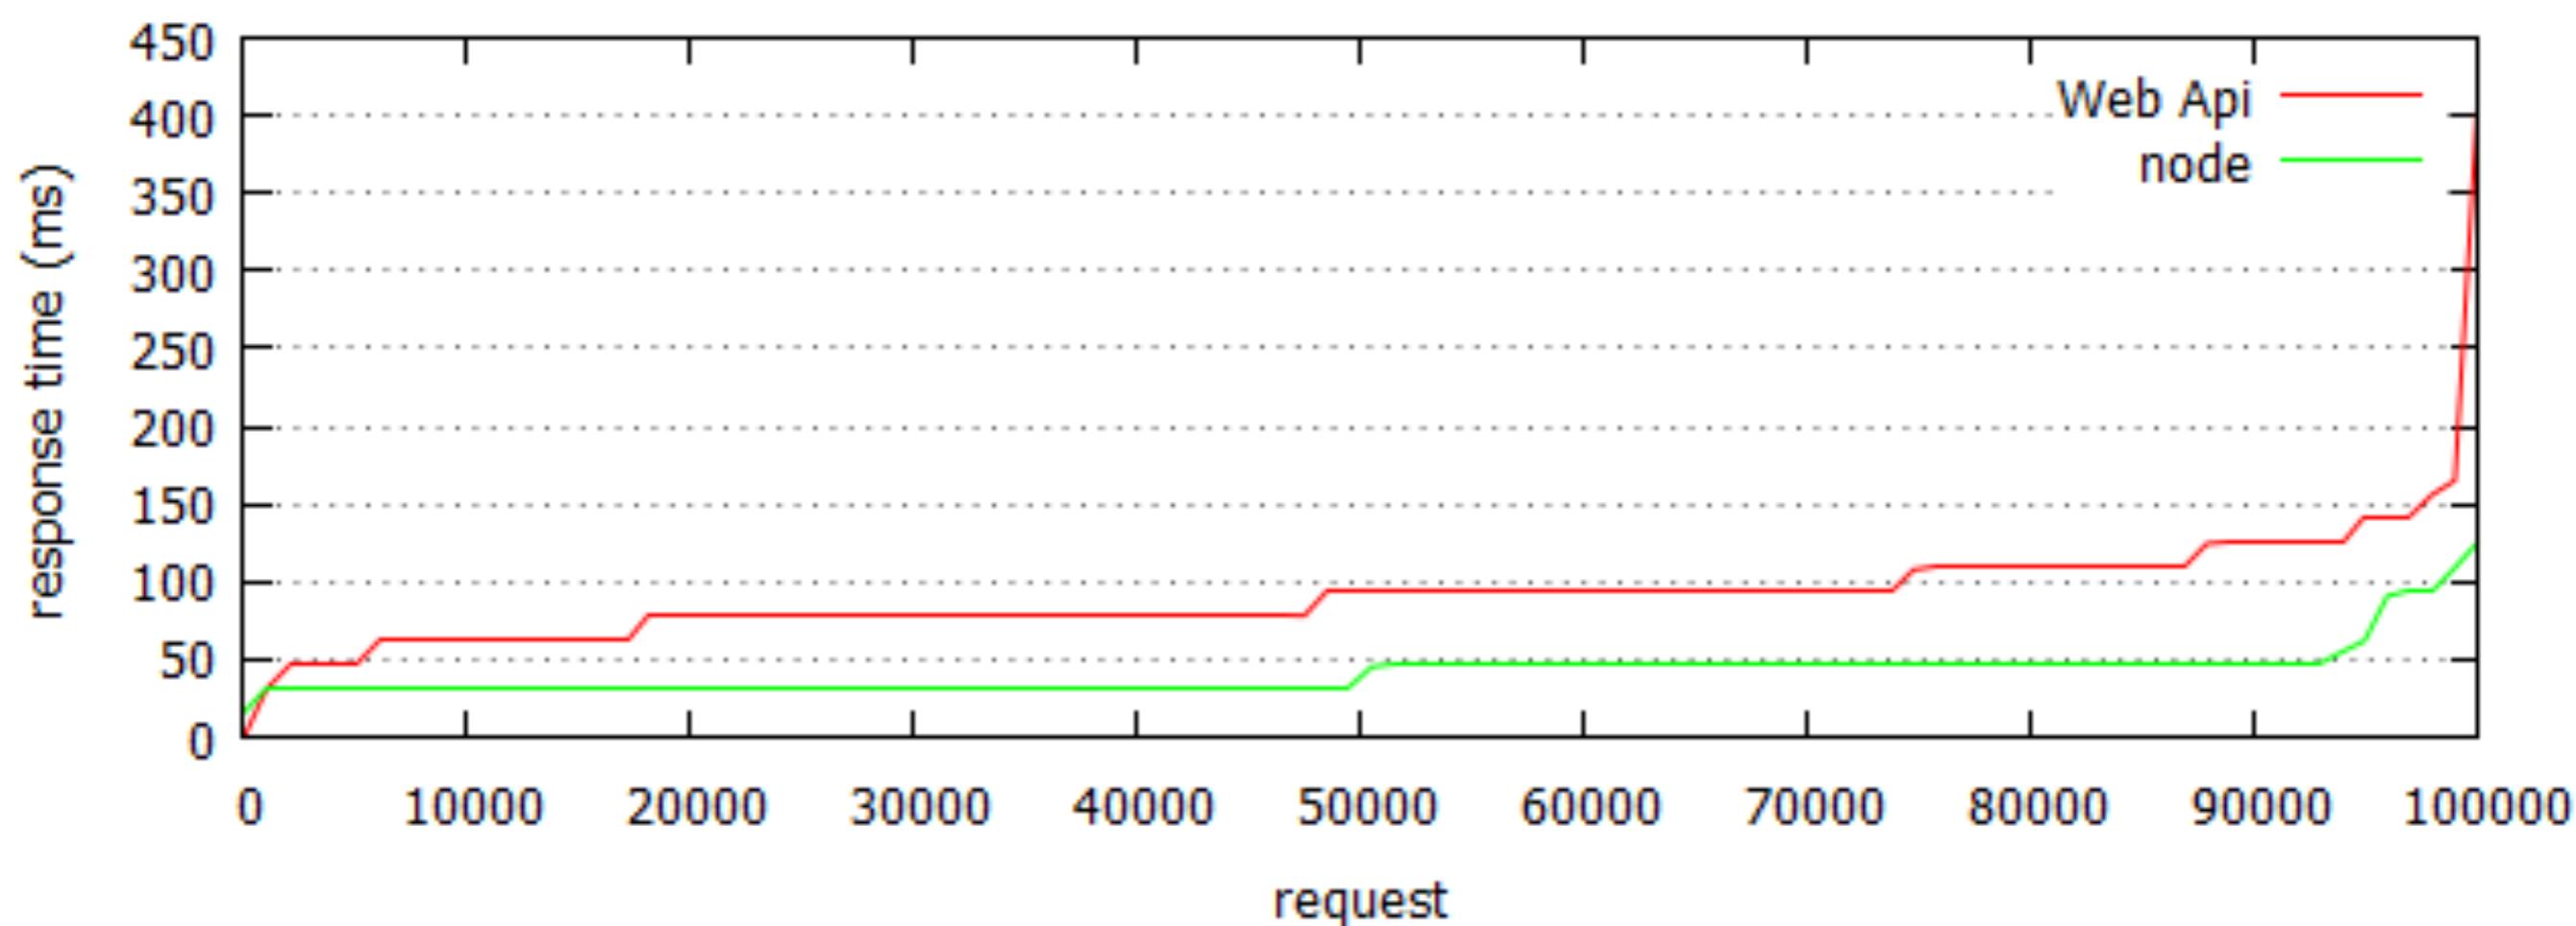


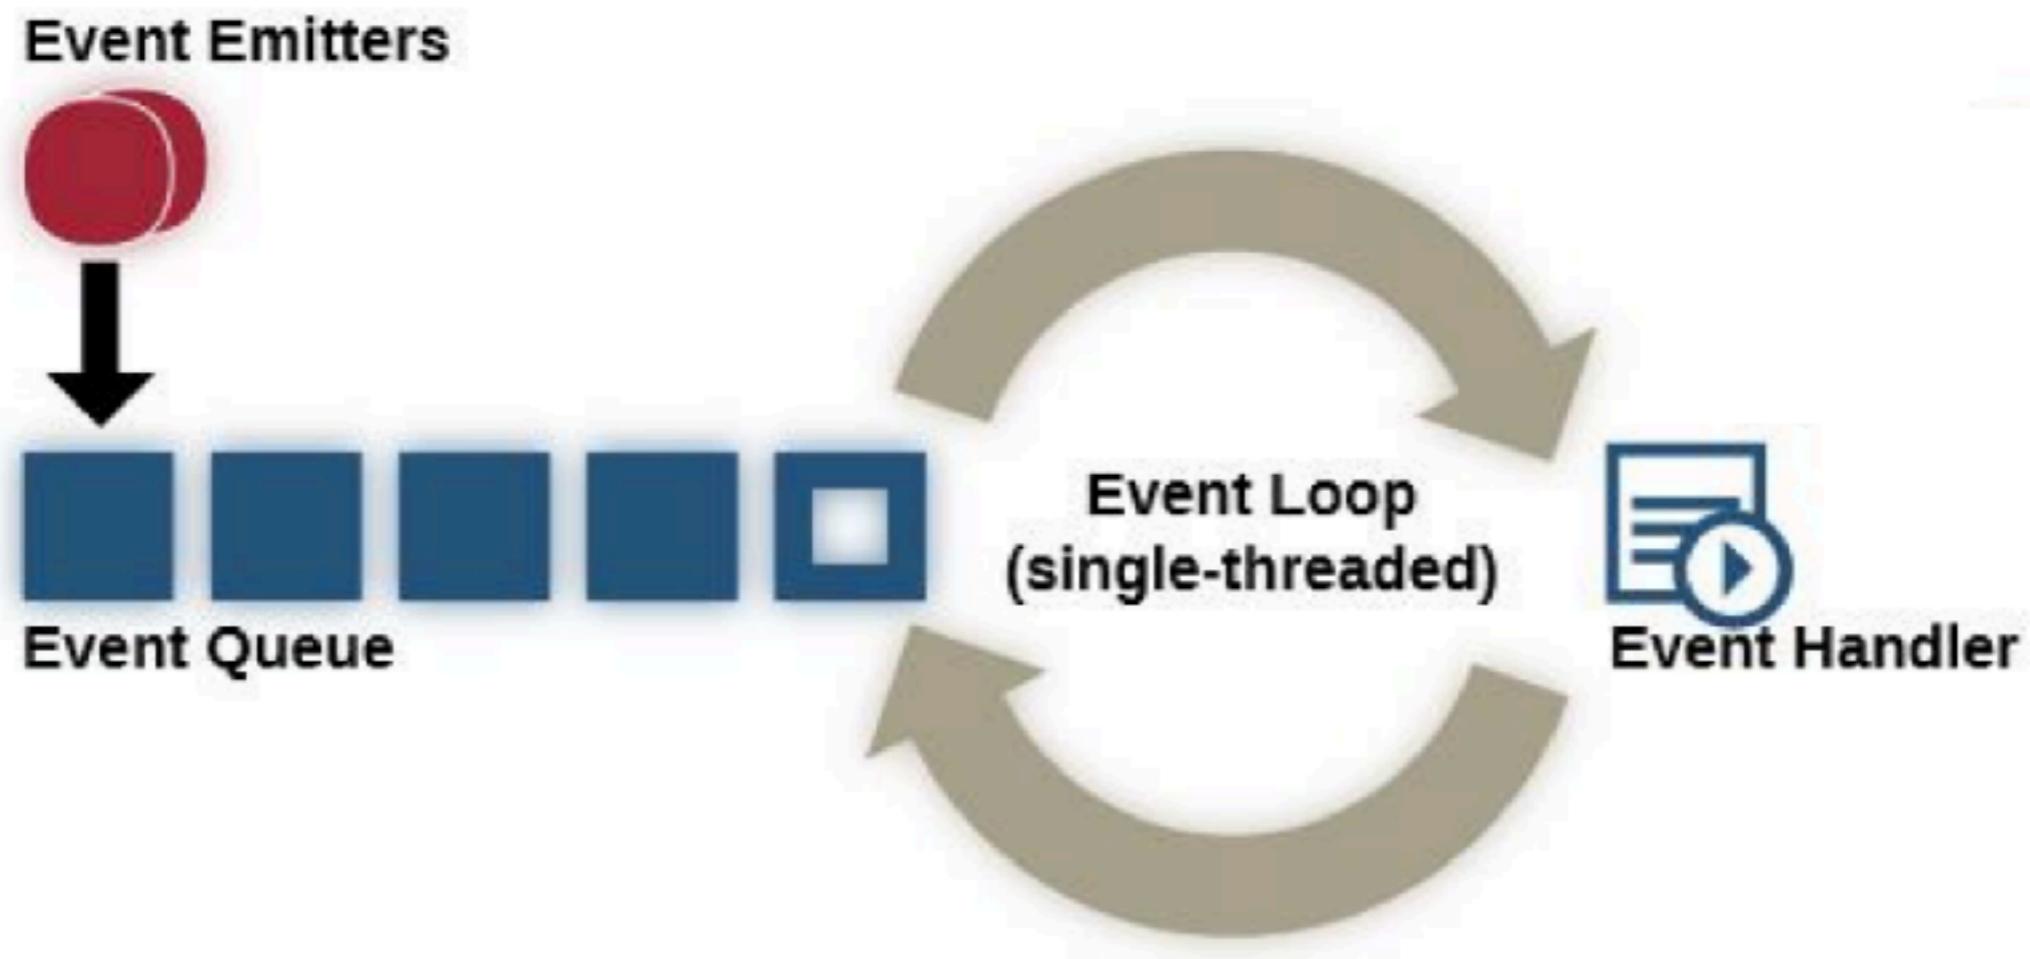


Node.js Server

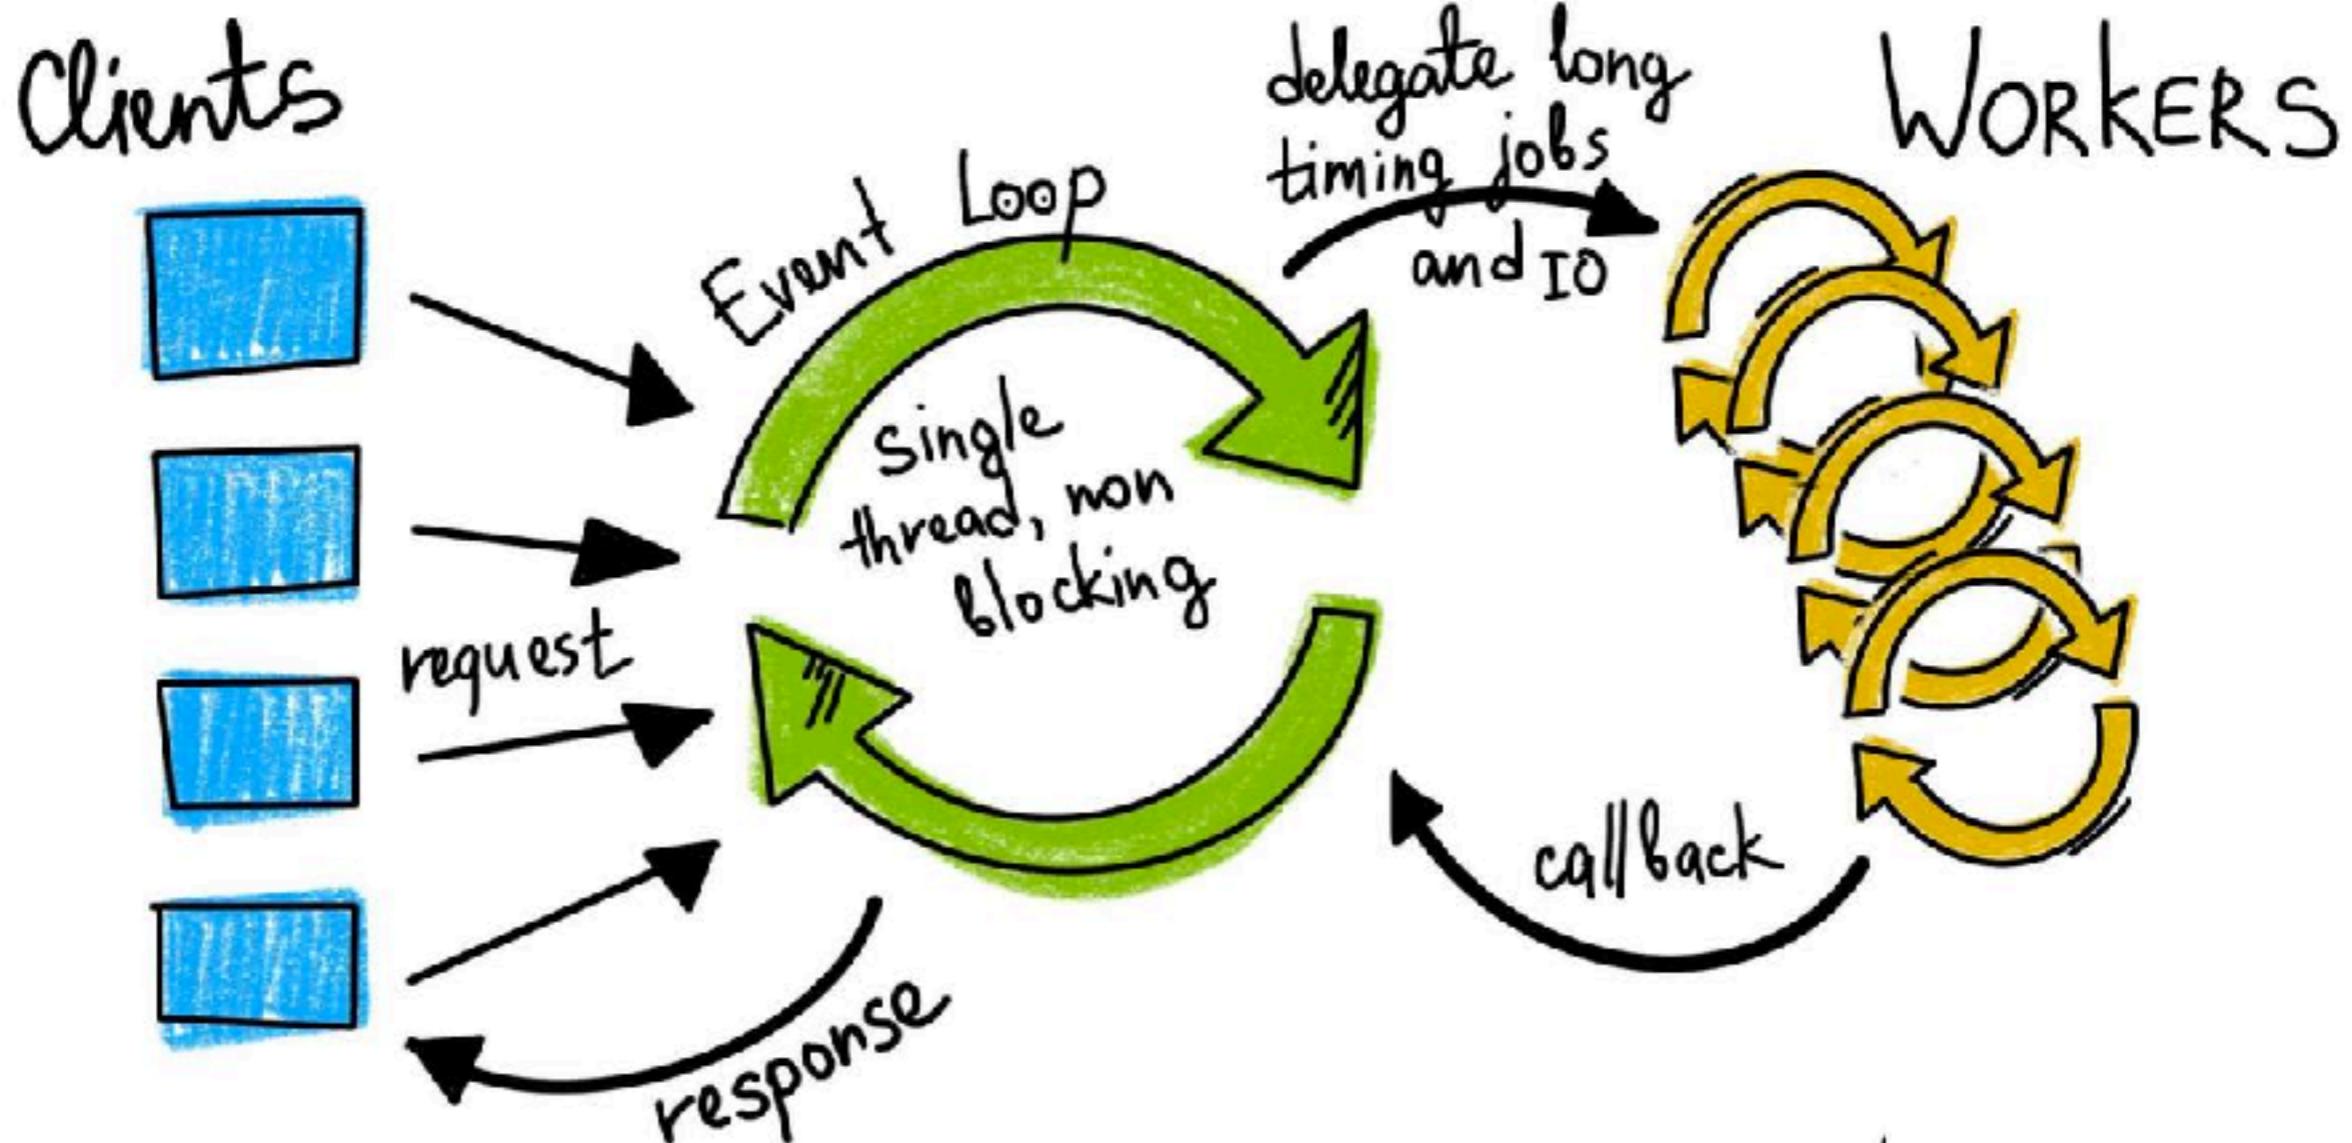


Web Api vs Node.js

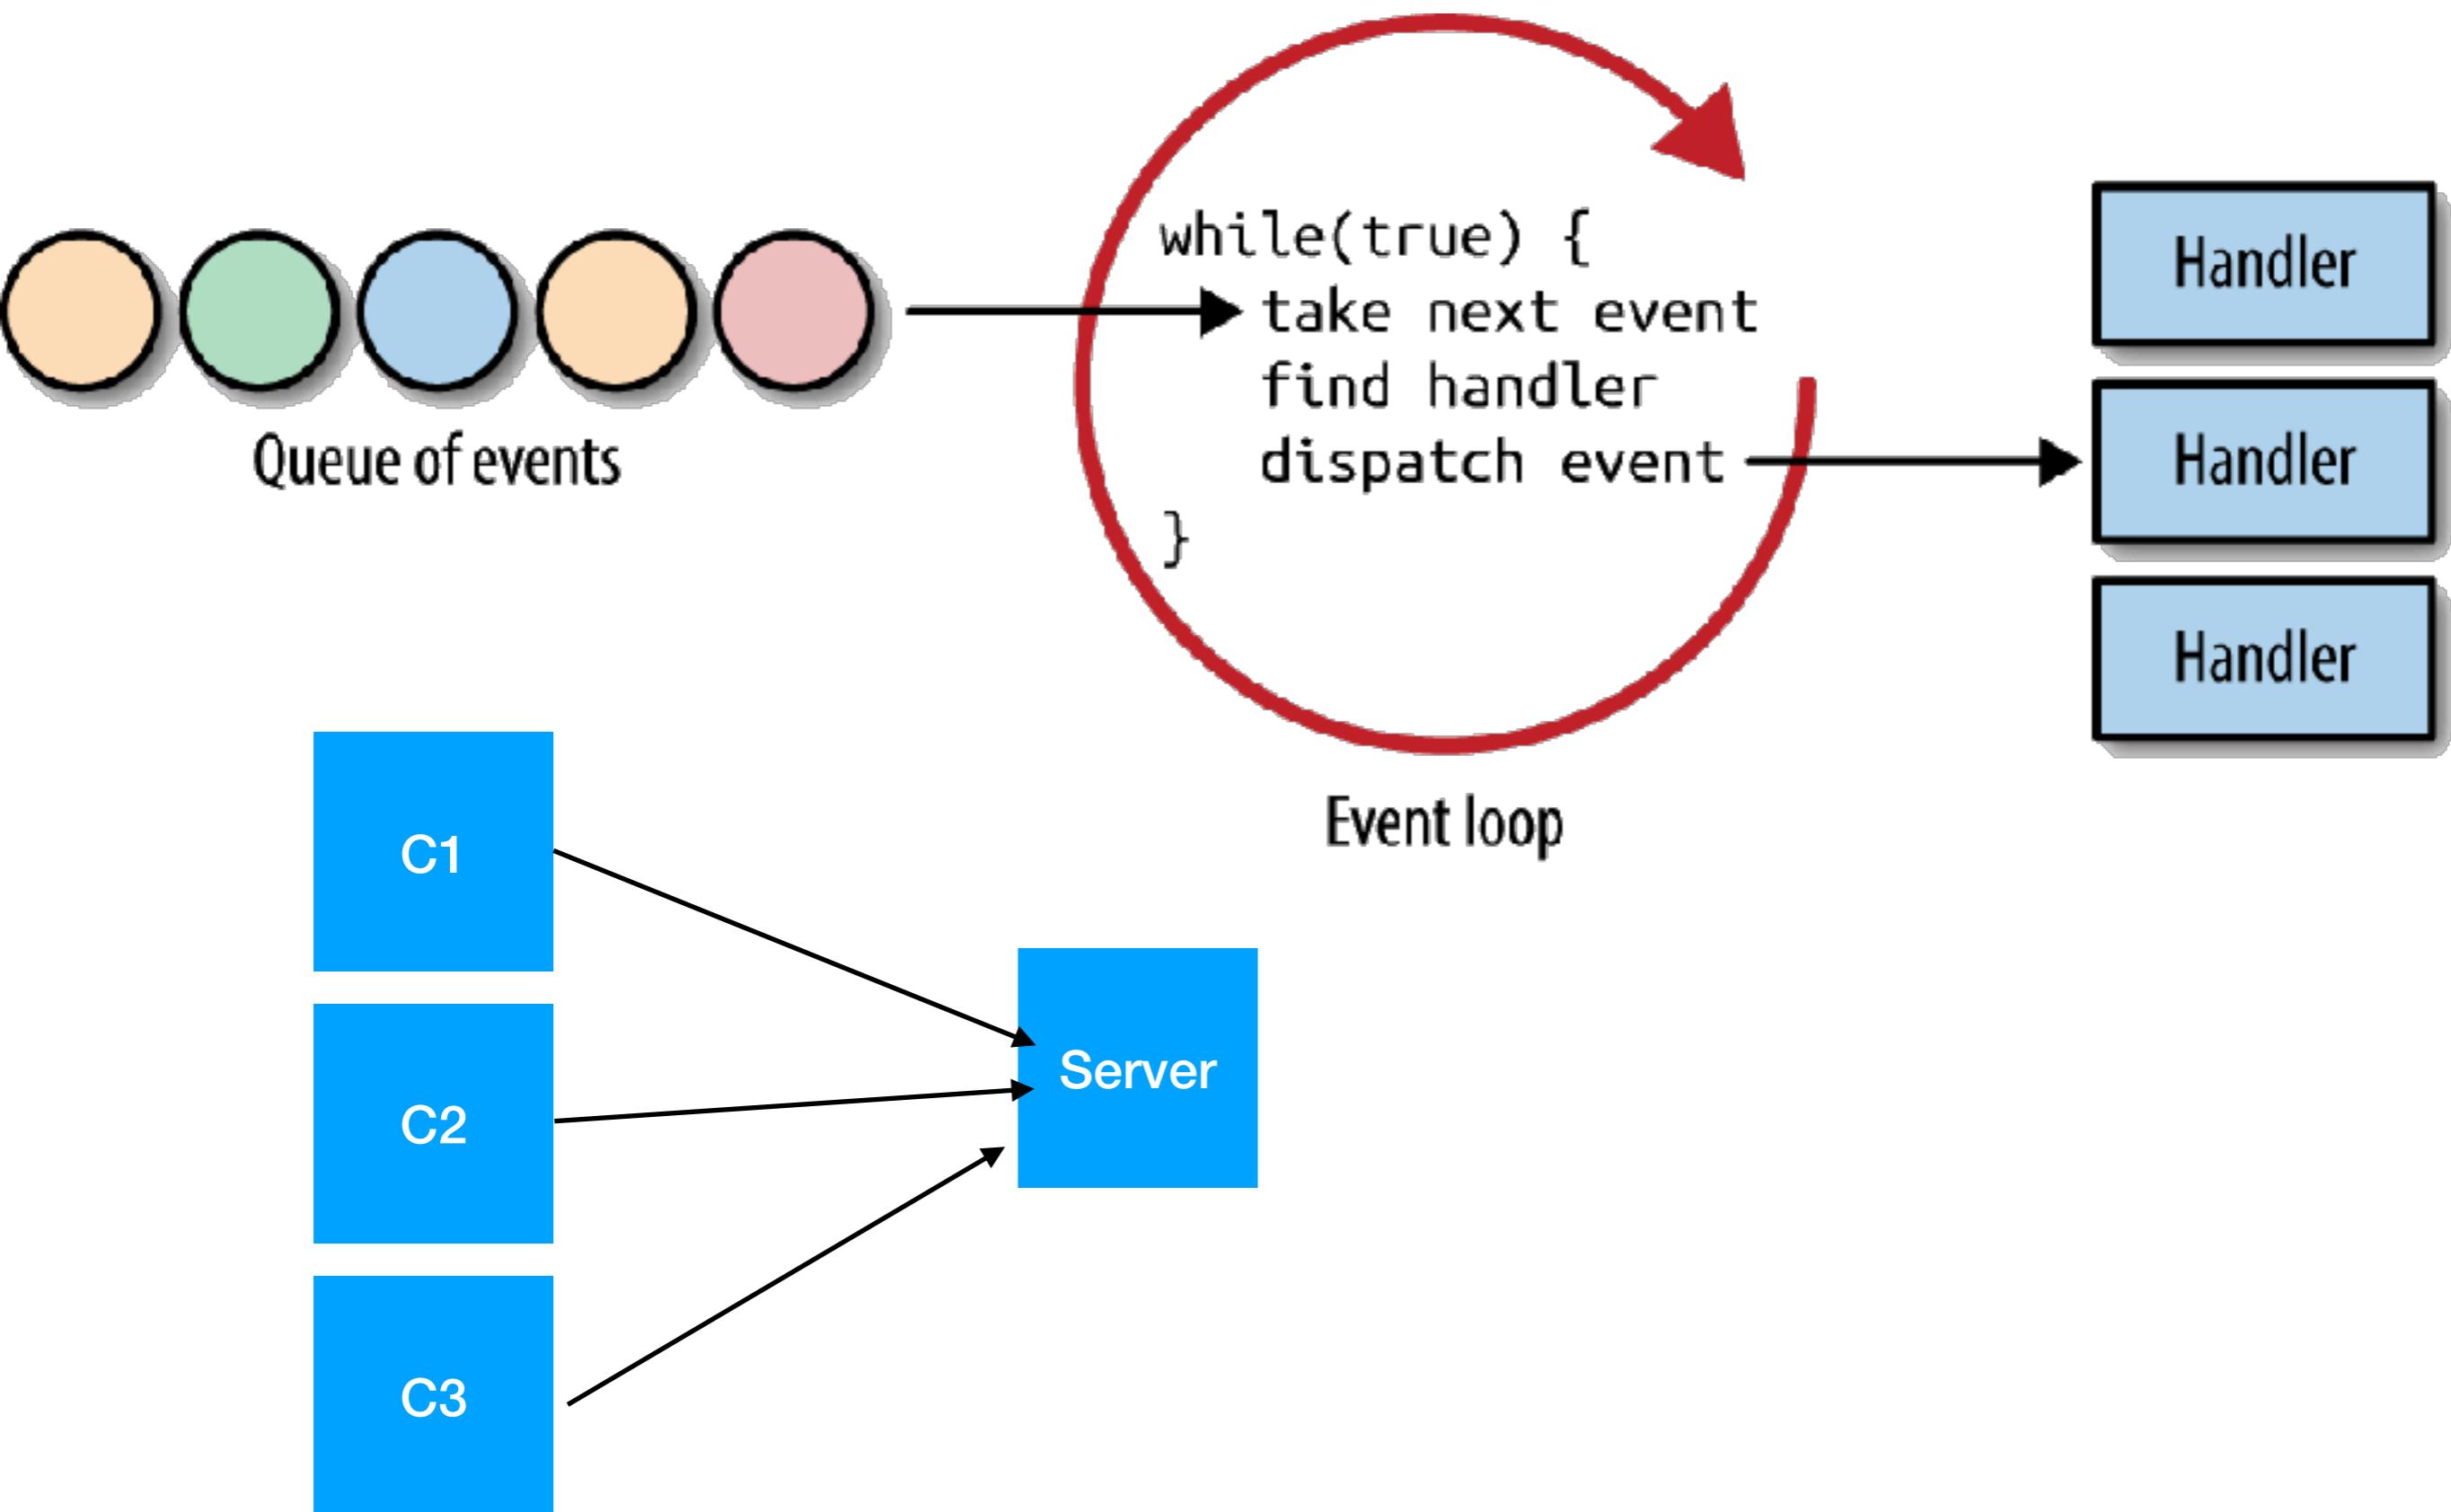


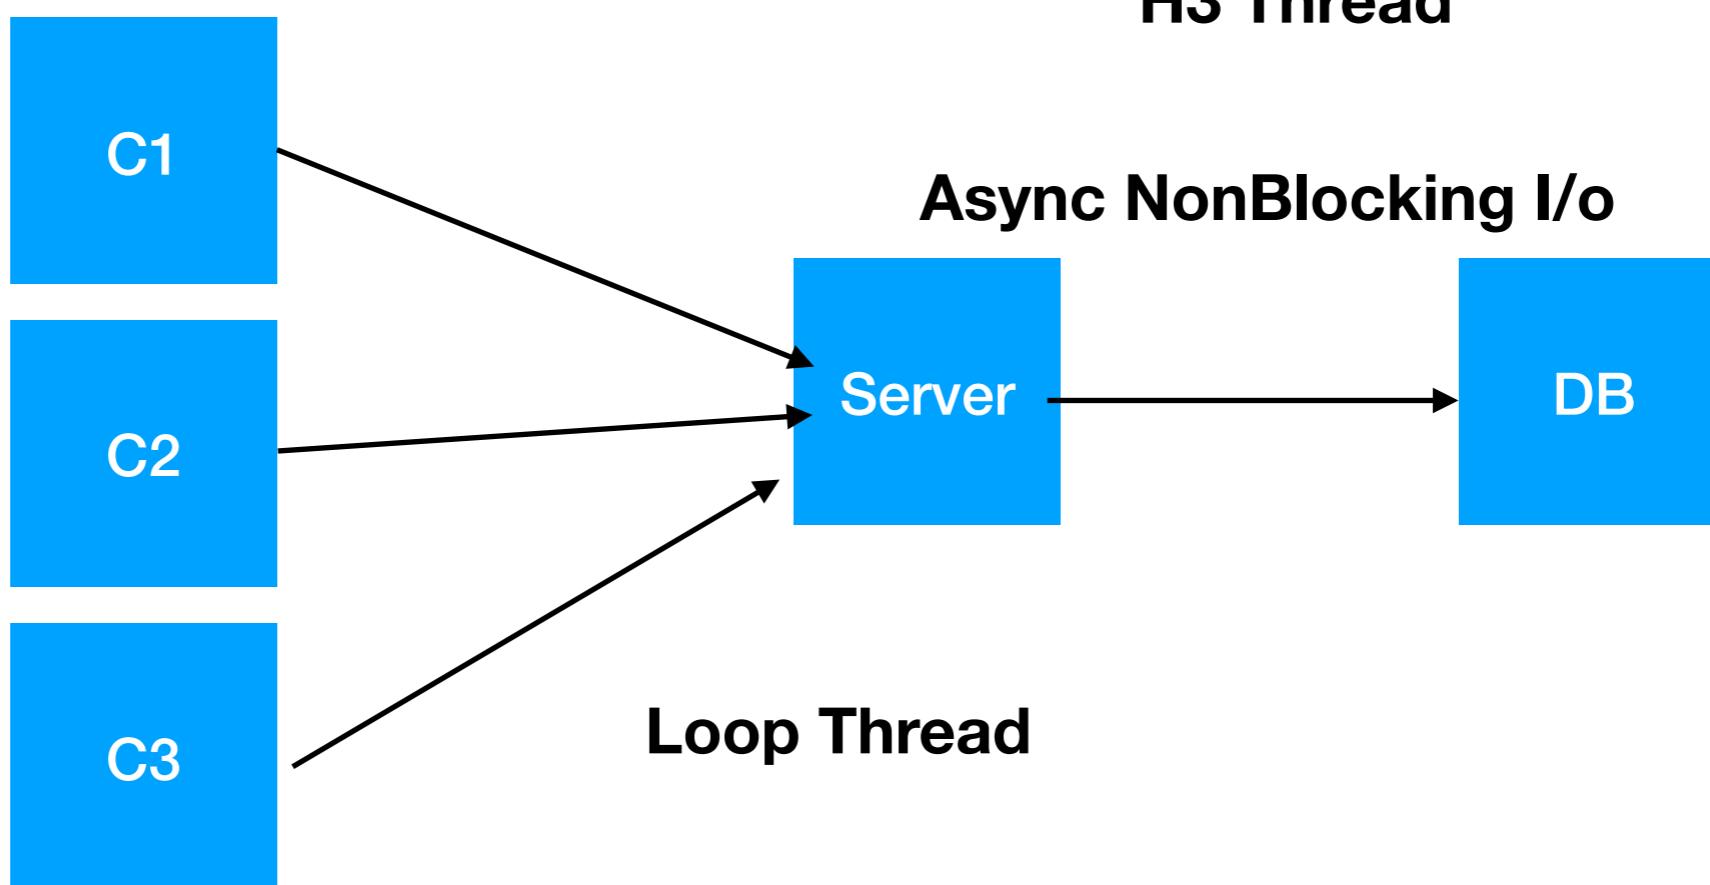
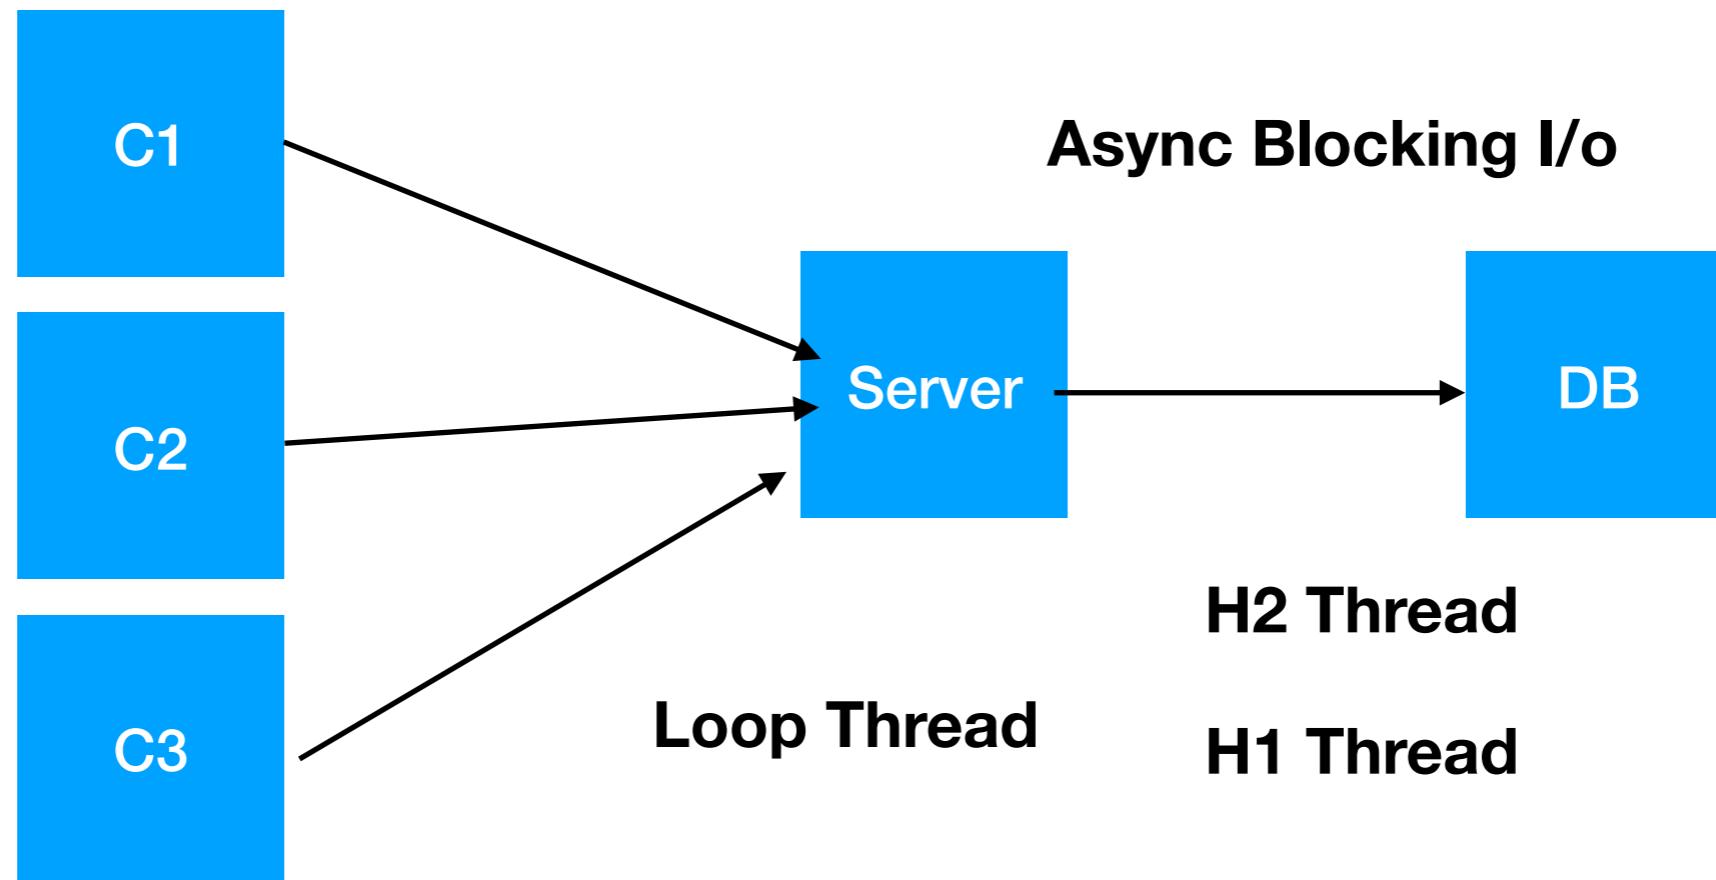


The event demultiplexor delivers the requested events to the event handlers.



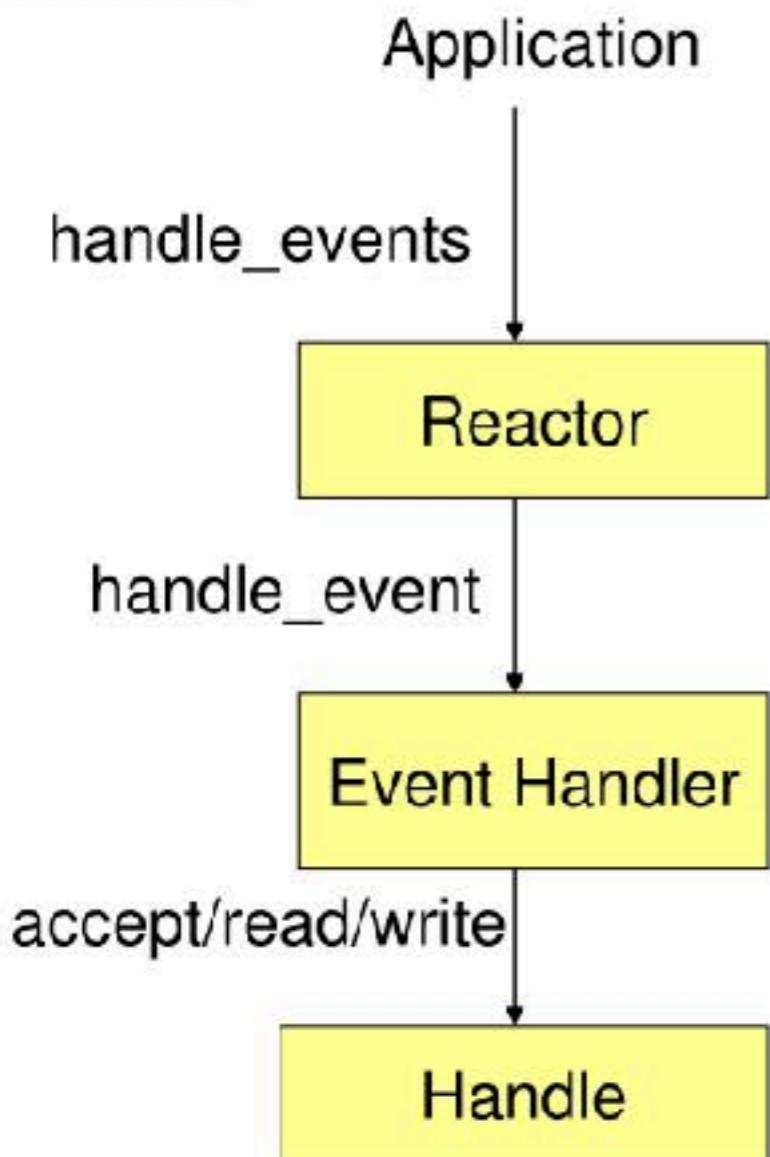
@luminousmen.com



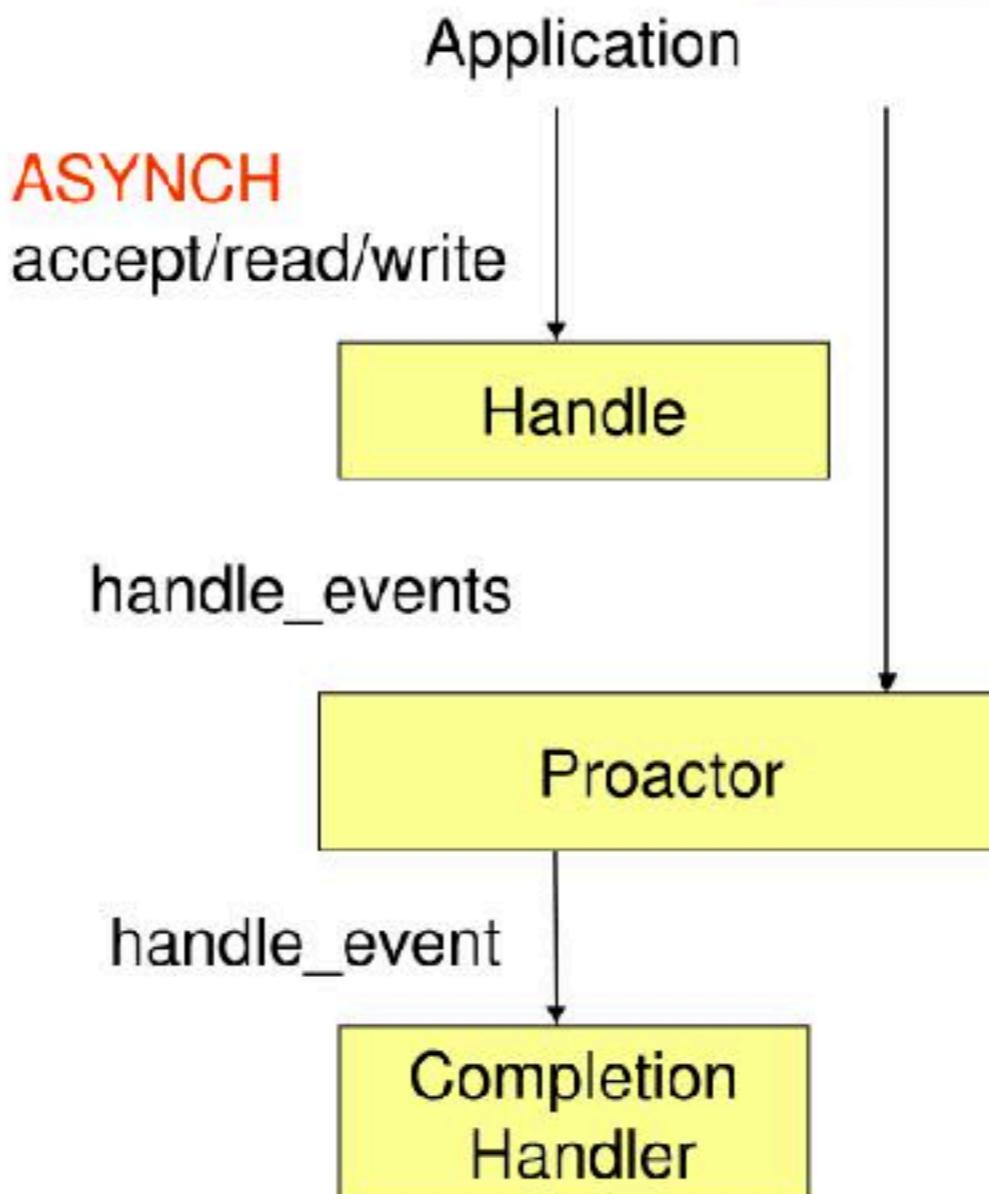


Compare Reactor vs. Proactor Side by Side

Reactor



Proactor



Master Slave

Master

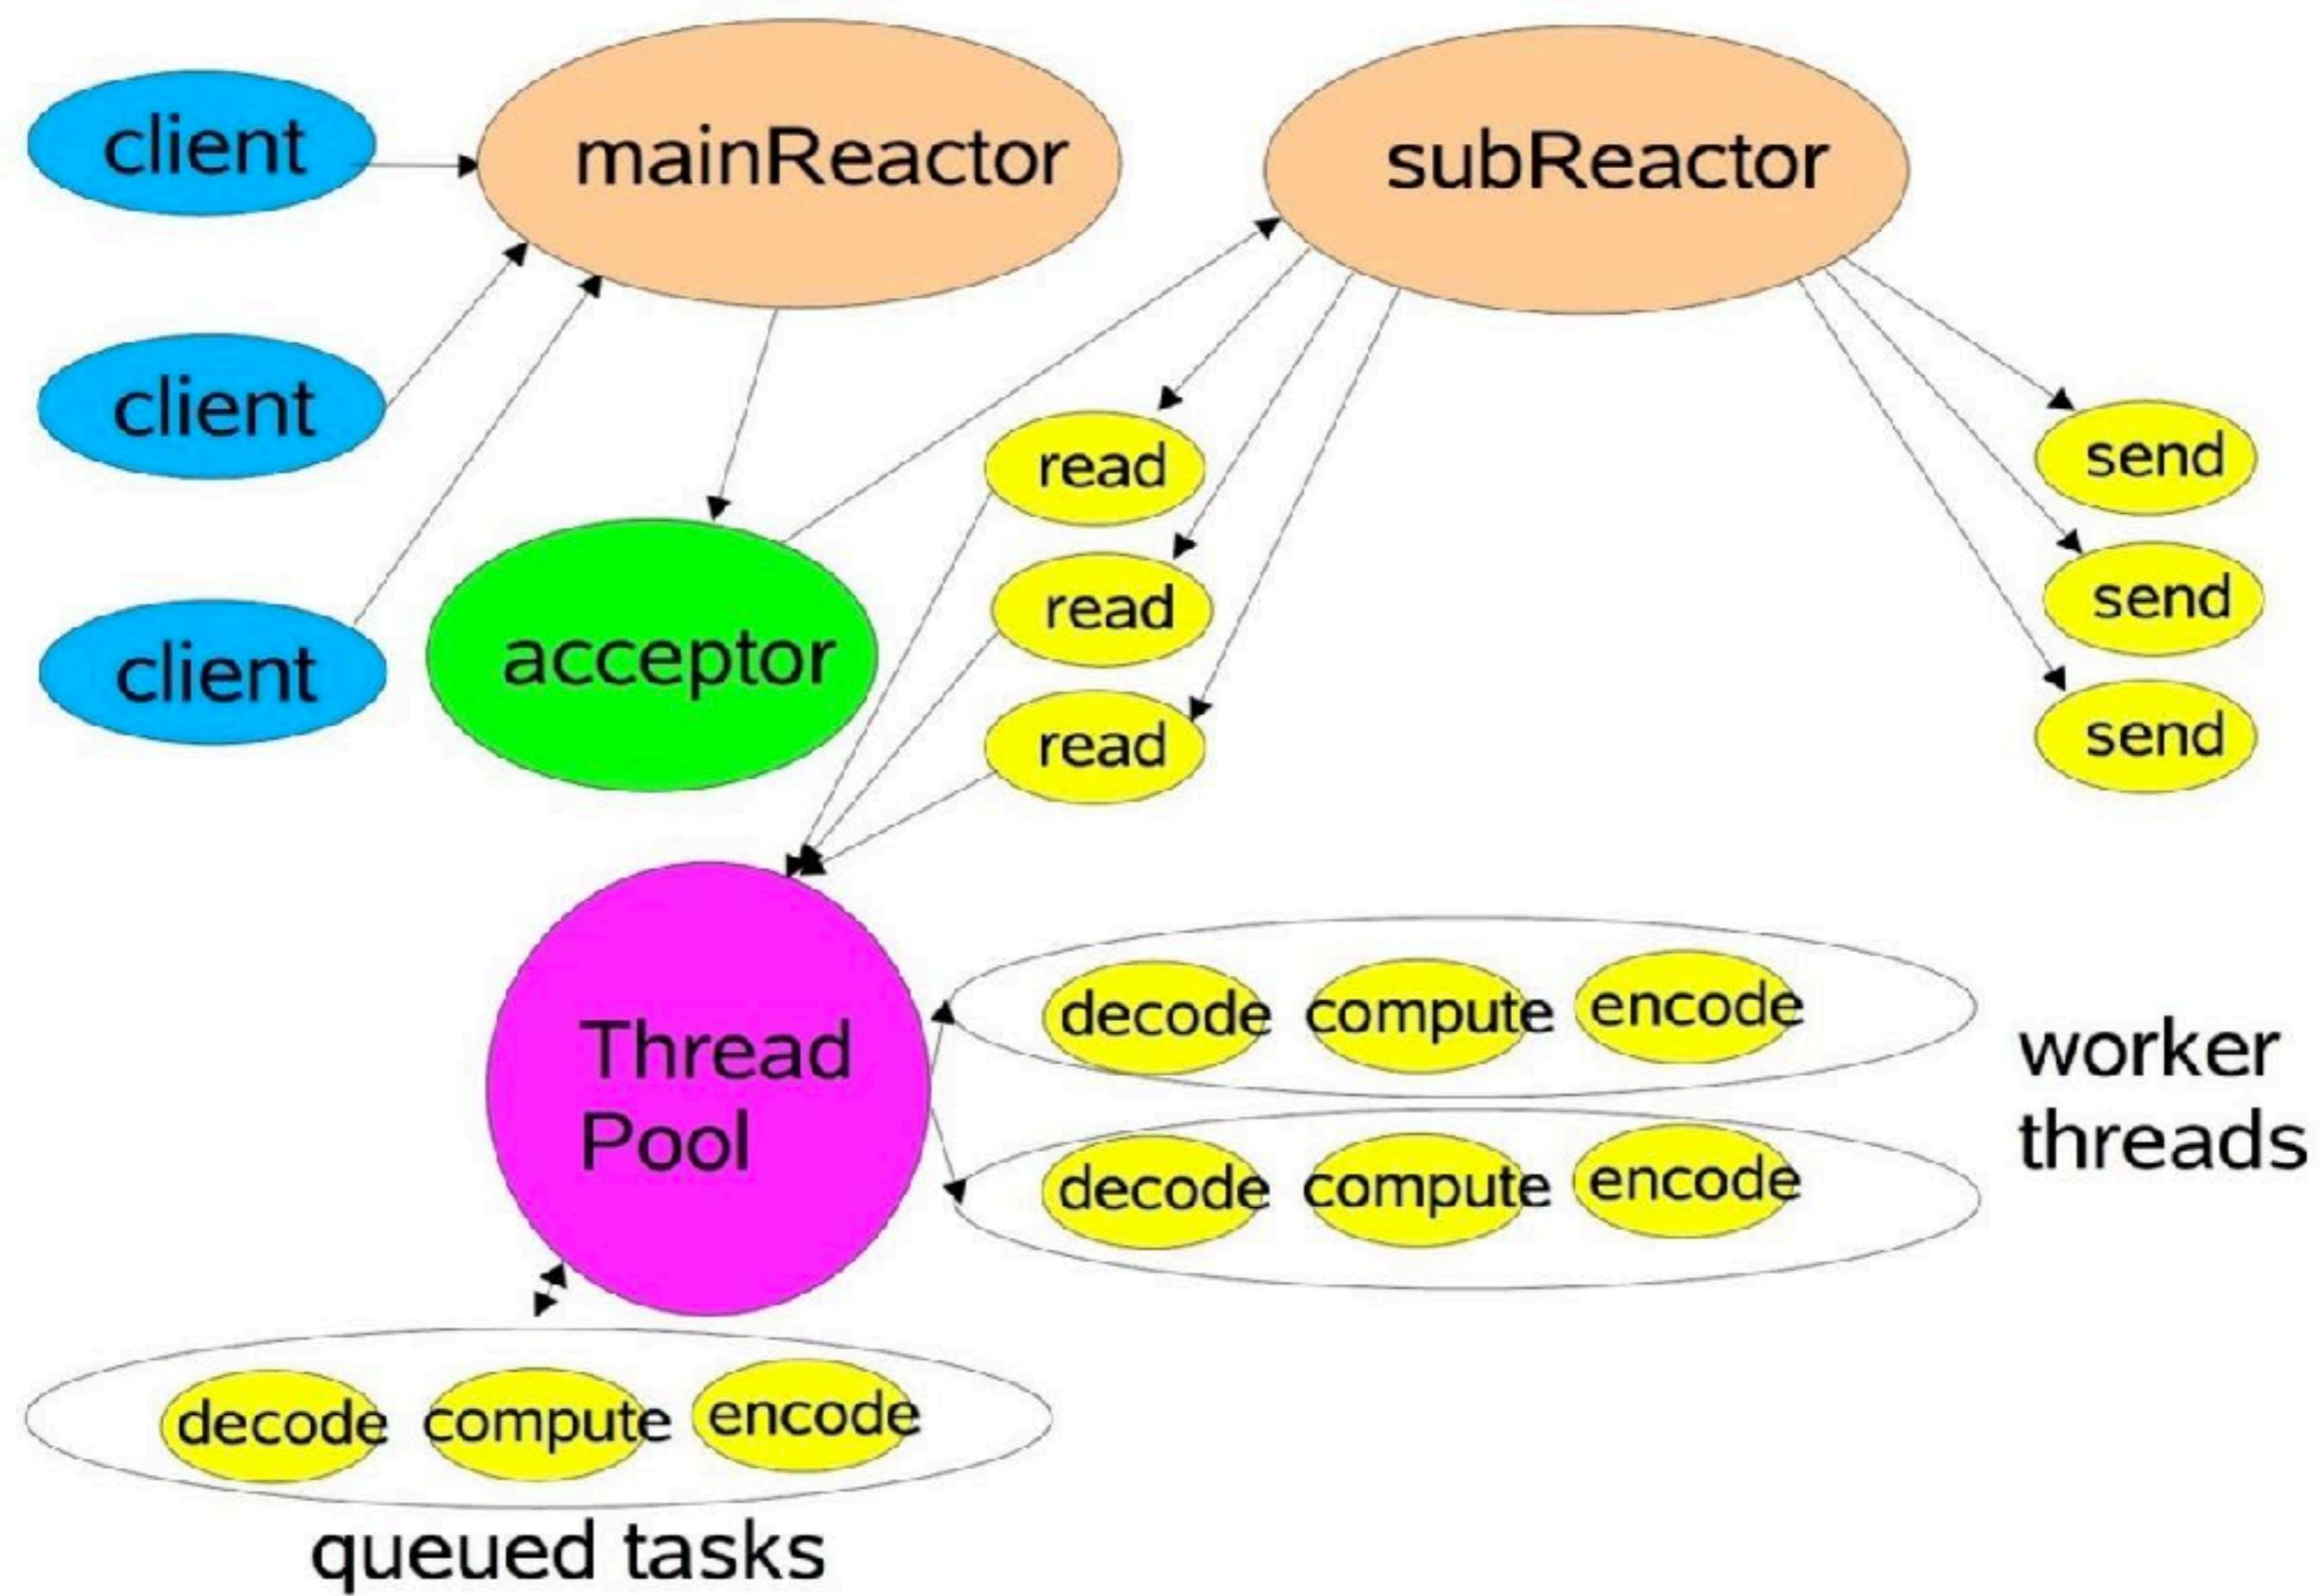
S1

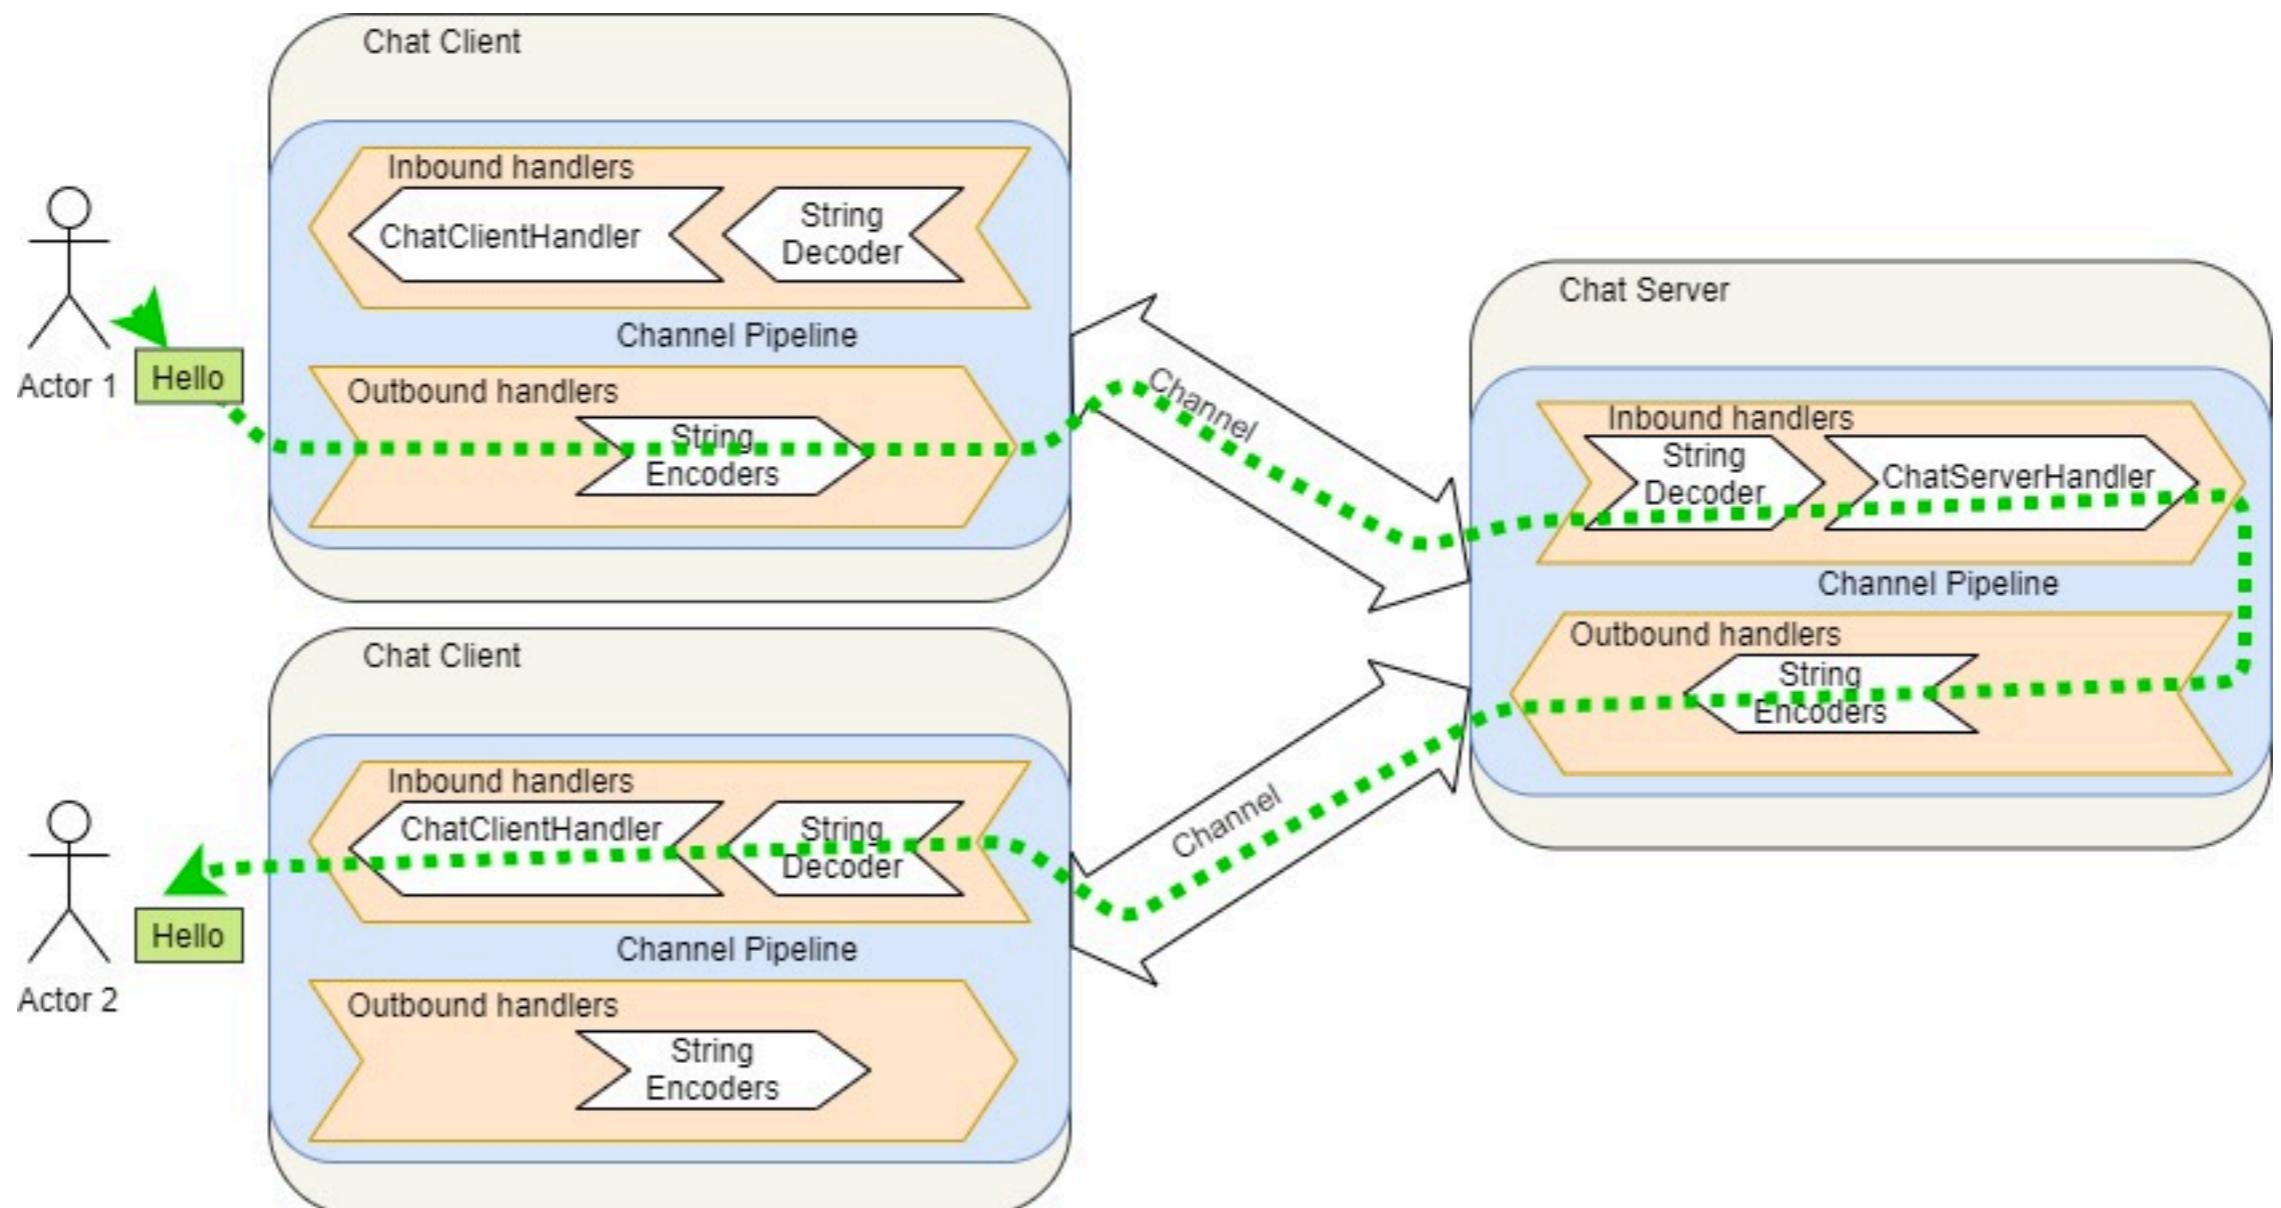
S2

S3

Fork and Join

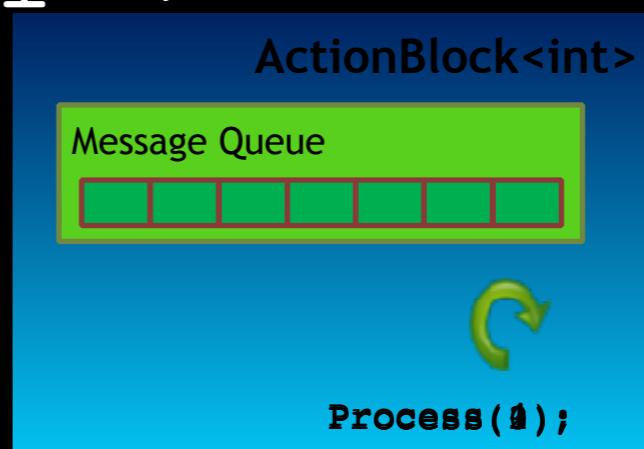




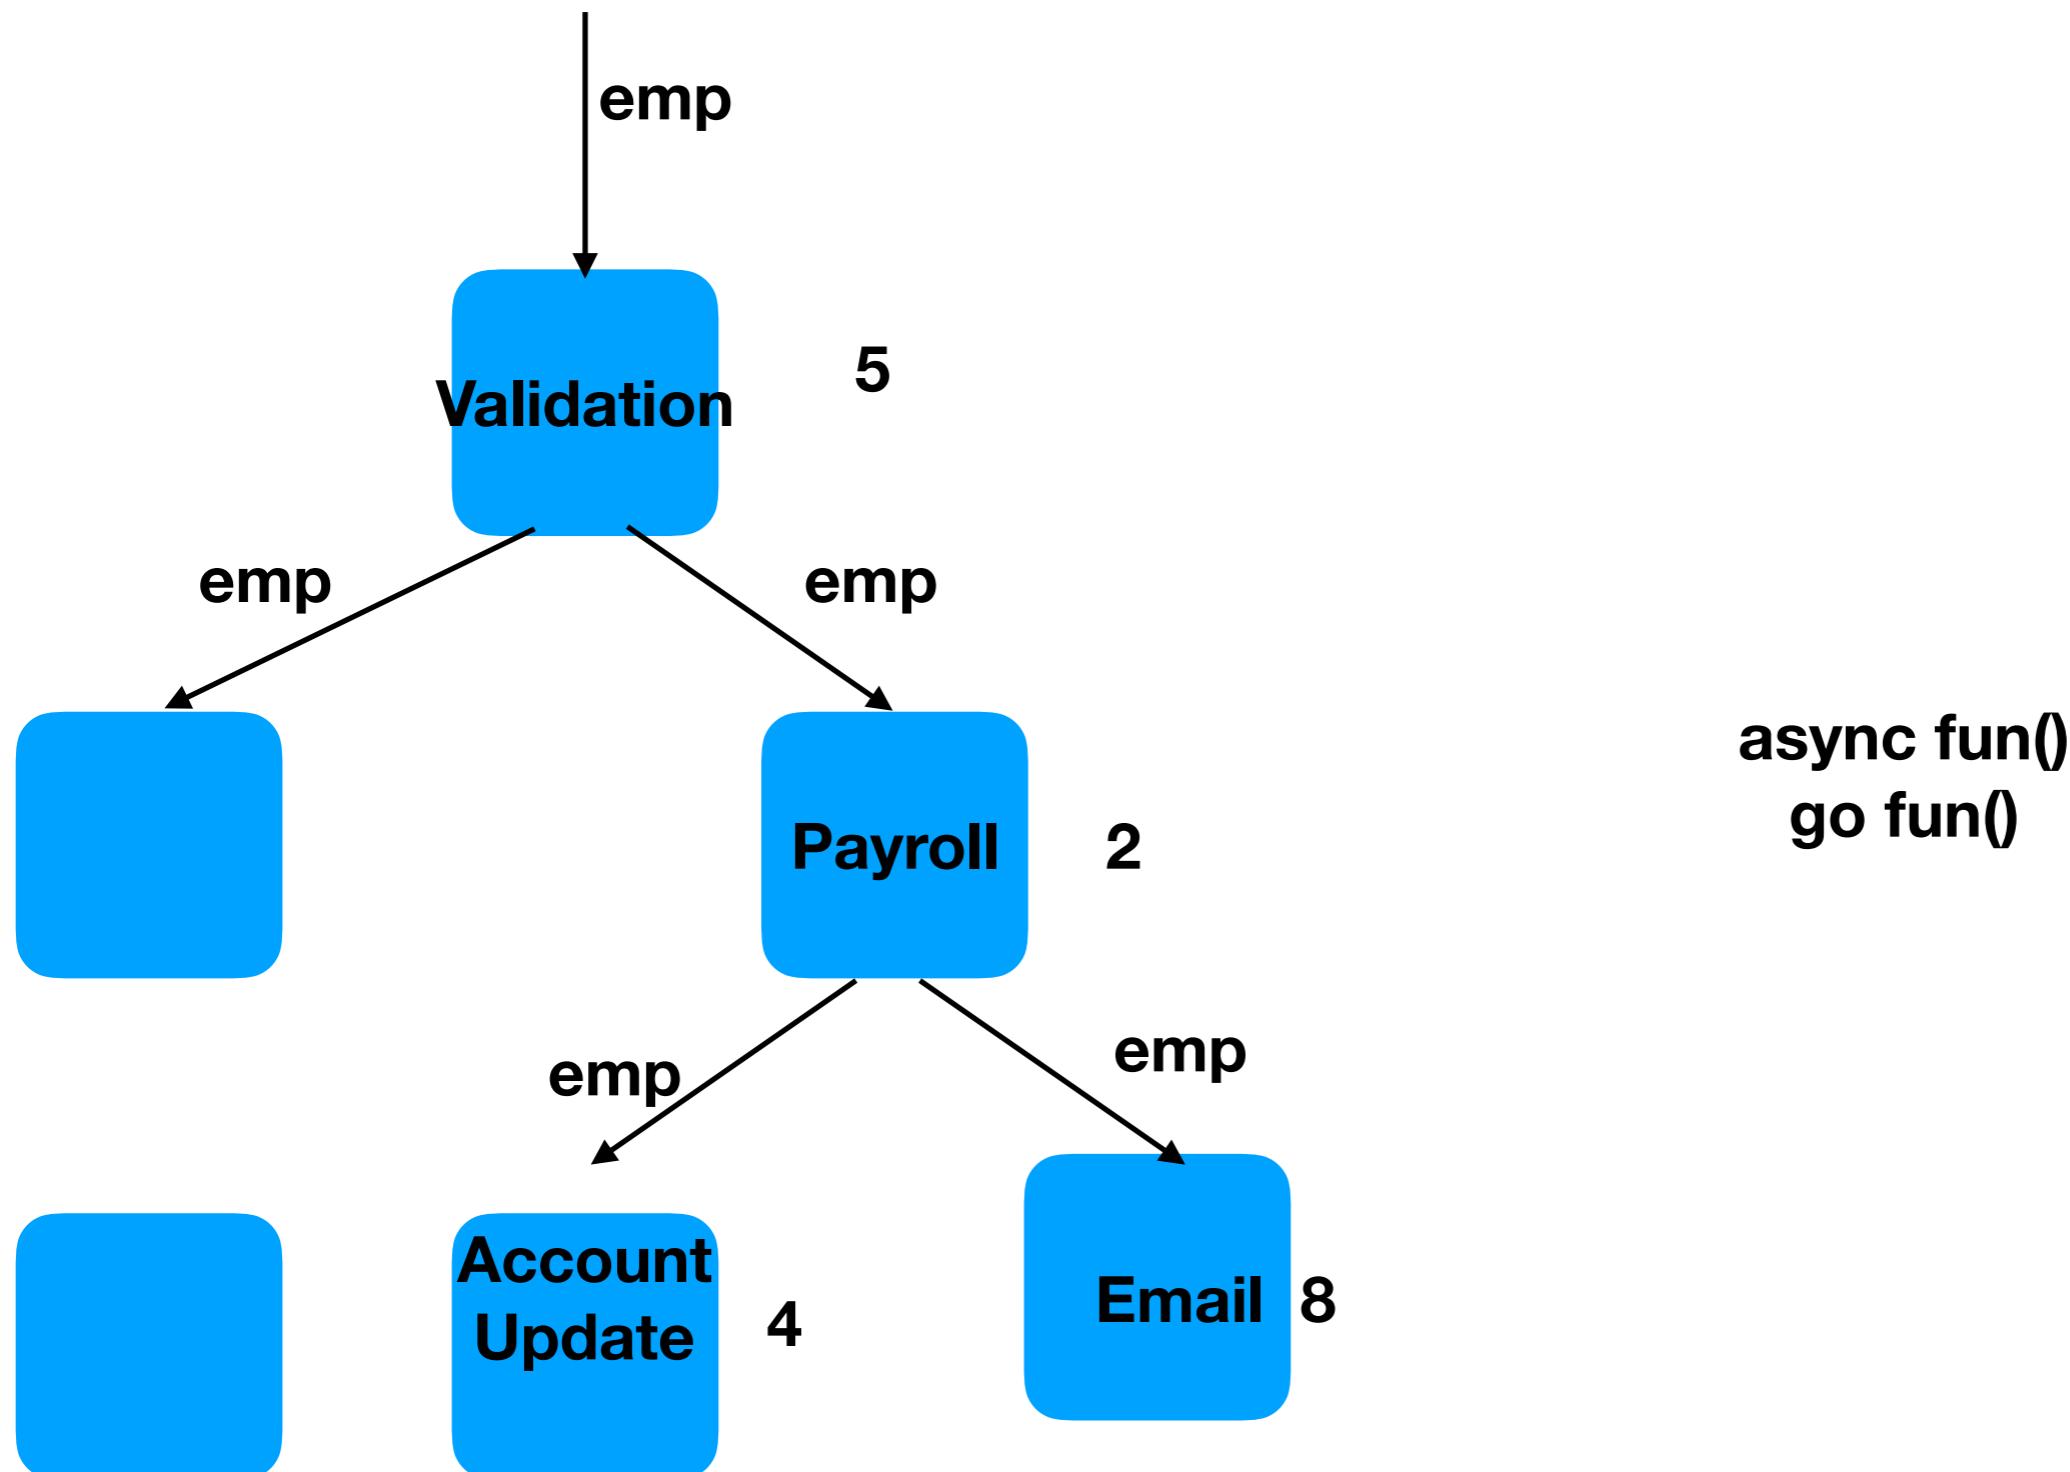


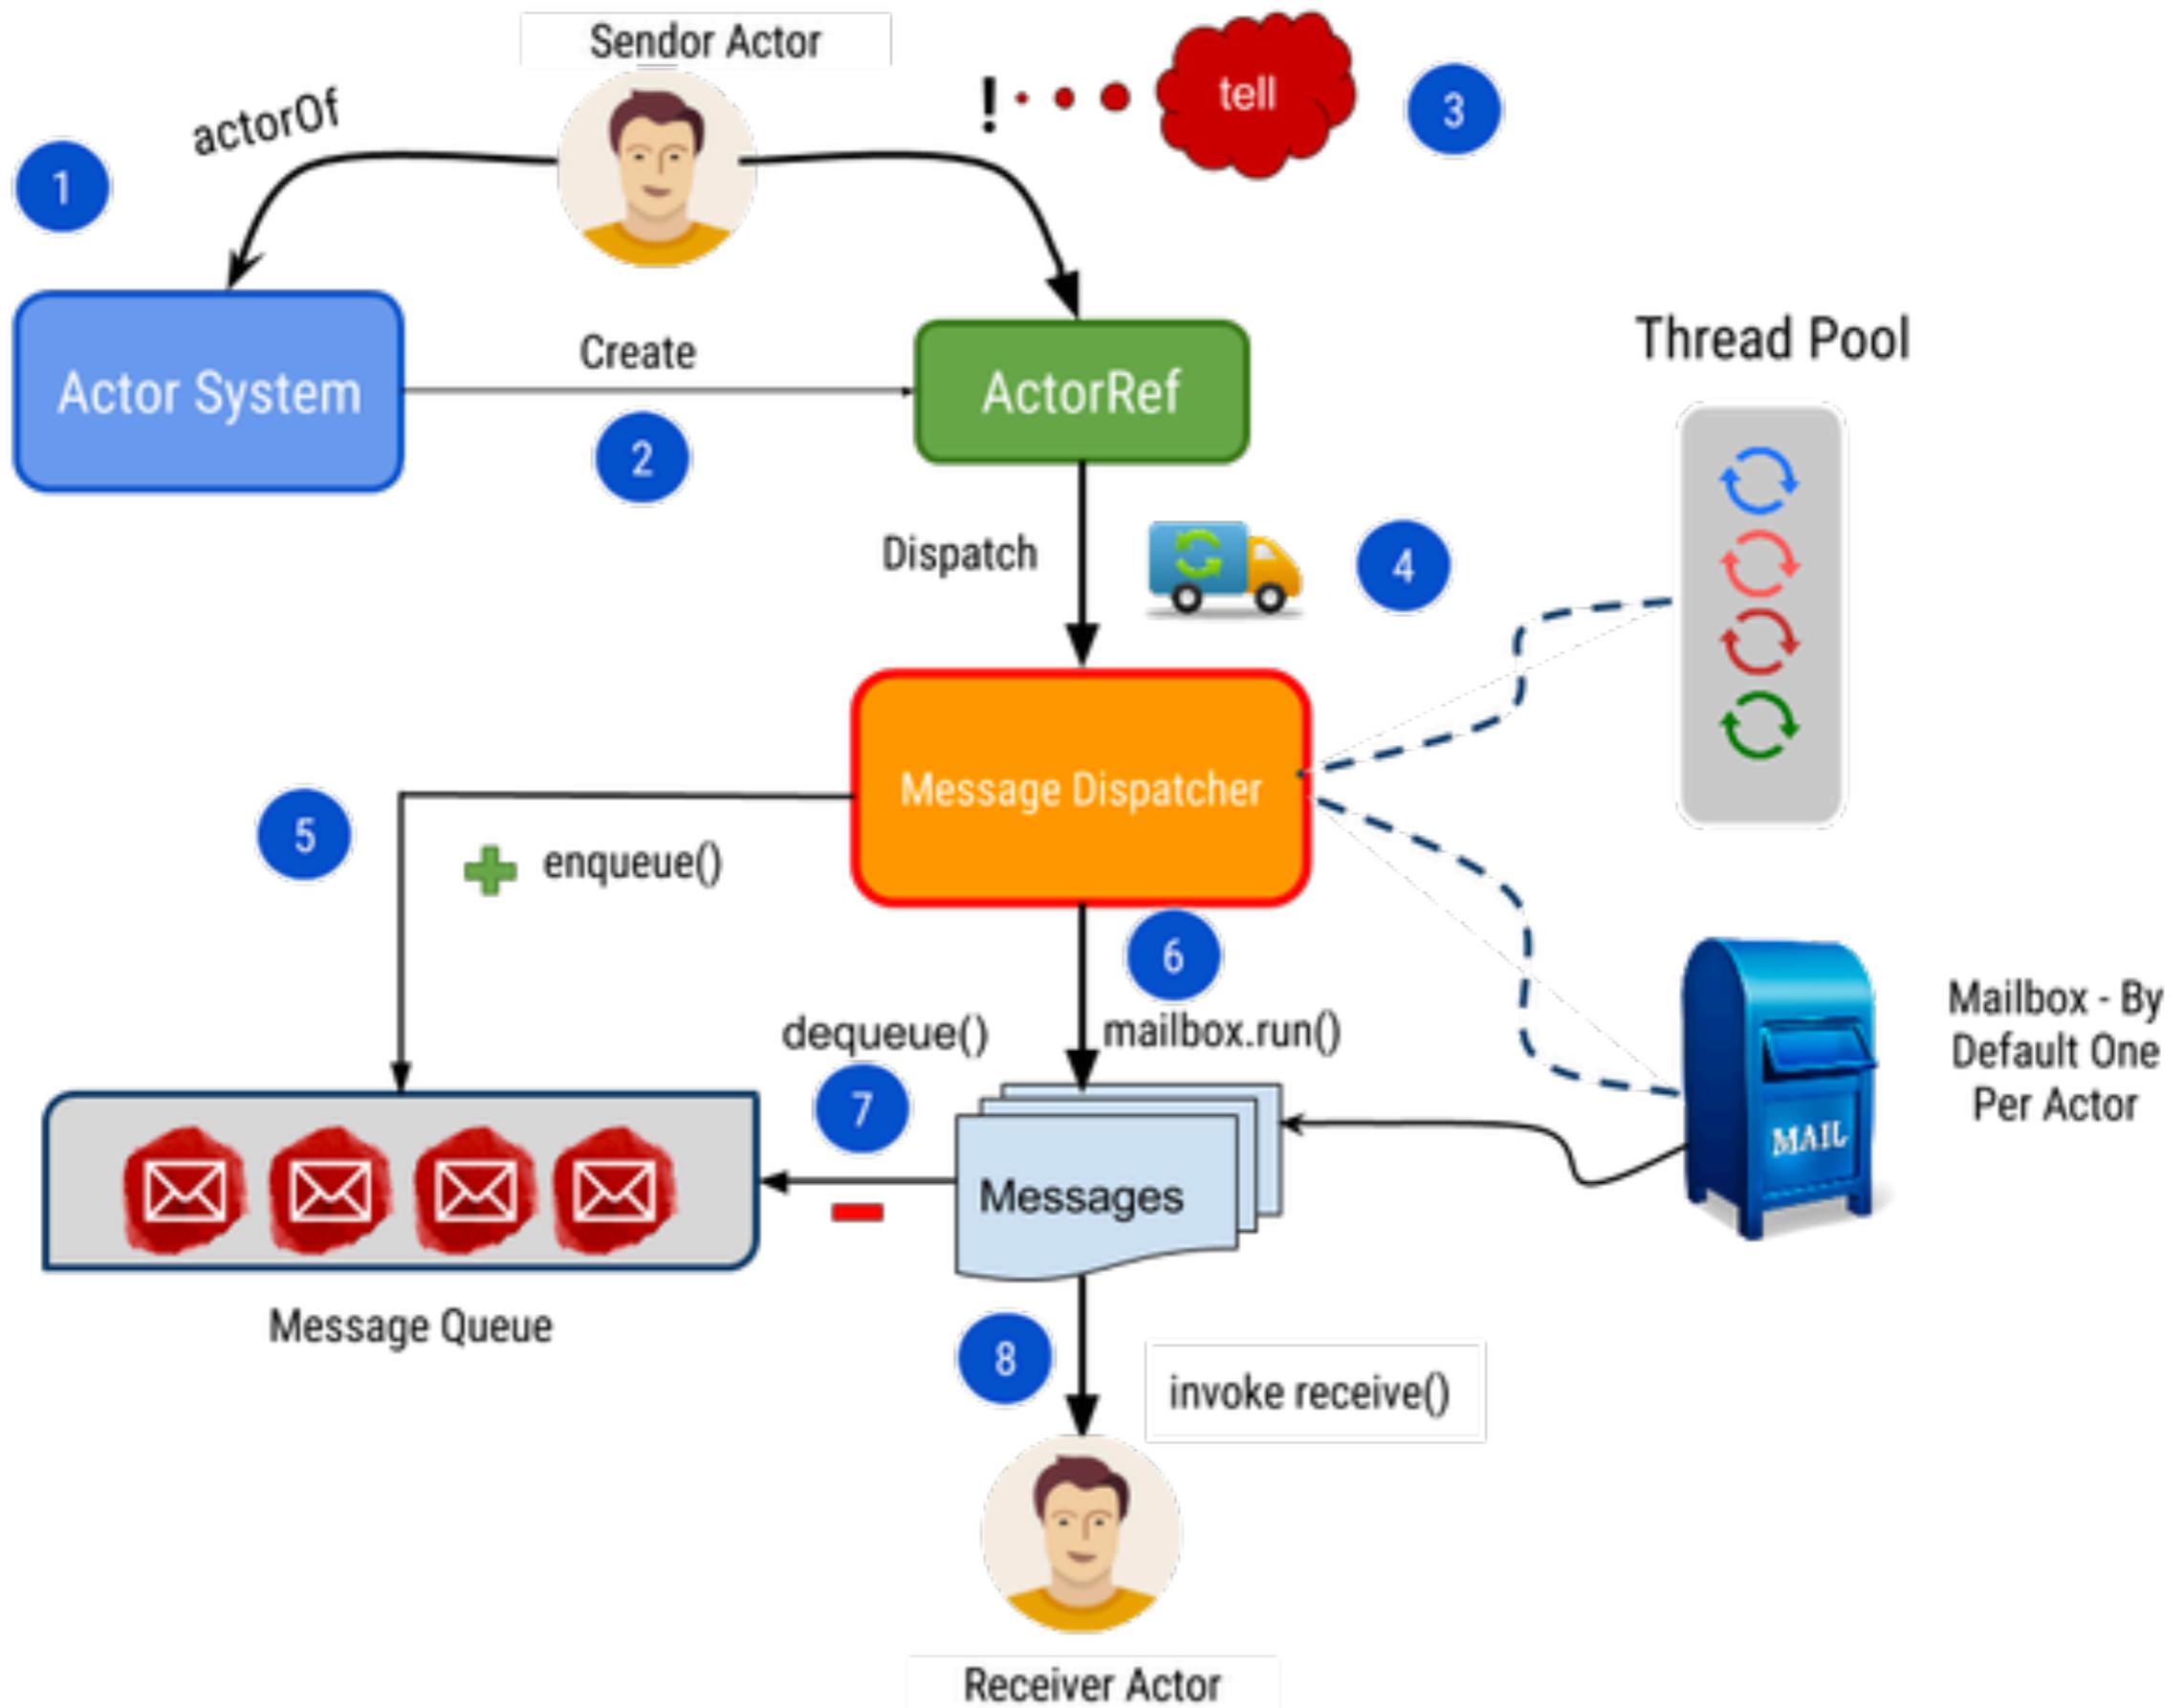
Netty is a low level non-blocking I/O library, used in most reactive solutions including Akka, Vert.x, HBaseAsync, PgAsync.

```
var c = new ActionBlock<int>(i =>
{
    Process(i);
}) ;
```



```
for(int i = 0; i < 5; i++)
{
    c.Post(i);
}
```





Netty is actually used within Akka to provide for distributed actors (although Akka is switching over to Spray for this).

Reactor and Proactor are two event handling patterns

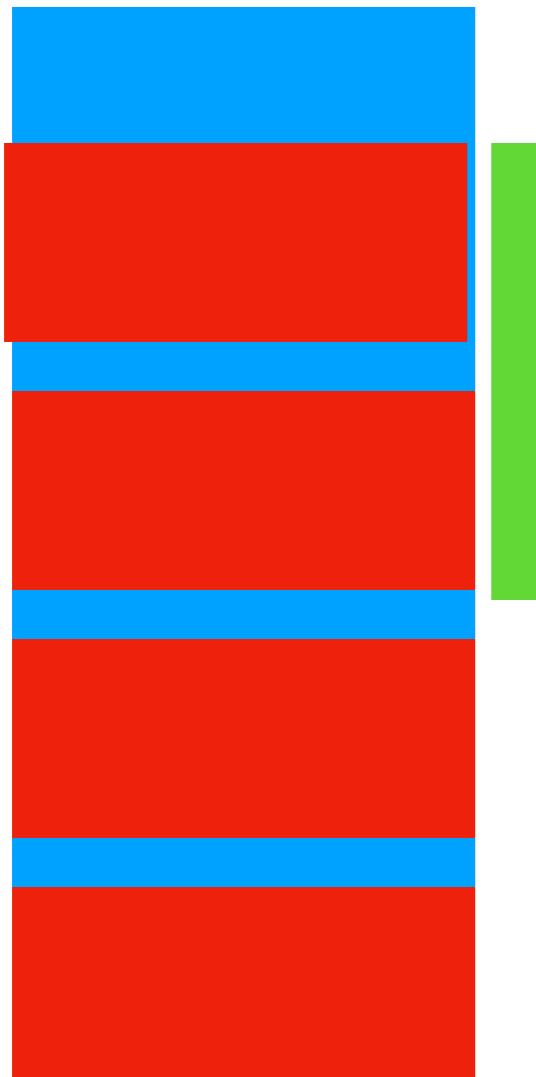
Proactor pattern is an asynchronous version of the Reactor pattern.

Lock-free updates

`x = x * 10;`



Atomic Operation



data ++;

1. Load from RAM into CPU register
2. Increment CPU register
3. Store from CPU register into RAM.



Reading and writing 64-bit fields is non atomic on 32-bit environments because it requires two separate instructions: one for each 32-bit memory location.

Interlocked

- Increment
- Decrement
- Add
 - value can be a negative number allowing subtraction
- Read
 - performs an atomic Read on a 64-bit field:
- `old val = Exchange(ref _sum, new val)`
 - Write a 64-bit field while reading previous value
- `CompareExchange(ref _sum,new val, old val)`
 - Update a field only if it matches a certain value
 - enables lock-free read-modify-write operations

Lock-free updates

1. Take a “snapshot” of x into a local variable
2. Calculate the new value
3. Write the calculated value back if the snapshot is still up-to-date using `Interlocked.CompareExchange`
4. If the snapshot was stale, *spin* and return to step 1

Thread-safe collections

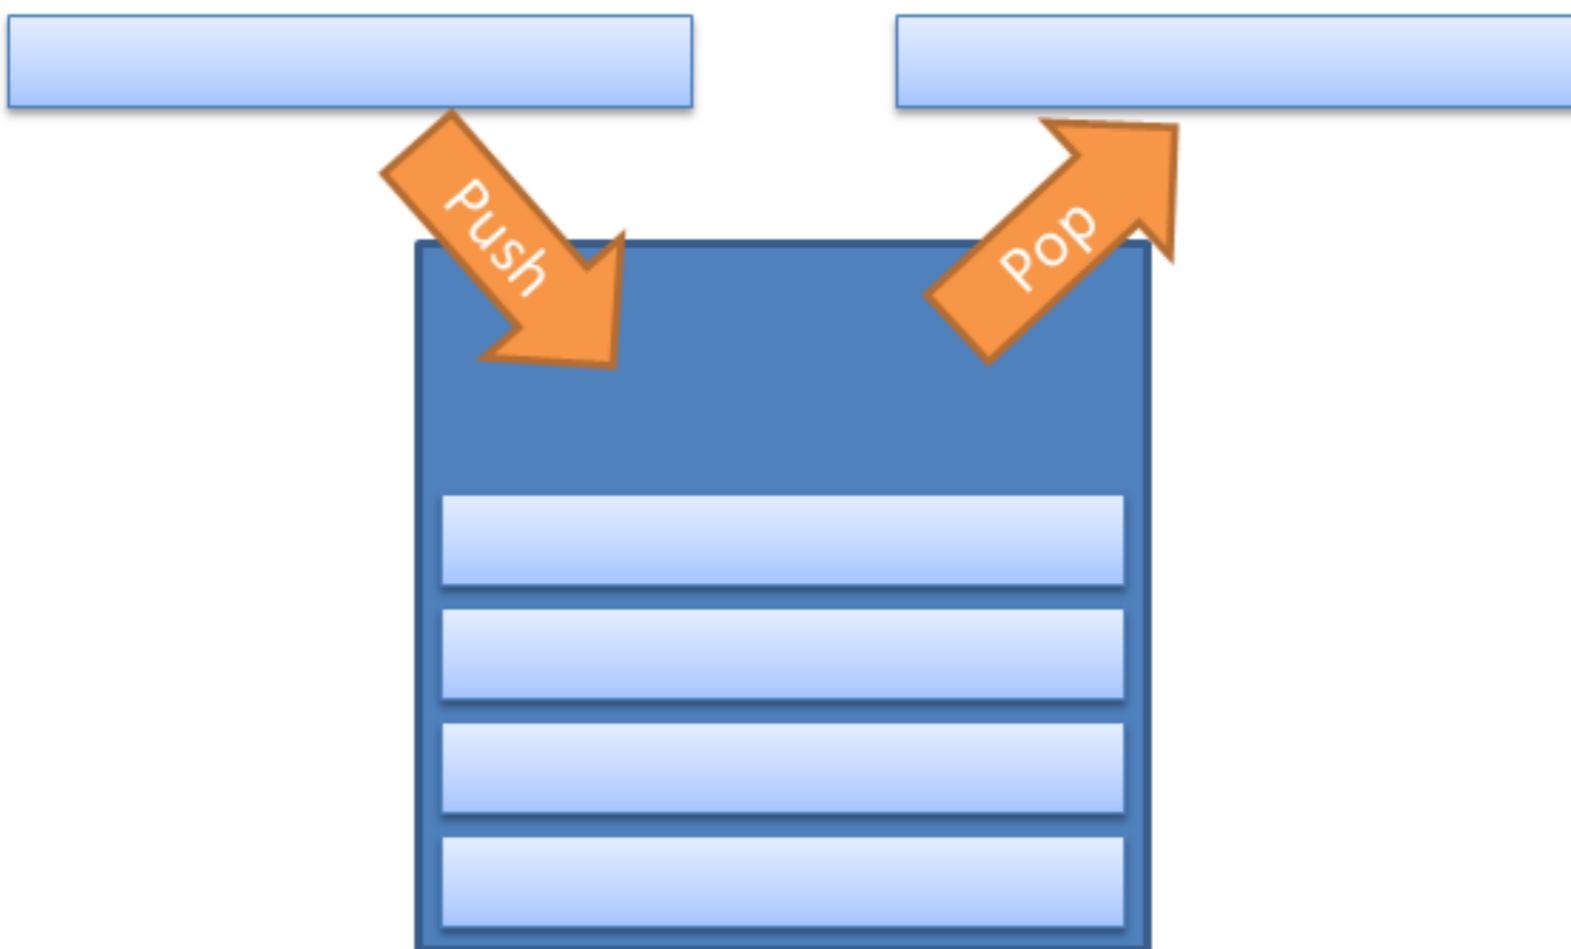
`ConcurrentBag<T>`

`ConcurrentStack<T>`

`ConcurrentQueue<T>`

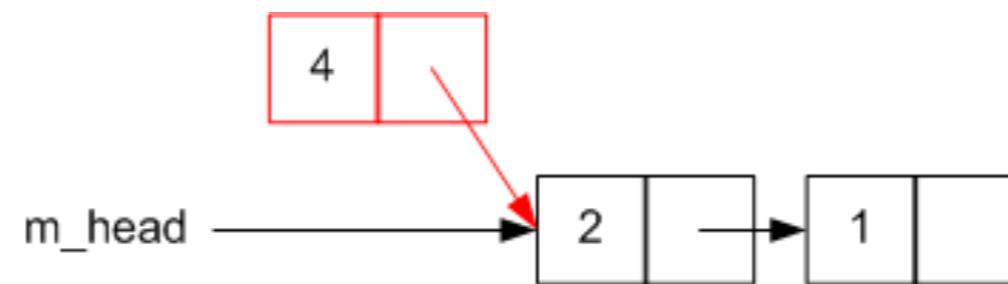
`ConcurrentDictionary< TKey, TValue >`

Concurrent Stack

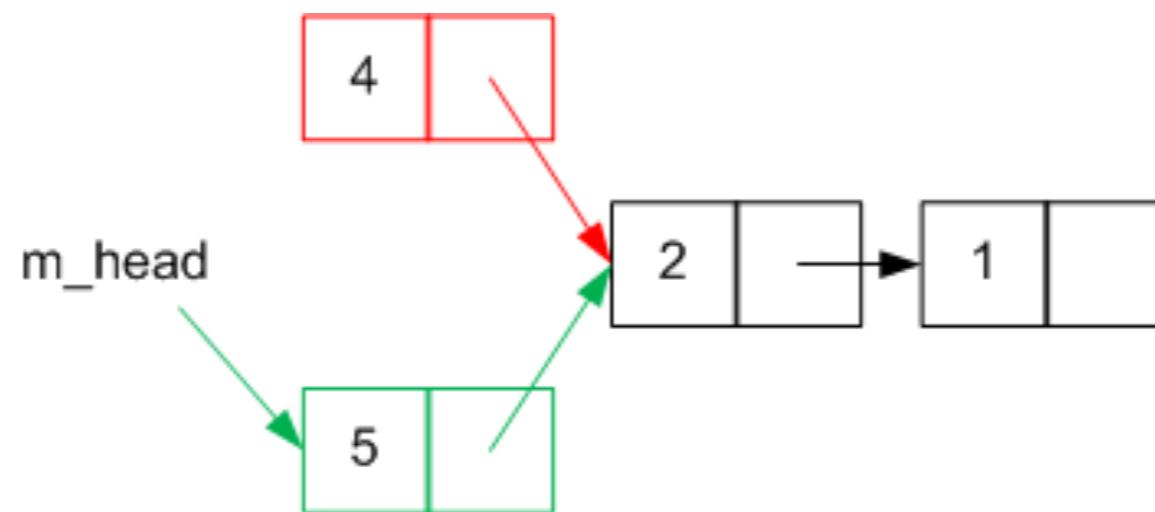


Adding a item to stack

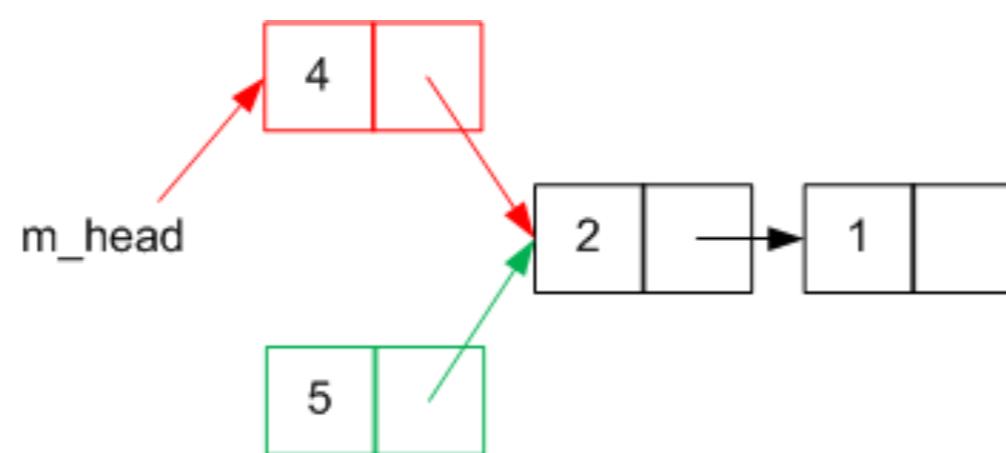
Thread 1, cycle 1 ,
executes 1,2



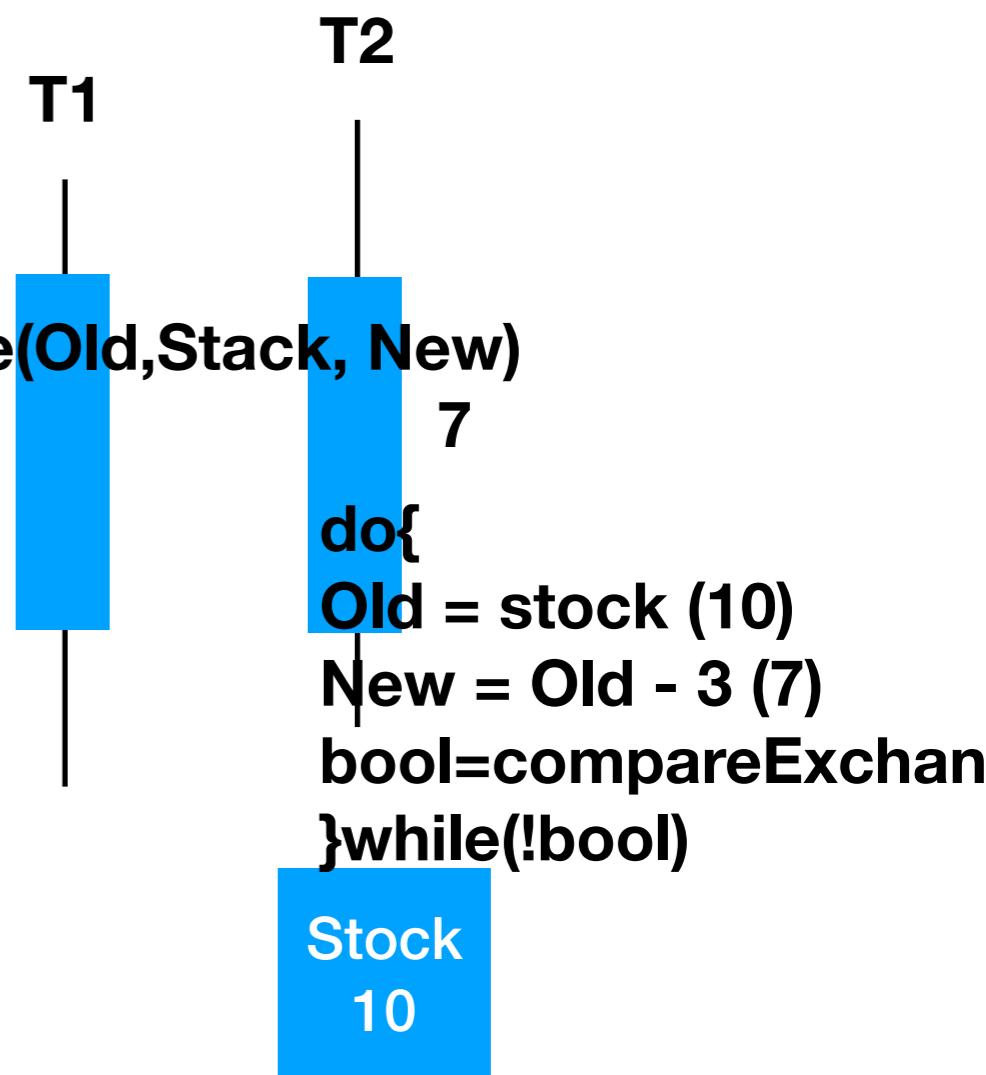
Thread 2, cycle 2 ,
executes 1,2,3



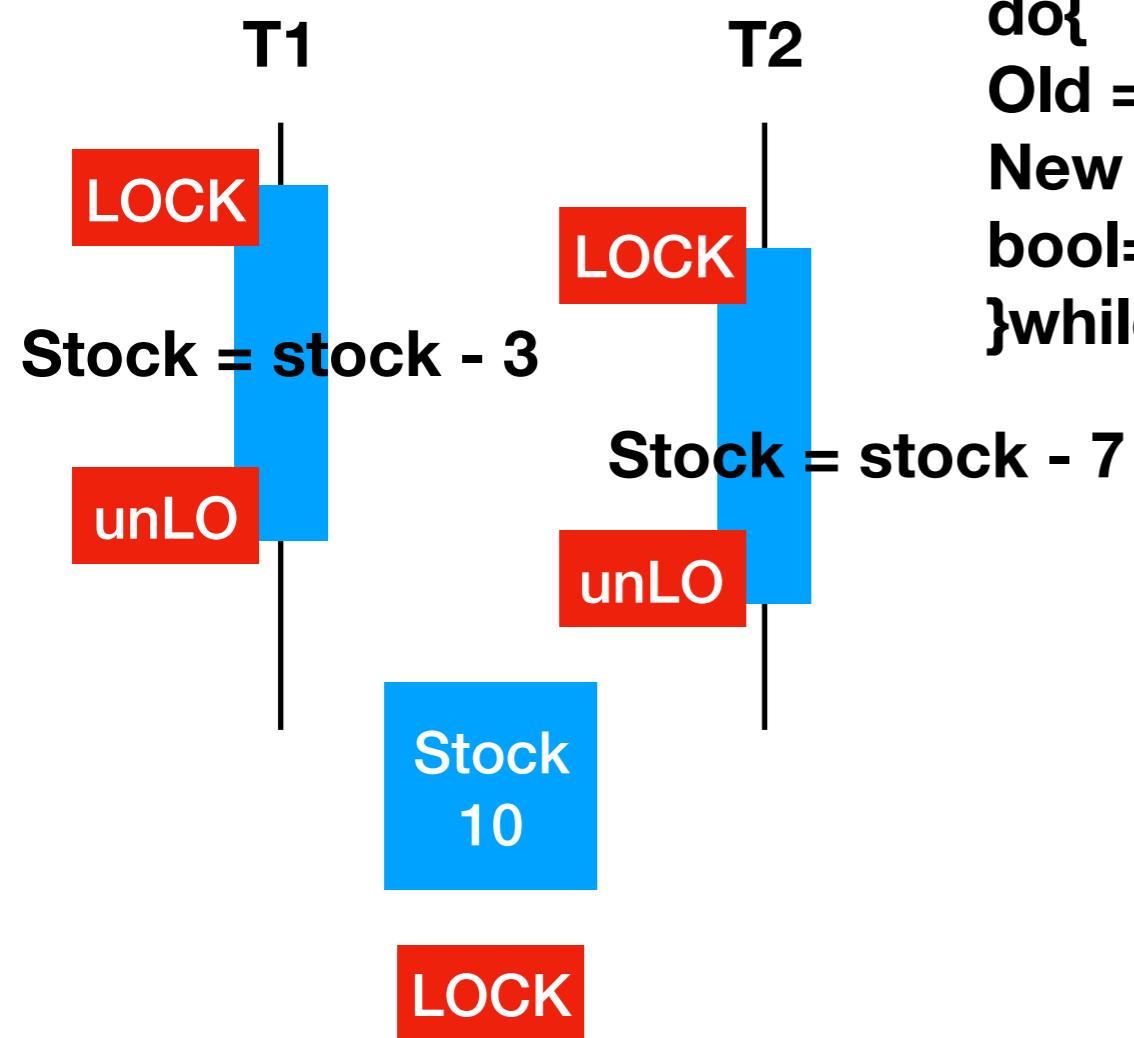
Thread 1, cycle 3 ,
executes 3



Optimistic Approach



Pessimistic Approach



Pessimistic Approach



Optimistic Approach

T1

3

```
do{
    Old = stock (10)
    New = Old - 3 (7)
    bool=compareExchange(Old,Stack, New)
    if(!bool)
        Spin(10000);
}while(!bool)
```

T2

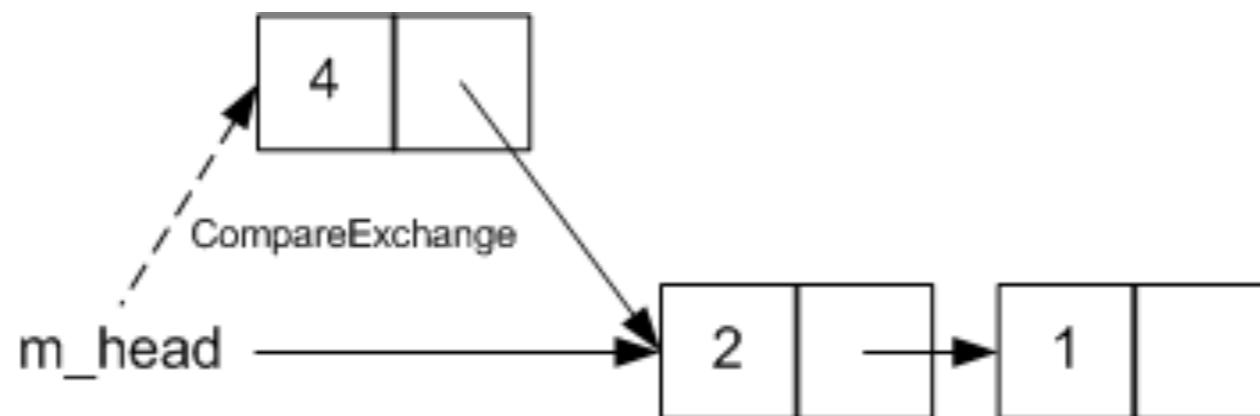
7

```
do{
    Old = stock (10)
    New = Old - 7 (3)
    bool=compareExchange(Old,Stack, New)
}while(!bool)
```



Adding a item to Concurrent Stack

```
public void Push(T item)
{
    Node node = new Node(item);      //step 1
    do {
        node.m_next = m_head;      //step 2
    } while (
        Interlocked.CompareExchange(ref m_head, node, node.m_next) !=
        node.m_next);               //step 3
}
```



Perform step 3 *only if the head hasn't changed in the meantime*. If it has been changed, then go back to step 2 to use the updated head node as the value of m_next.

ConcurrentQueue

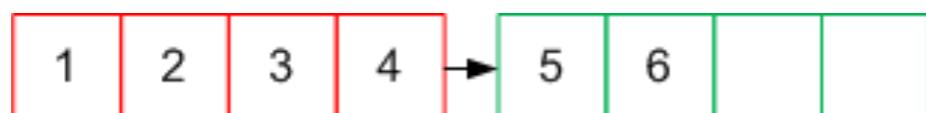


ConcurrentQueue



`m_head`
`m_tail`

three items are enqueued



`m_head`

`m_tail`

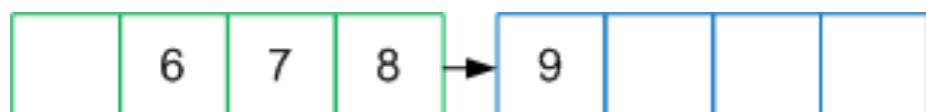
Three more items are then enqueue



`m_head`

`m_tail`

And another three



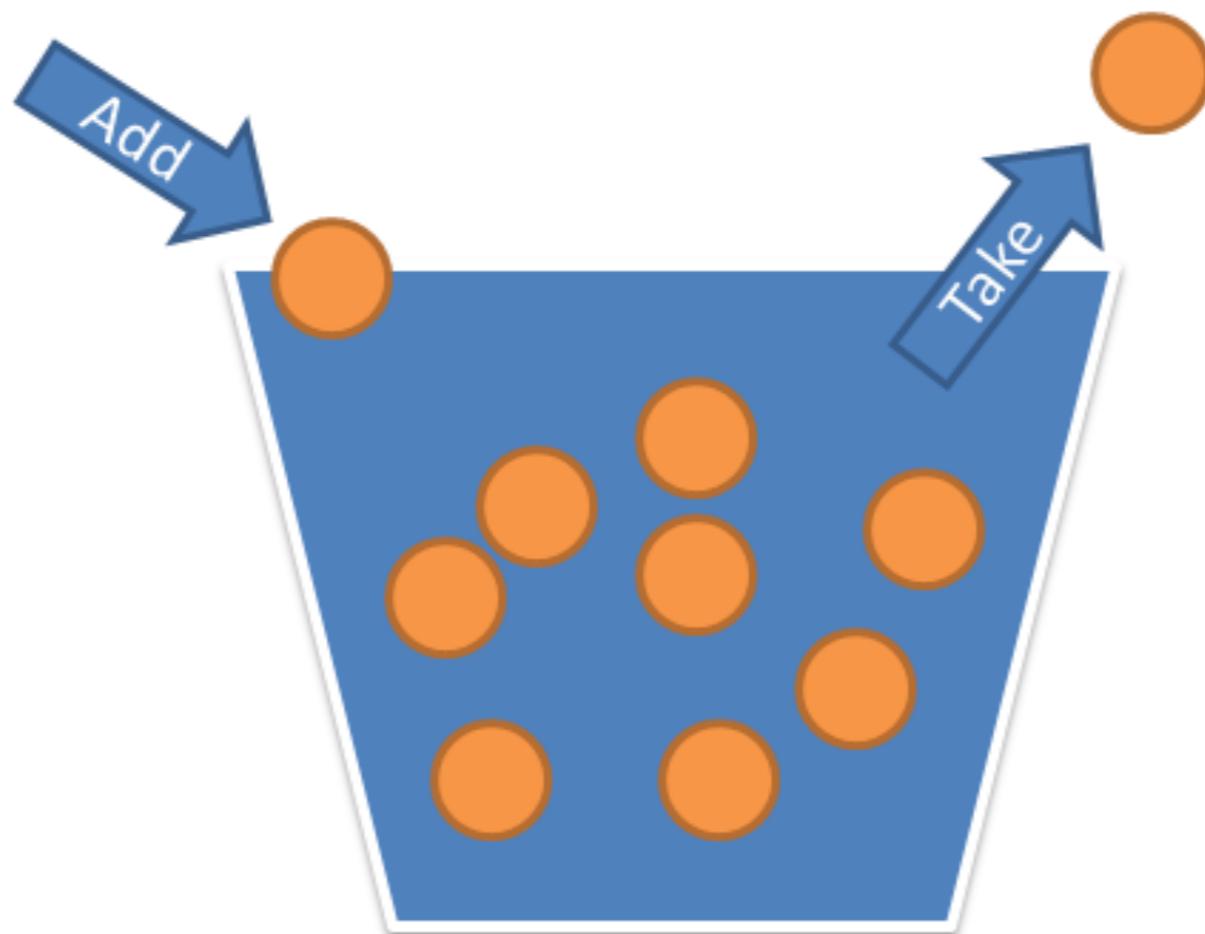
`m_head`

`m_tail`

Five items are then dequeued from the head.

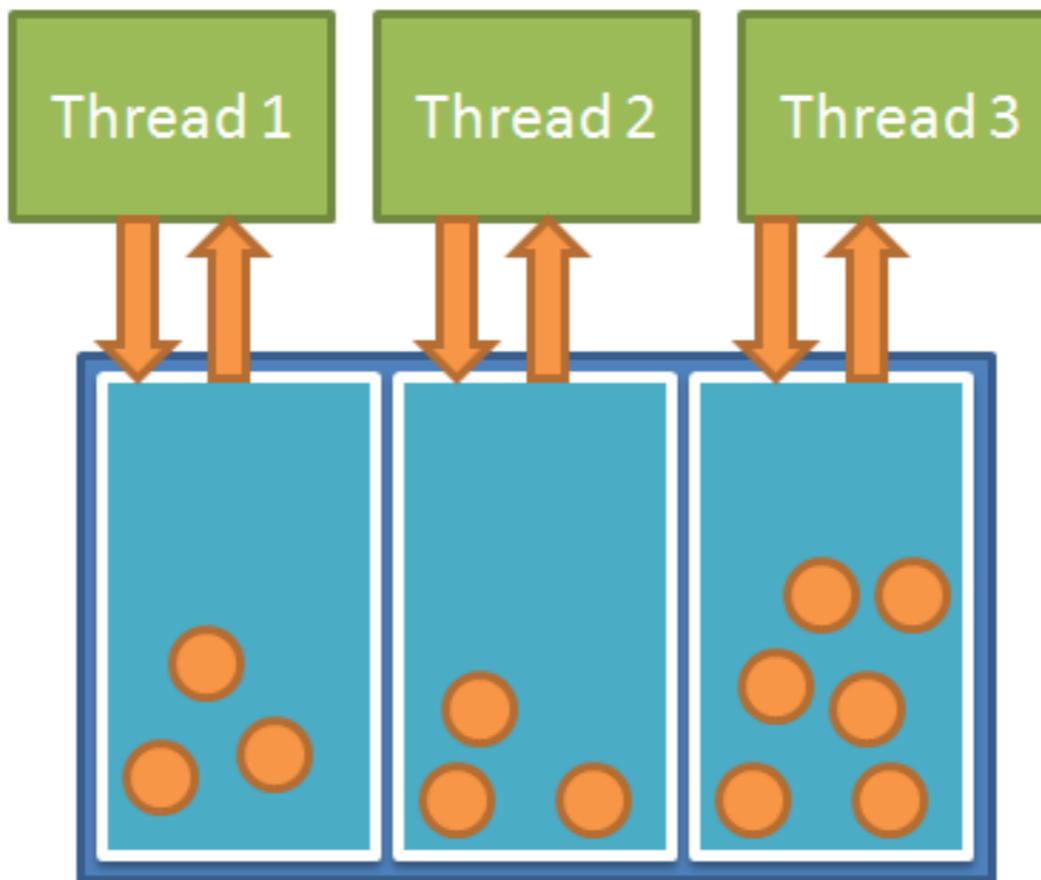
ConcurrentBag

ConcurrentHashMultiset



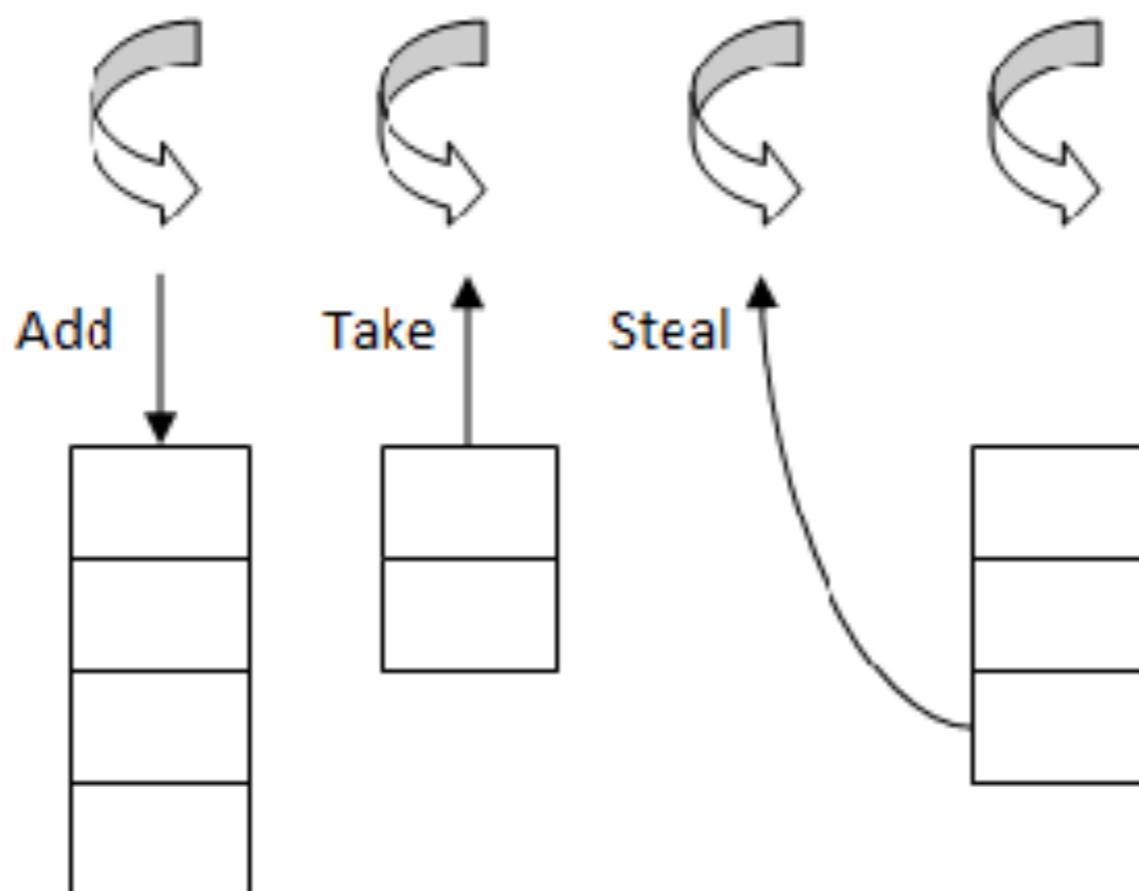
a thread-safe, unordered collection of objects

concurrent bag



if local queue is empty, it will then perform some locking and look at one of the queues from another thread.

concurrent bag



```
static void Main()
{
    bool complete = false;
    var t = new Thread(() =>
    {
        int i=0;
        while (!complete)
            i++;
        Console.WriteLine(i);
    });
    t.Start();
    Thread.Sleep(1000);
    complete = true;
    t.Join();
}
```

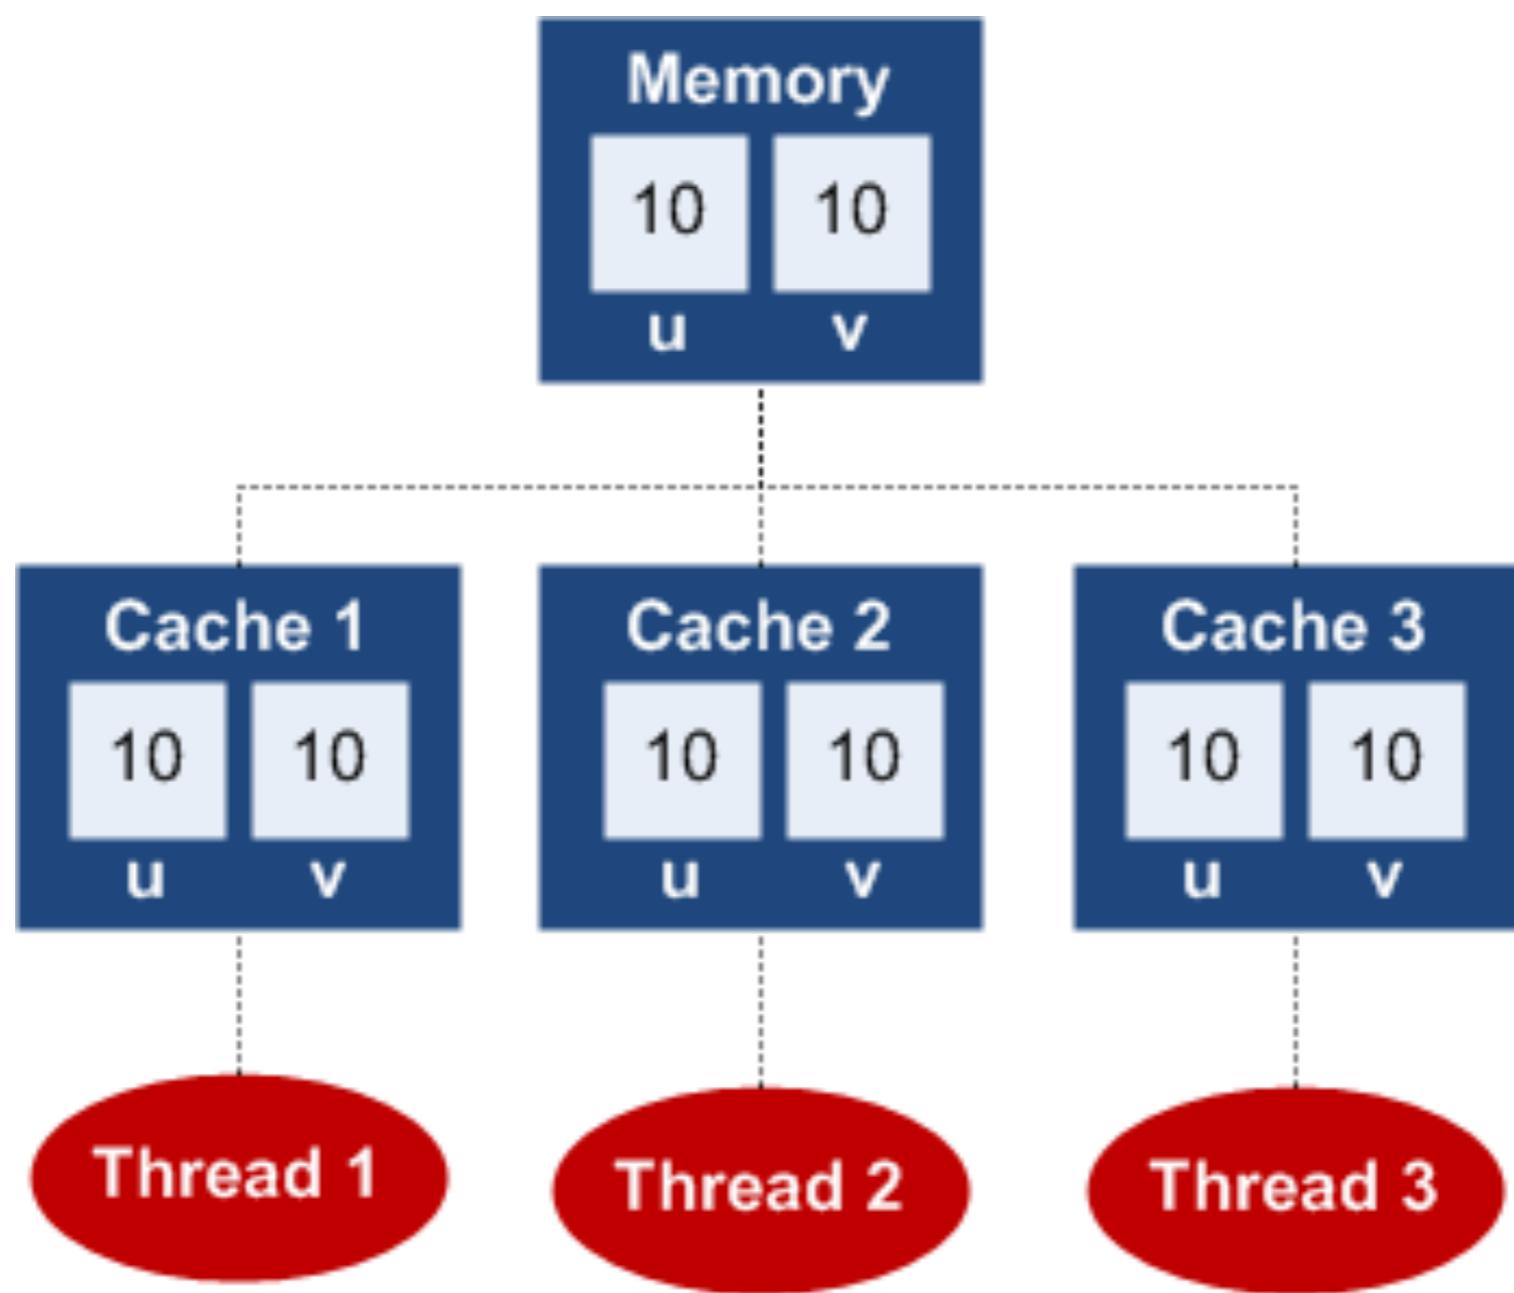
```
class Foo
{
    int _answer=100;
    bool _complete=false;
    void A()
    {
        _answer = 123;
        _complete = true;
    }
    void B()
    {
        if(_complete)
        {
            Console.WriteLine (_ answer);
        }
    }
}
```

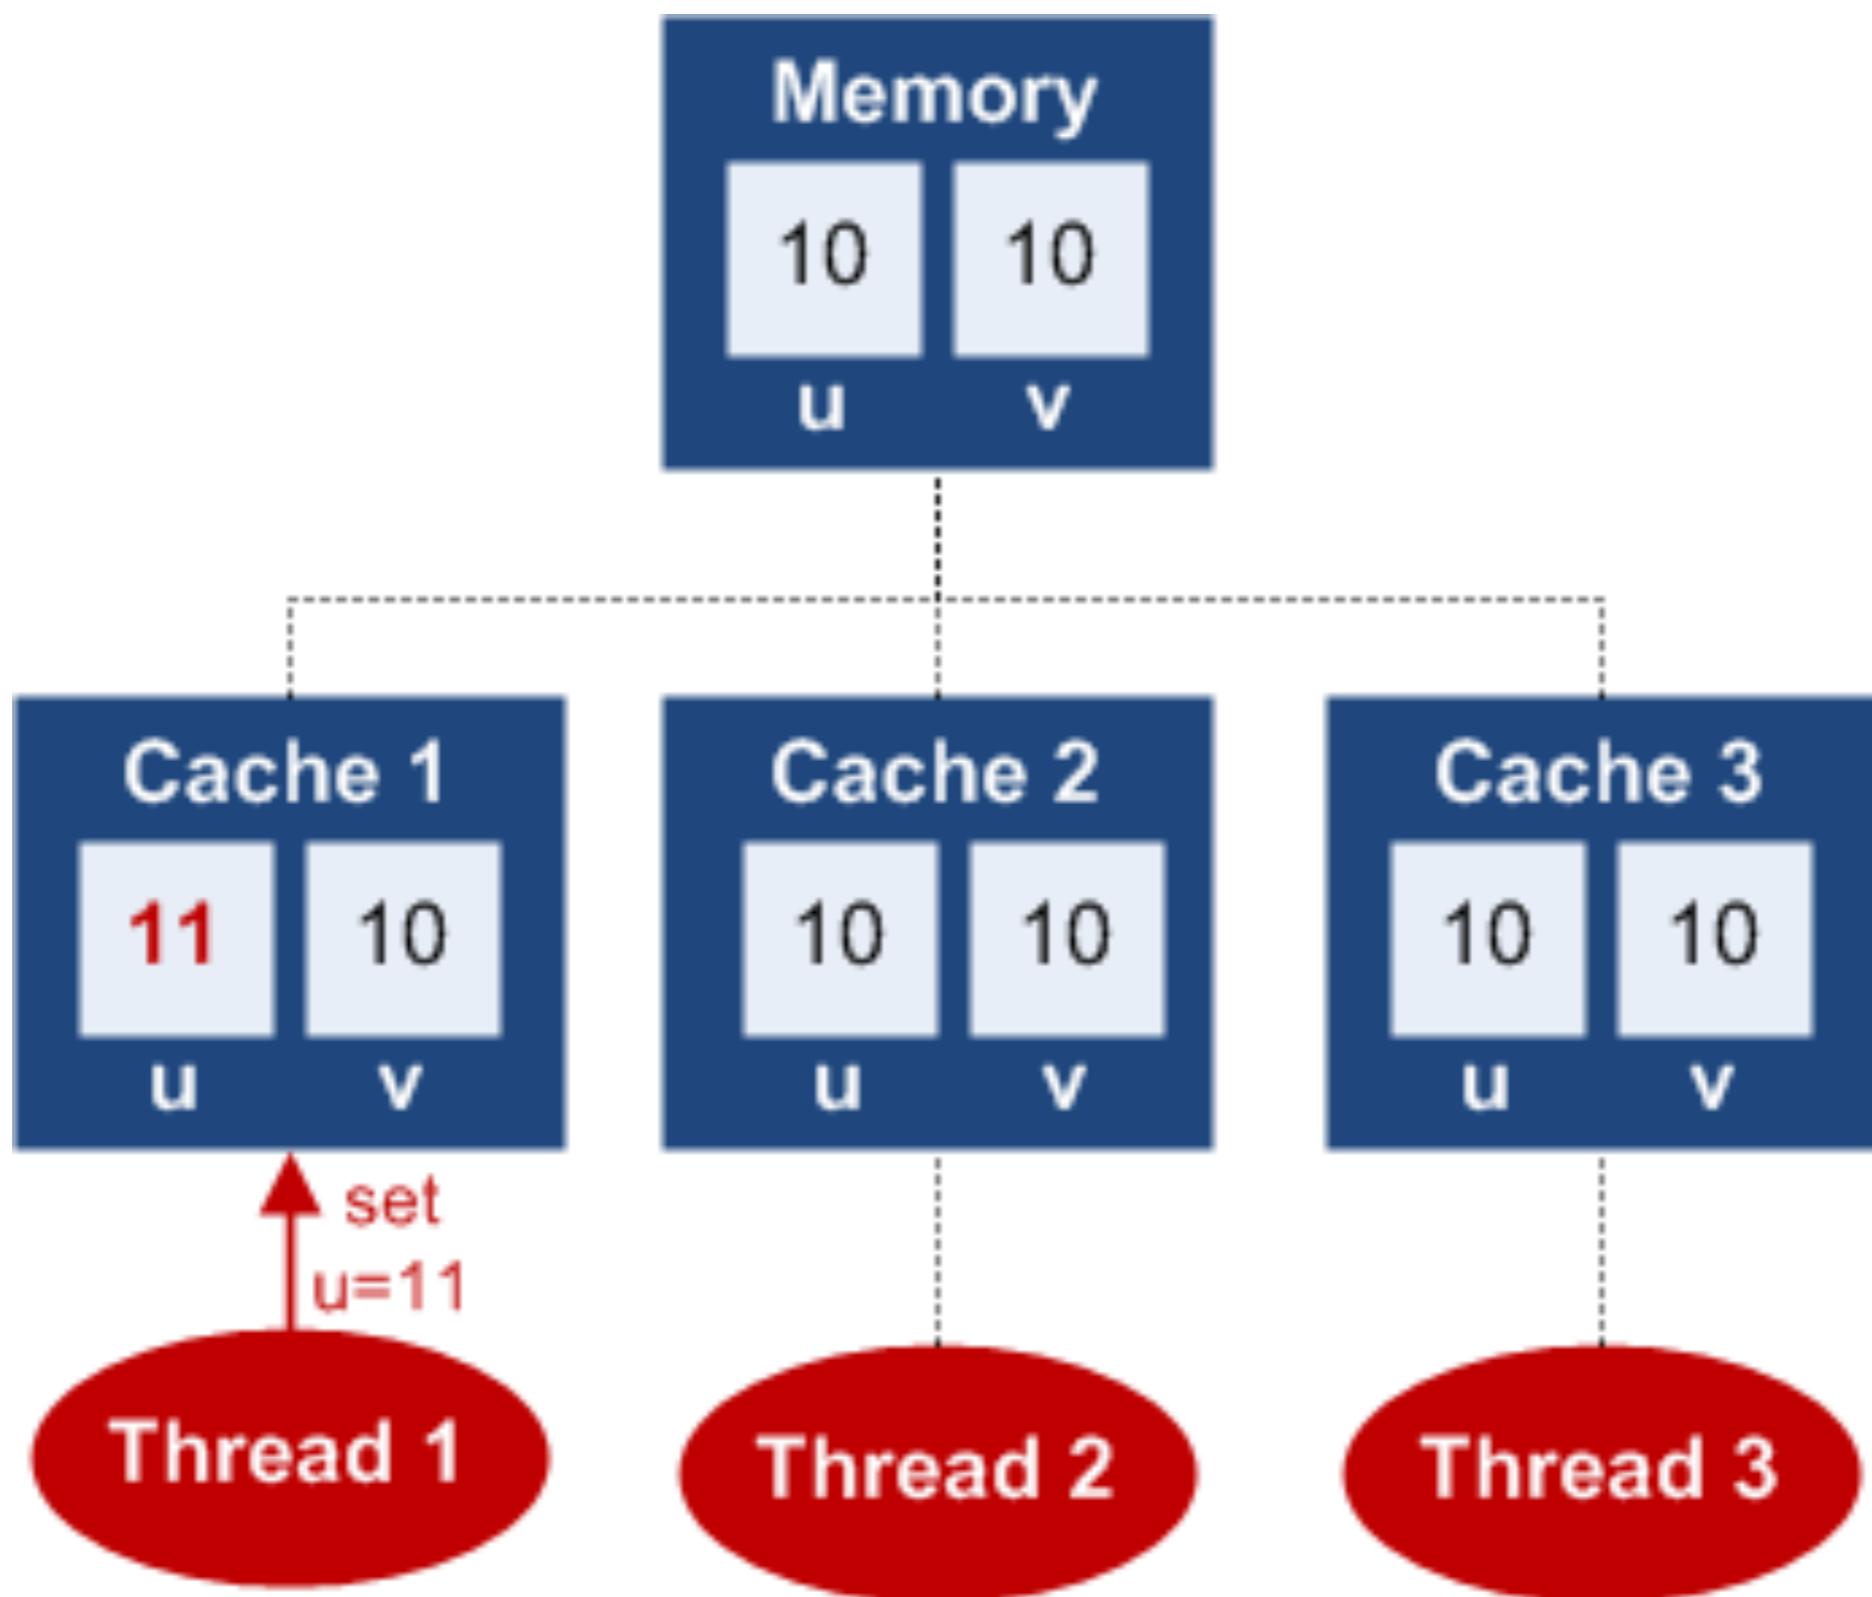
```
class Foo
{
    int _answer;
    bool _complete;
    void A()
    {
        _answer = 123;
        _complete = true; //← instructions could get reordered causing B to write
        "0"
    }
    void B()
    {
        if (_complete)
            Console.WriteLine (_answer);
    }
}
```

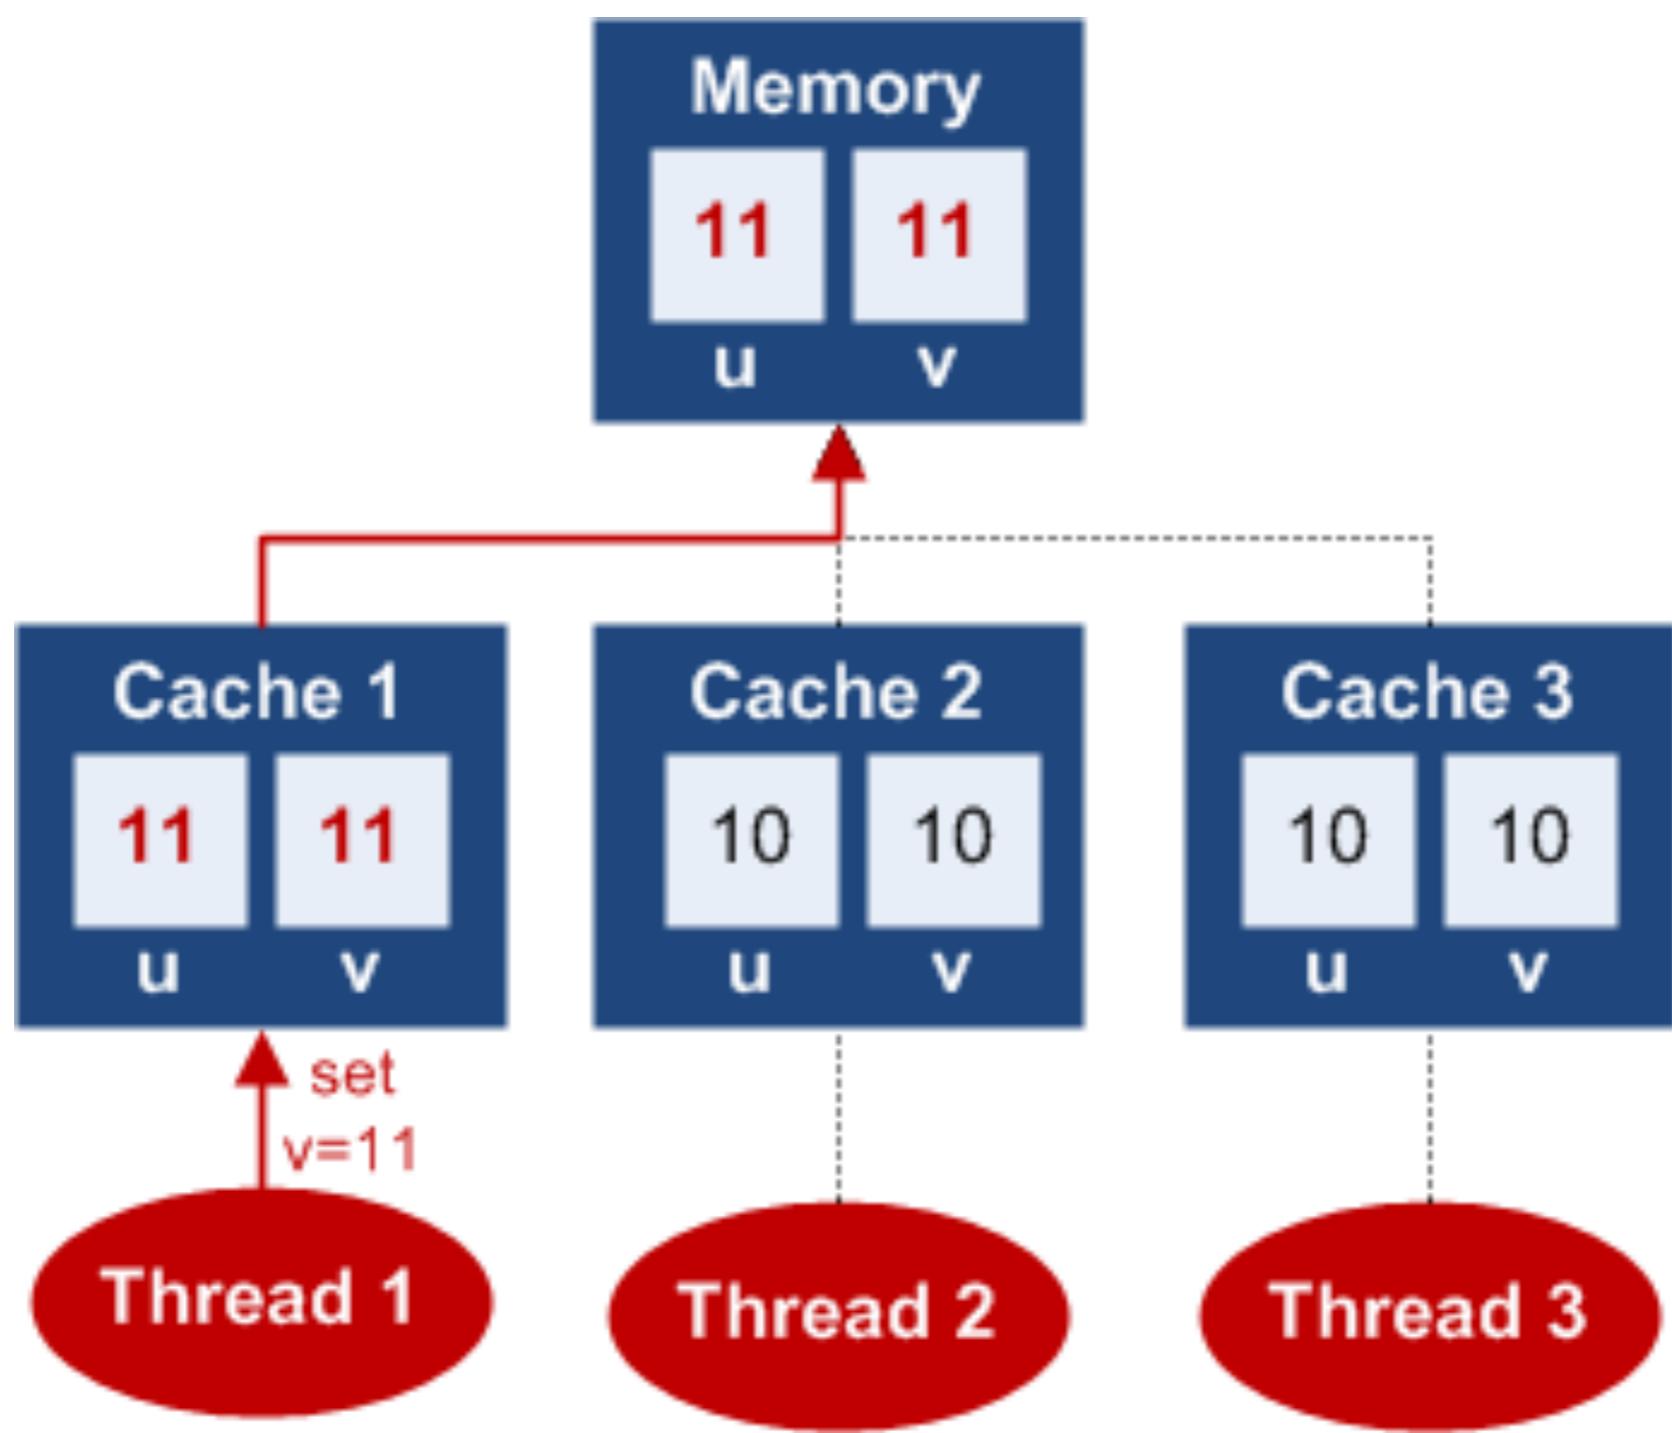
```
static void Main()
{
    bool complete = false;
    var t = new Thread(() =>
    {
        int i=0;
        while (!complete)
        {
            Thread.MemoryBarrier();
            i++;
        }

        Console.WriteLine(i);
    });
    t.Start();
    Thread.Sleep(1000);
    complete = true;
    t.Join();
}
```

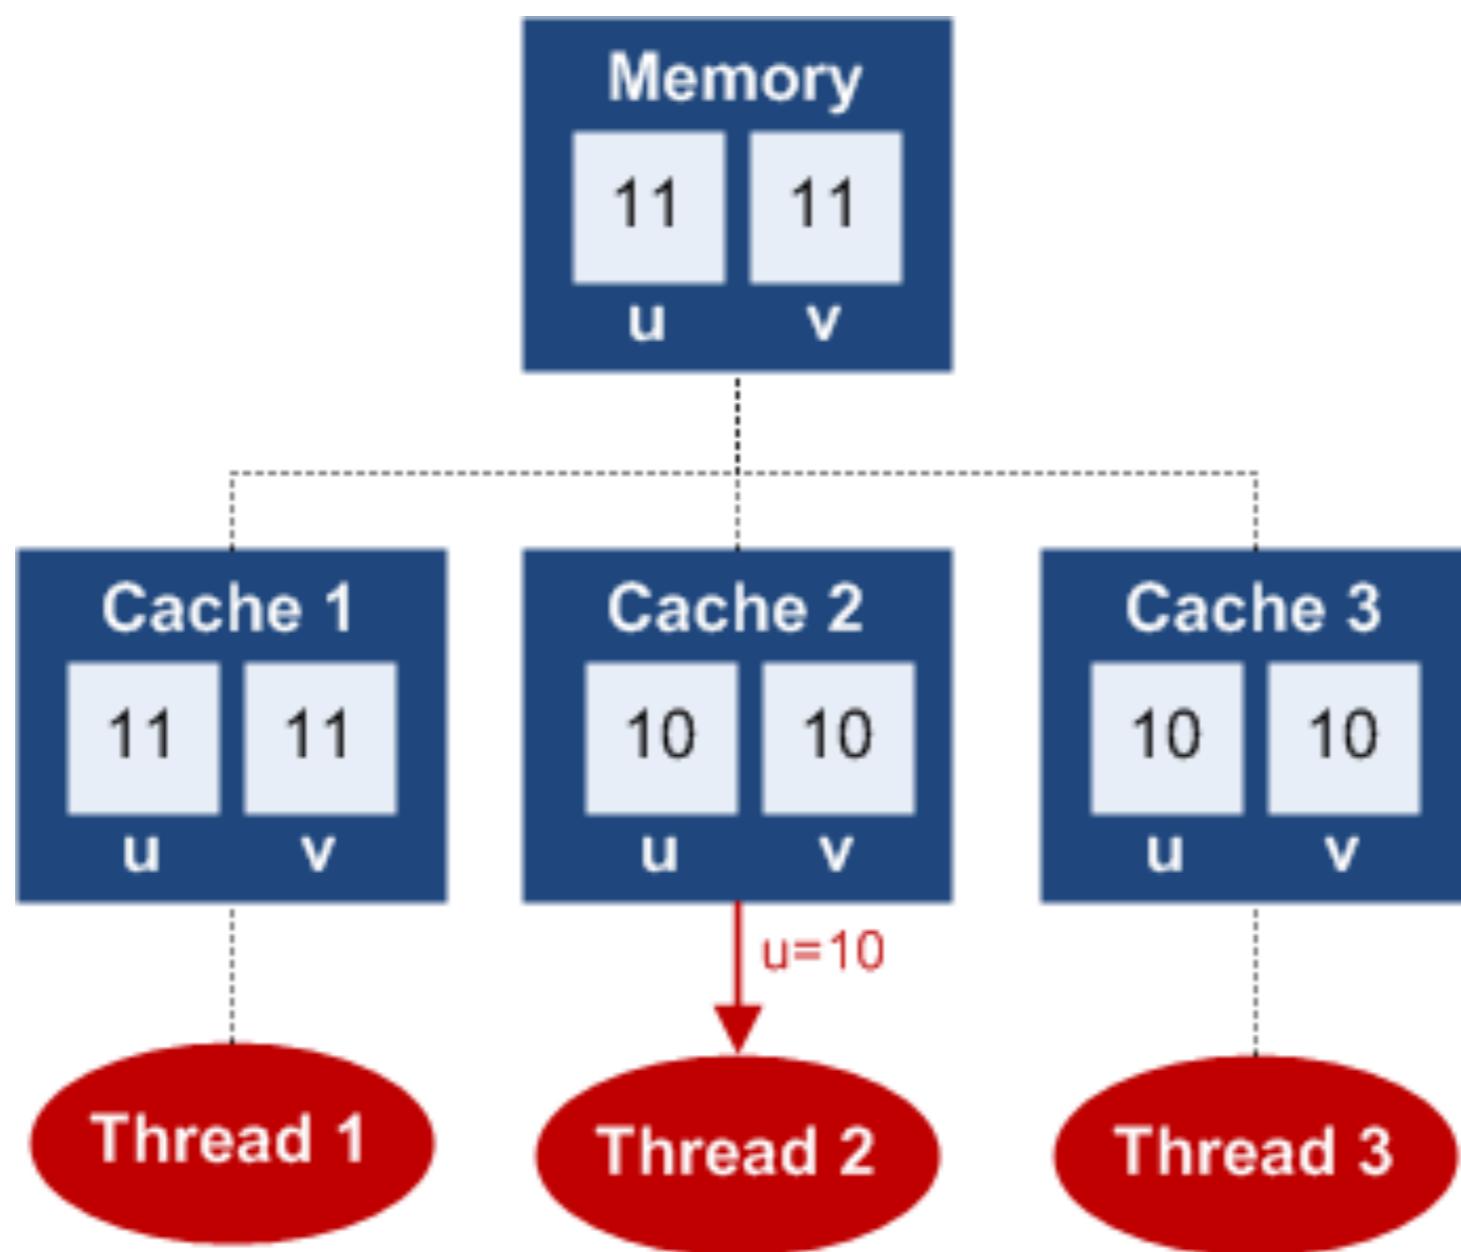
run it with optimizations enabled and Release Mode

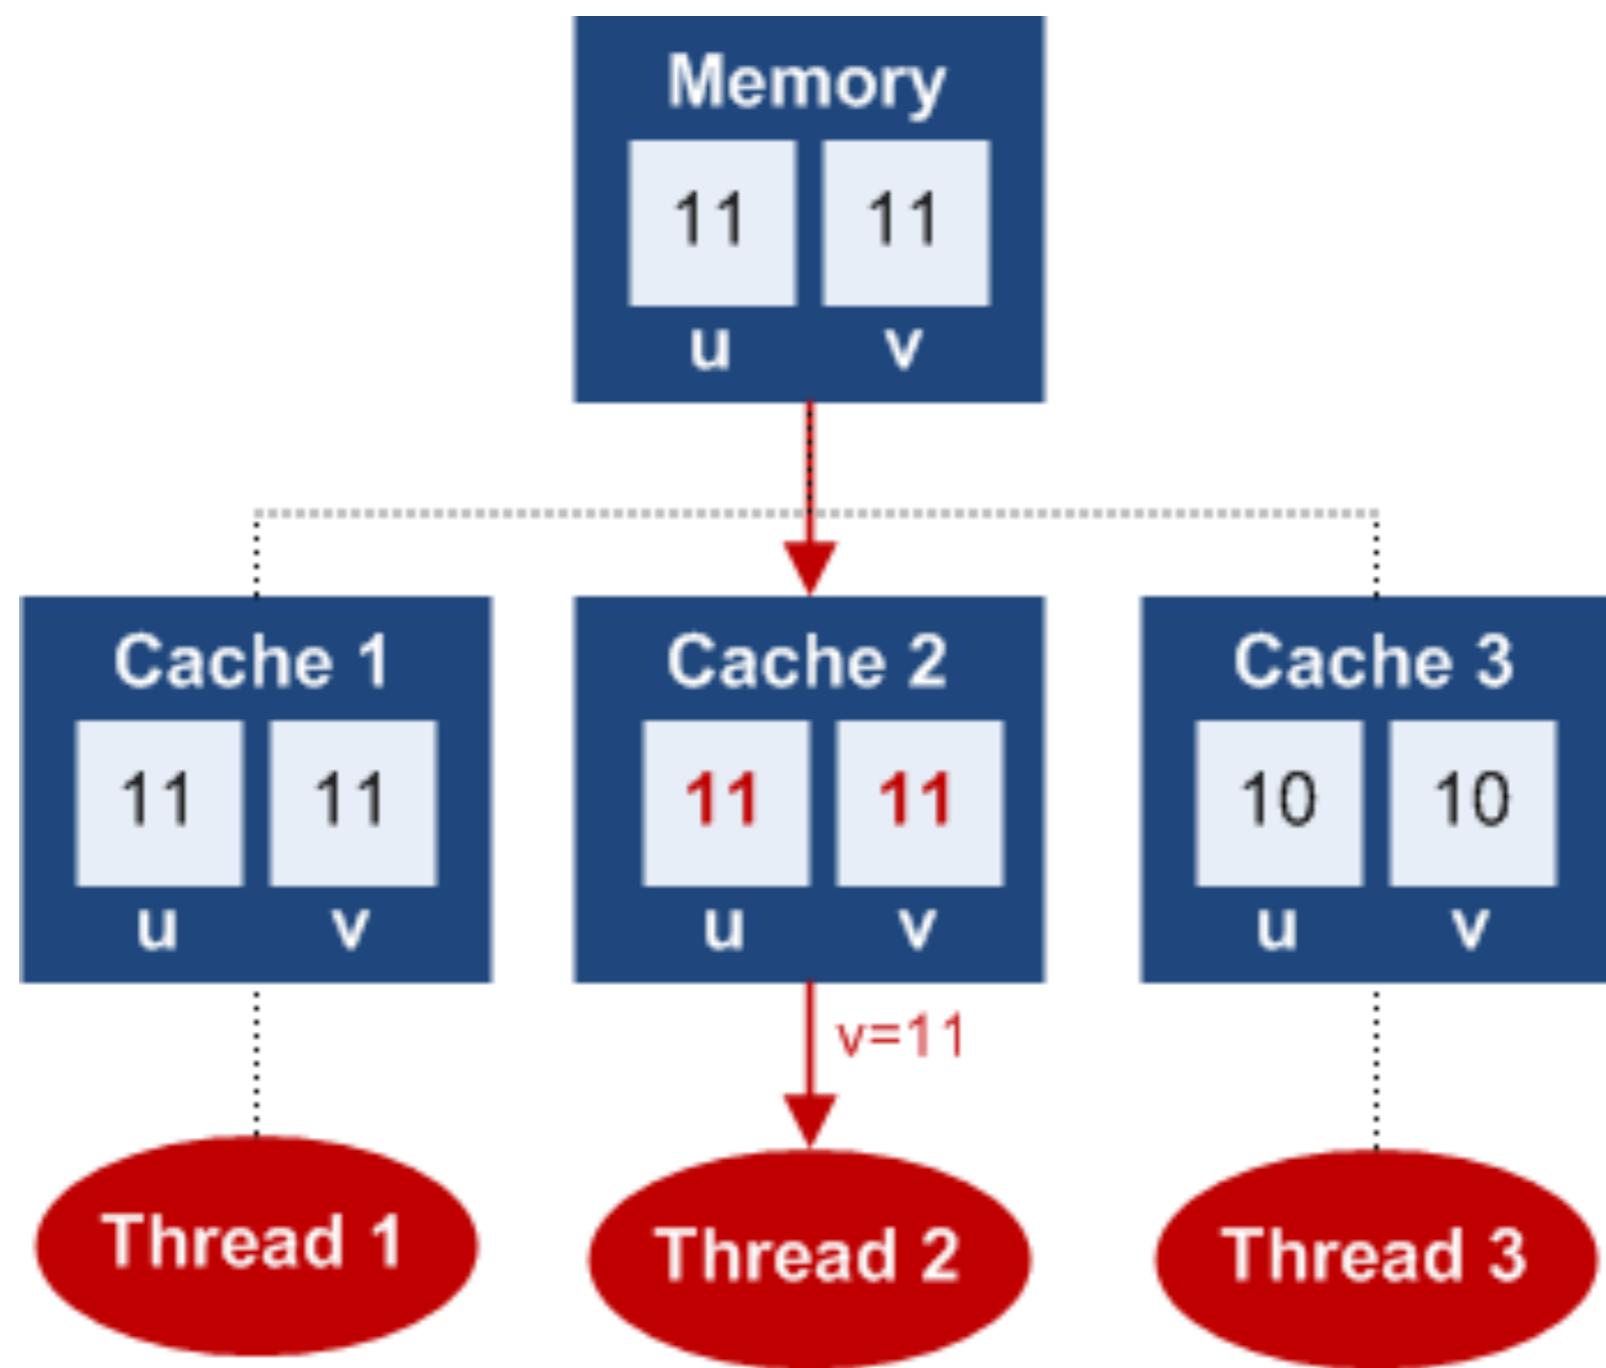






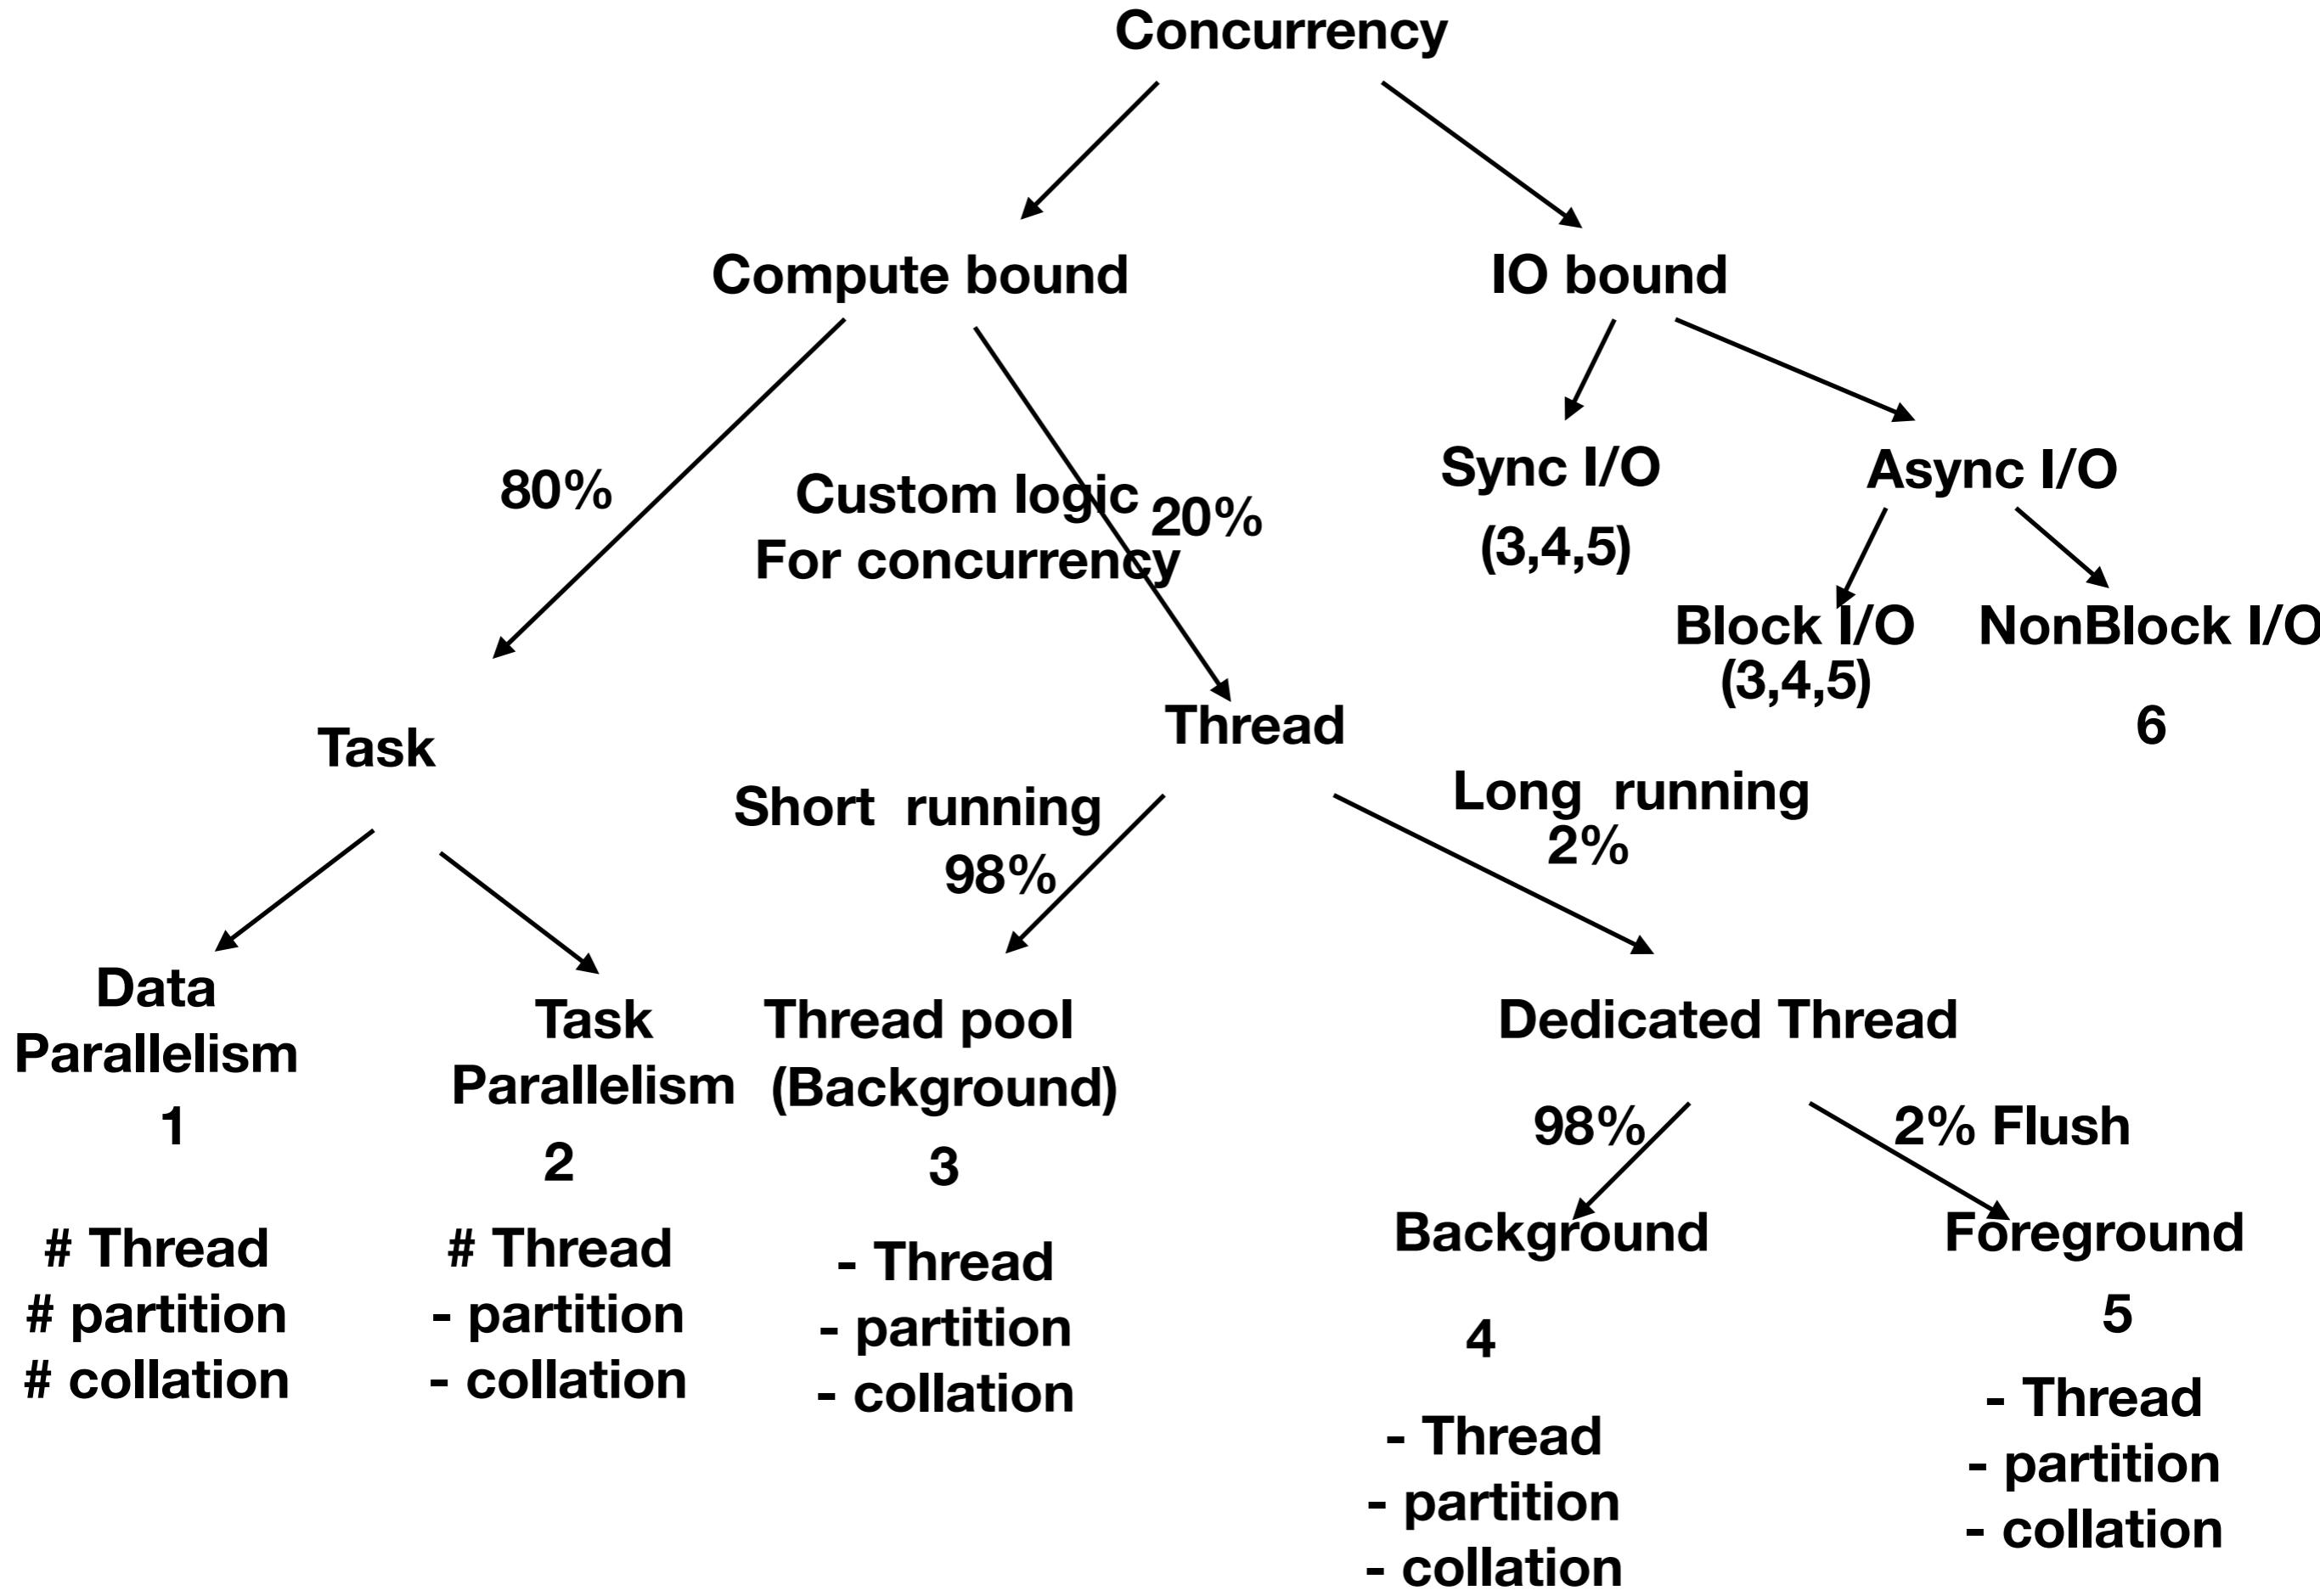
all C# writes are volatile





Synchronization constructs

- **Simple blocking methods**
 - These wait for another thread to finish or for a period of time to elapse.
- **Locking constructs**
 - These limit the number of threads that can perform some activity or execute a section of code at a time.
- **Signaling constructs**
 - These allow a thread to pause until receiving a notification from another
- **Nonblocking synchronization constructs**
 - These protect access to a common field by calling upon processor primitives.

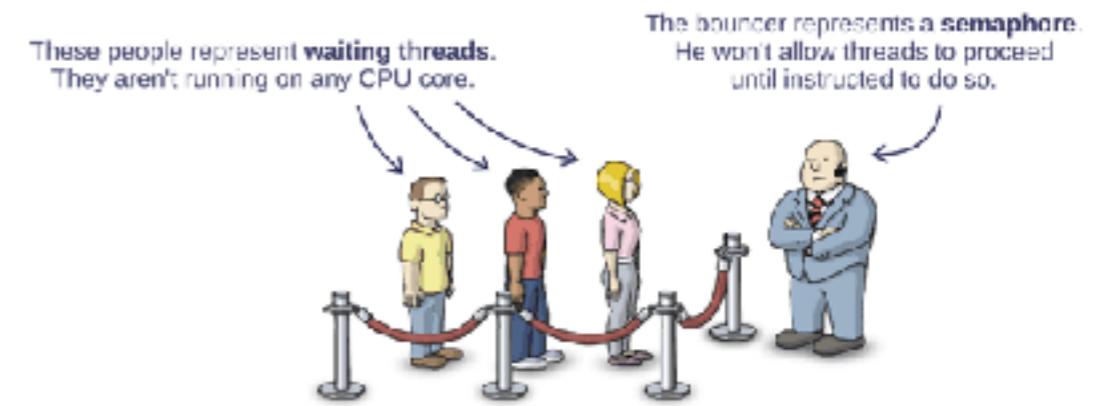


Synchronization

- Nonblocking synchronization constructs
 - volatile
 - Atomic.
 - Simple blocking methods
 - Sleep (bad)
 - Join
 - Spin
 - Locking constructs
 - *Exclusive locking*
 - nonexclusive locking
 - Semaphore
 - reader/writer locks.
 - Signaling constructs
 - event wait
 - Latch
 - Barrier
- Spin(10000)

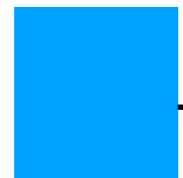


- **Nonblocking synchronization constructs**
 - volatile
 - Atomic.
- **Simple blocking methods (For a Period)**
 - Sleep (bad)
 - Join
 - Spin
- **Locking constructs**
 - *Exclusive* locking
 - nonexclusive locking
 - Semaphore
 - reader/writer locks.
- **Signaling constructs**
 - event wait
 - Latch
 - Barrier



**Subscribe to
Keyboard events**

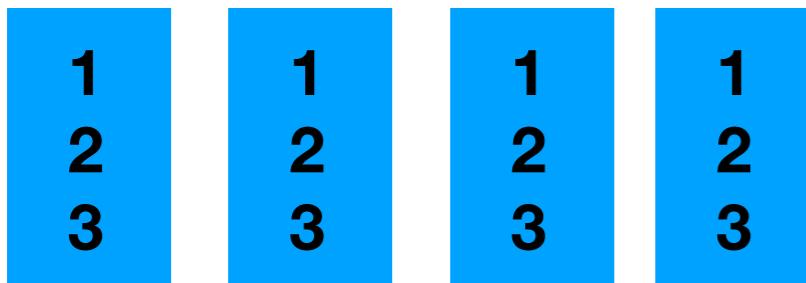
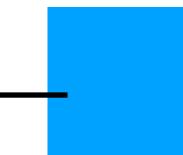
Thread1



Buffer

**Subscribe to
Window events**

Thread2



When user press Enter key,

- 1# read all window events from buffer**
- 2# take a screen shot(image) of the desktop**
- 3# collect os info of active window**

@ PID

@ title

@ type of app

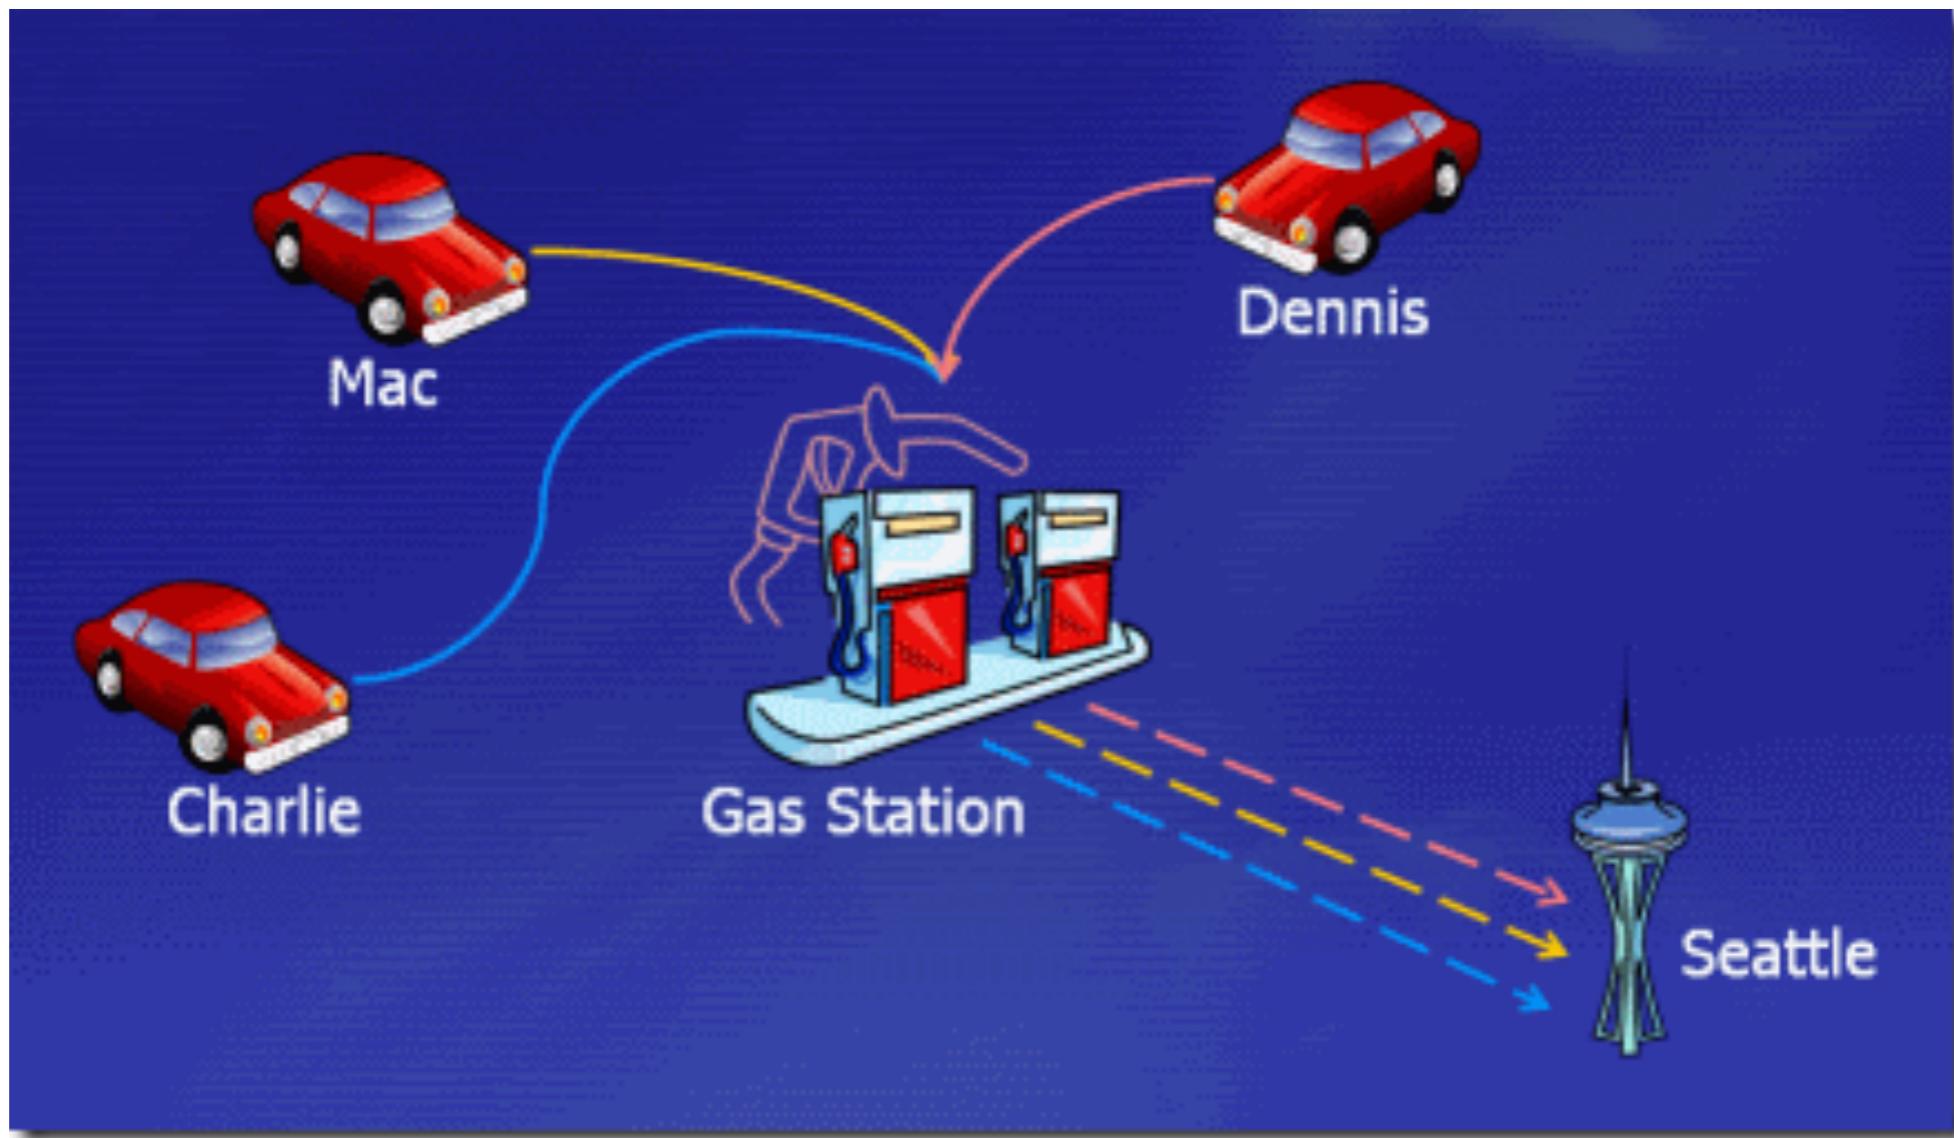
- 4# extract text from the image (OCR)**

- 5# mask the image**

- 6# store the image and data on the disk/nw**

6 Async IO

Latch vs Barrier



Latches are for waiting for events; barriers are for waiting for other threads.

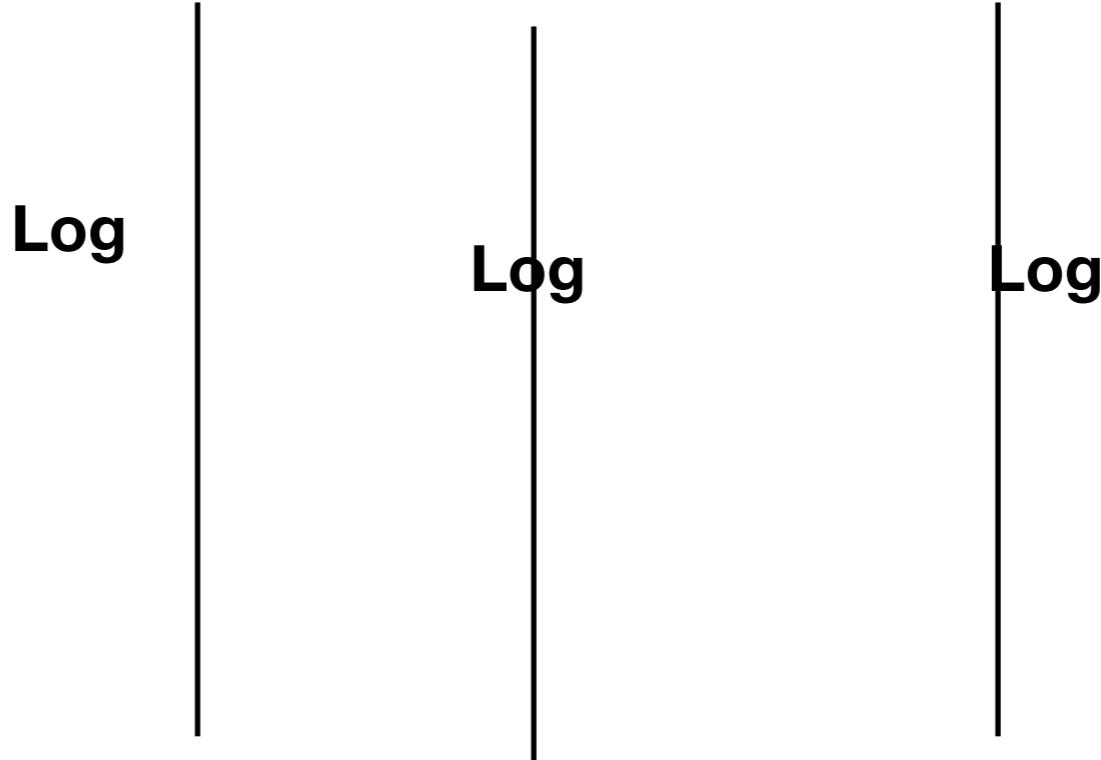
CountDownLatch, no. of calls(tasks) to the countdown() method

CyclicBarrier, no. of threads that should call await()

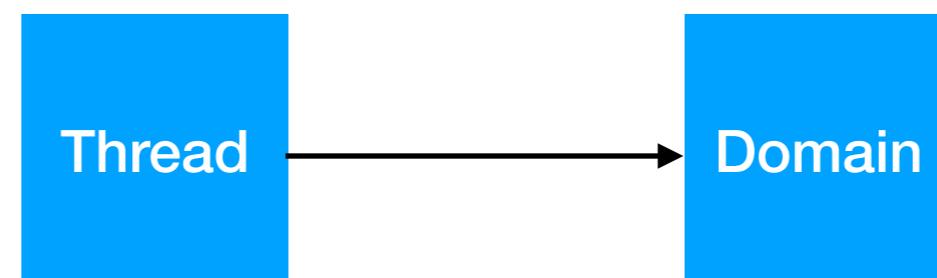


The deprecated Suspend and Resume
methods have two modes: dangerous and
useless!

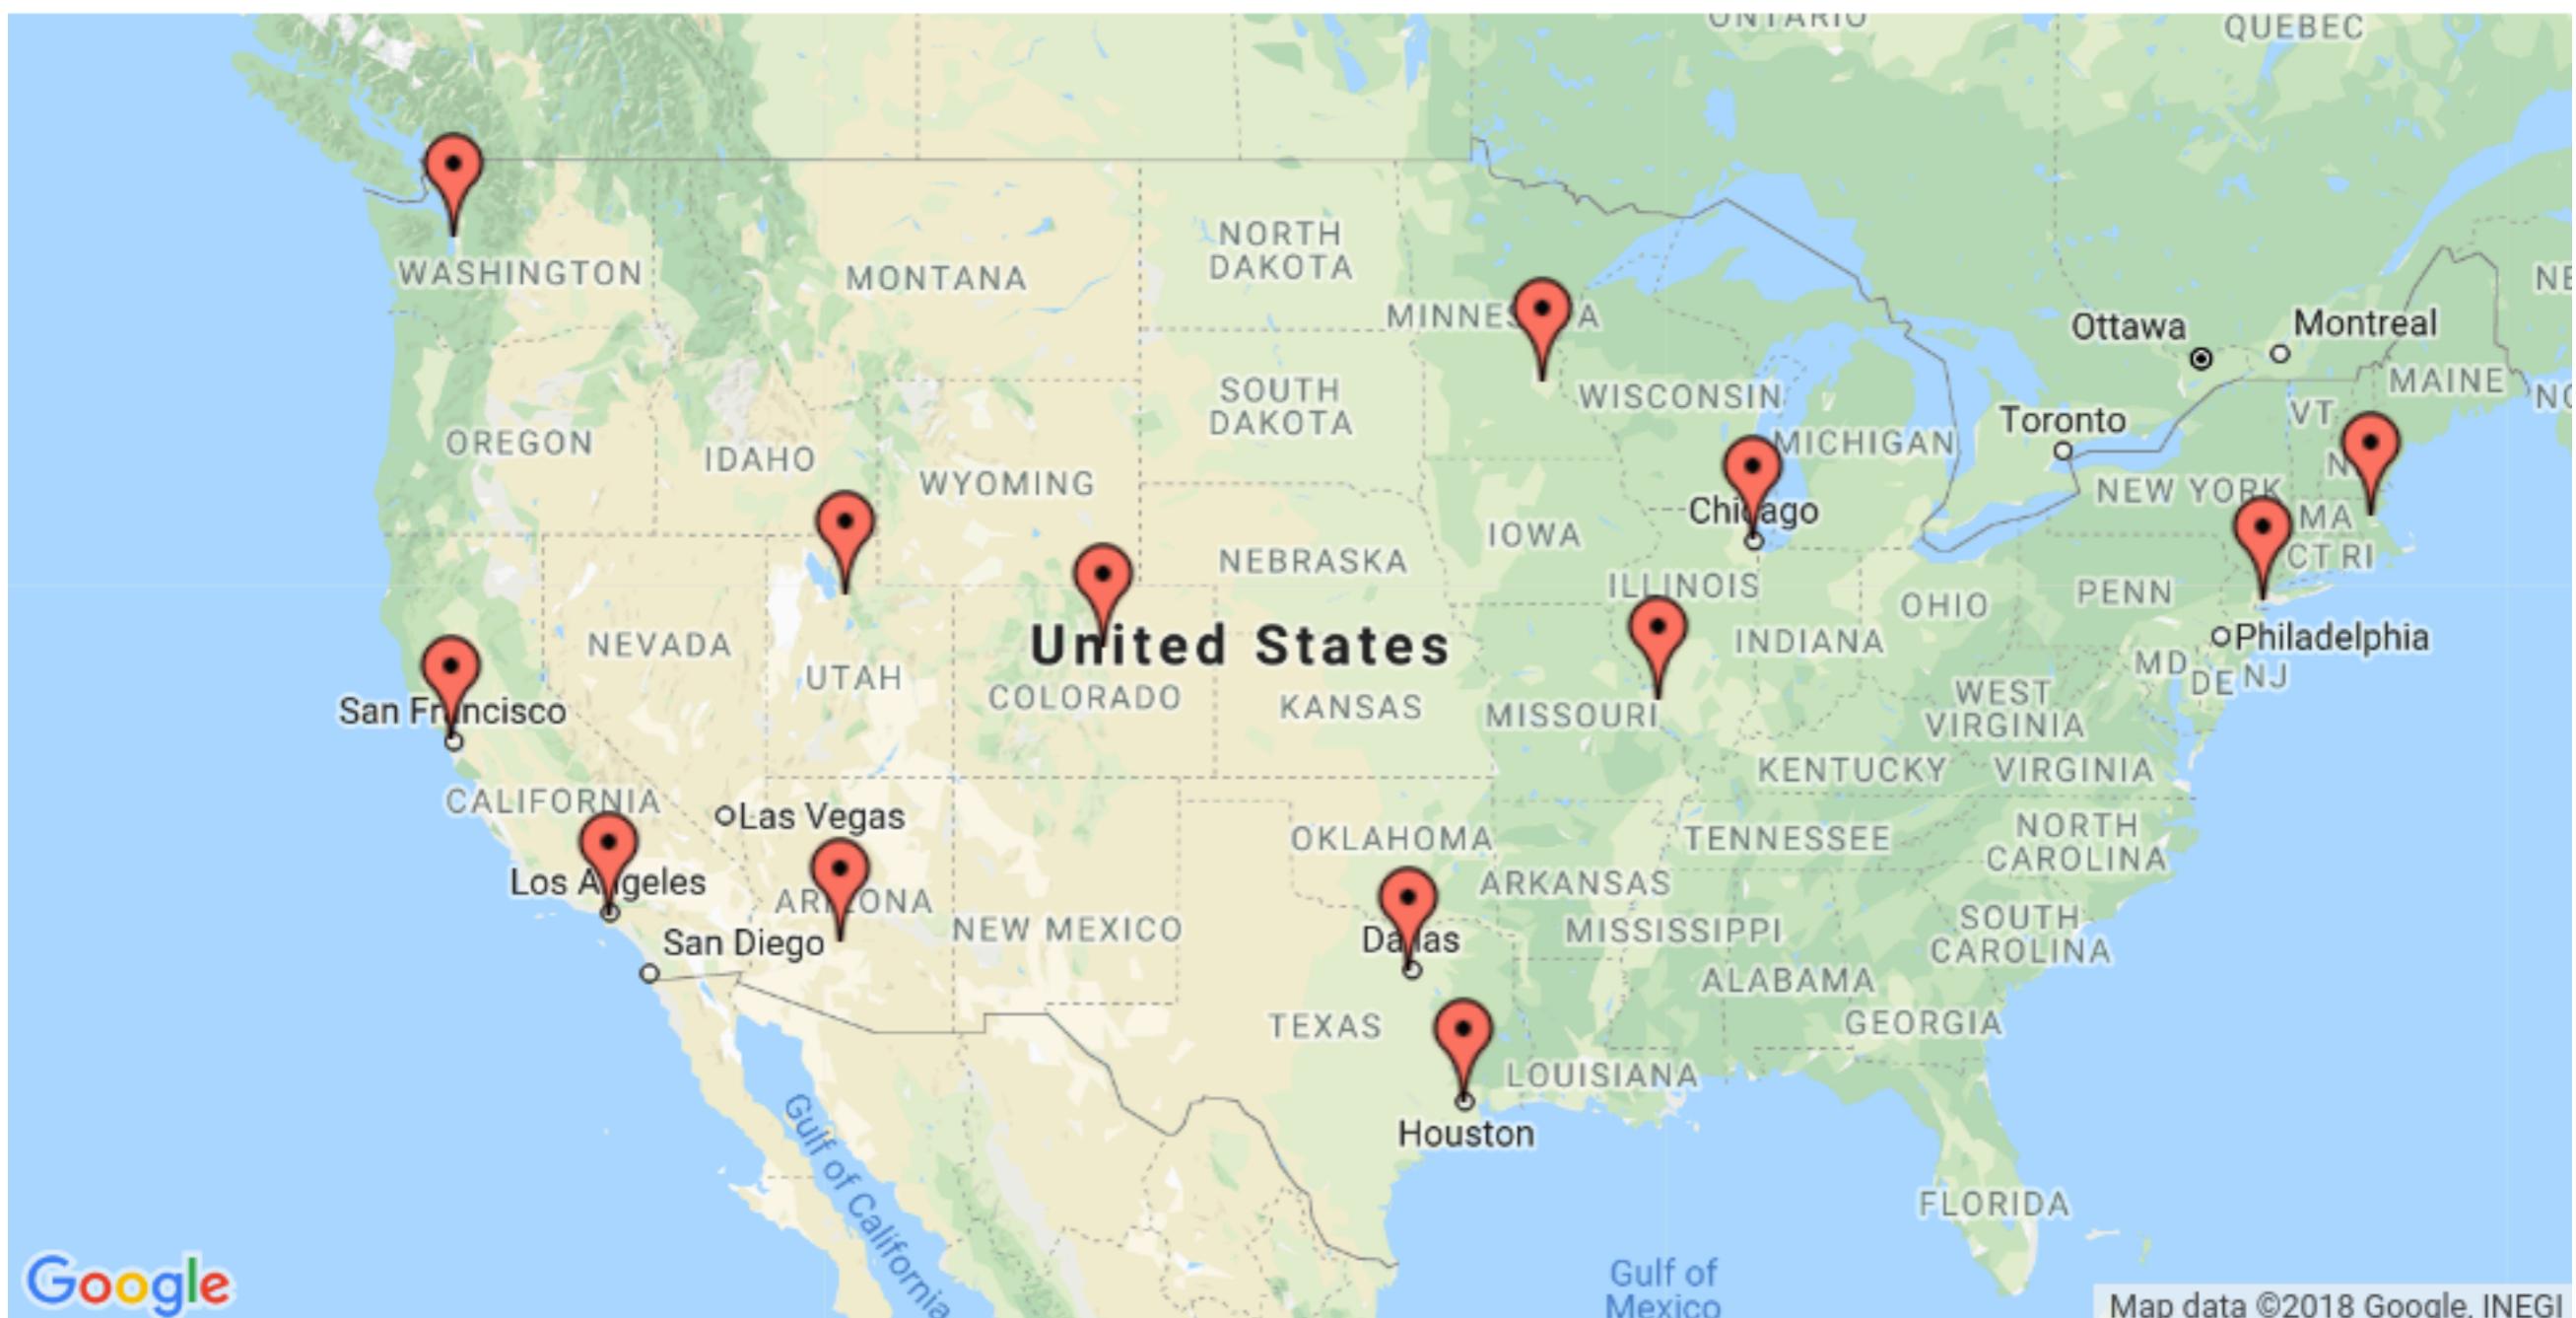
CPU profiler
Memory profiler
Thread Profiler



Humble Objects







0. New York - 1. Los Angeles - 2. Chicago - 3. Minneapolis - 4. Denver - 5. Dallas - 6. Seattle - 7. Boston - 8. San Francisco - 9. St. Louis - 10. Houston - 11. Phoenix - 12. Salt Lake City

```

static class DataModel {
    public final long[][][] distanceMatrix = {
        {0, 2451, 713, 1018, 1631, 1374, 2408, 213, 2571, 875, 1420, 2145, 1972},
        {2451, 0, 1745, 1524, 831, 1240, 959, 2596, 403, 1589, 1374, 357, 579},
        {713, 1745, 0, 355, 920, 803, 1737, 851, 1858, 262, 940, 1453, 1260},
        {1018, 1524, 355, 0, 700, 862, 1395, 1123, 1584, 466, 1056, 1280, 987},
        {1631, 831, 920, 700, 0, 663, 1021, 1769, 949, 796, 879, 586, 371},
        {1374, 1240, 803, 862, 663, 0, 1681, 1551, 1765, 547, 225, 887, 999},
        {2408, 959, 1737, 1395, 1021, 1681, 0, 2493, 678, 1724, 1891, 1114, 701},
        {213, 2596, 851, 1123, 1769, 1551, 2493, 0, 2699, 1038, 1605, 2300, 2099},
        {2571, 403, 1858, 1584, 949, 1765, 678, 2699, 0, 1744, 1645, 653, 600},
        {875, 1589, 262, 466, 796, 547, 1724, 1038, 1744, 0, 679, 1272, 1162},
        {1420, 1374, 940, 1056, 879, 225, 1891, 1605, 1645, 679, 0, 1017, 1200},
        {2145, 357, 1453, 1280, 586, 887, 1114, 2300, 653, 1272, 1017, 0, 504},
        {1972, 579, 1260, 987, 371, 999, 701, 2099, 600, 1162, 1200, 504, 0},
    };
    public final int vehicleNumber = 1;
    public final int depot = 0;
}

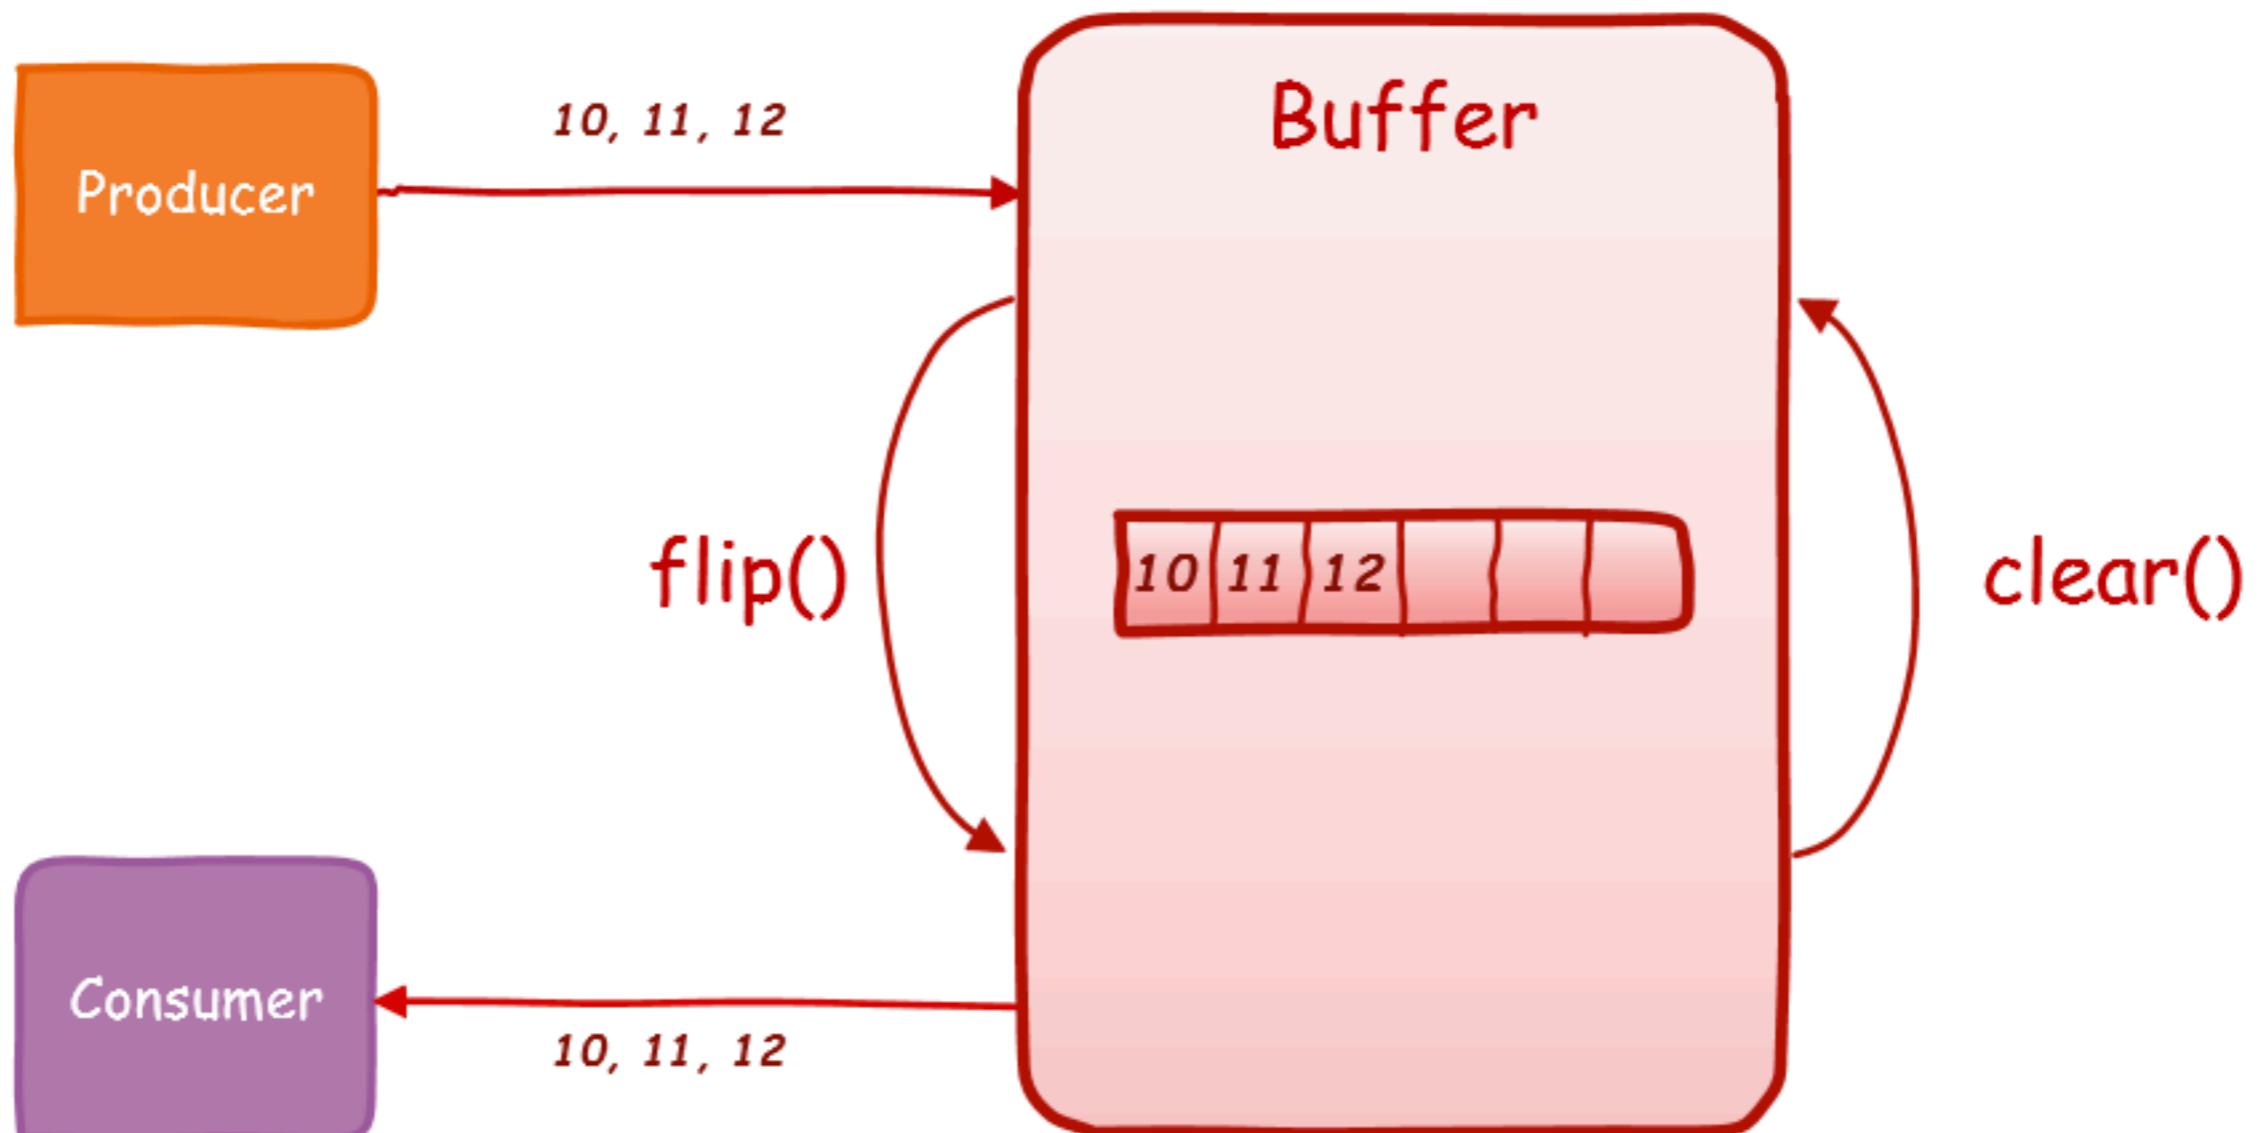
```

- The number of vehicles in the problem, which is 1 because this is a TSP. (For a vehicle routing problem (VRP), the number of vehicles can be greater than 1.)
- The *depot*: the start and end location for the route. In this case, the depot is 0, which corresponds to New York.

TSP with 10 cities can be solved by a DP method in almost 0.2 seconds using intel core i7. This number increases to almost 13 seconds (~60 times greater) with 15 cities.

Appendix

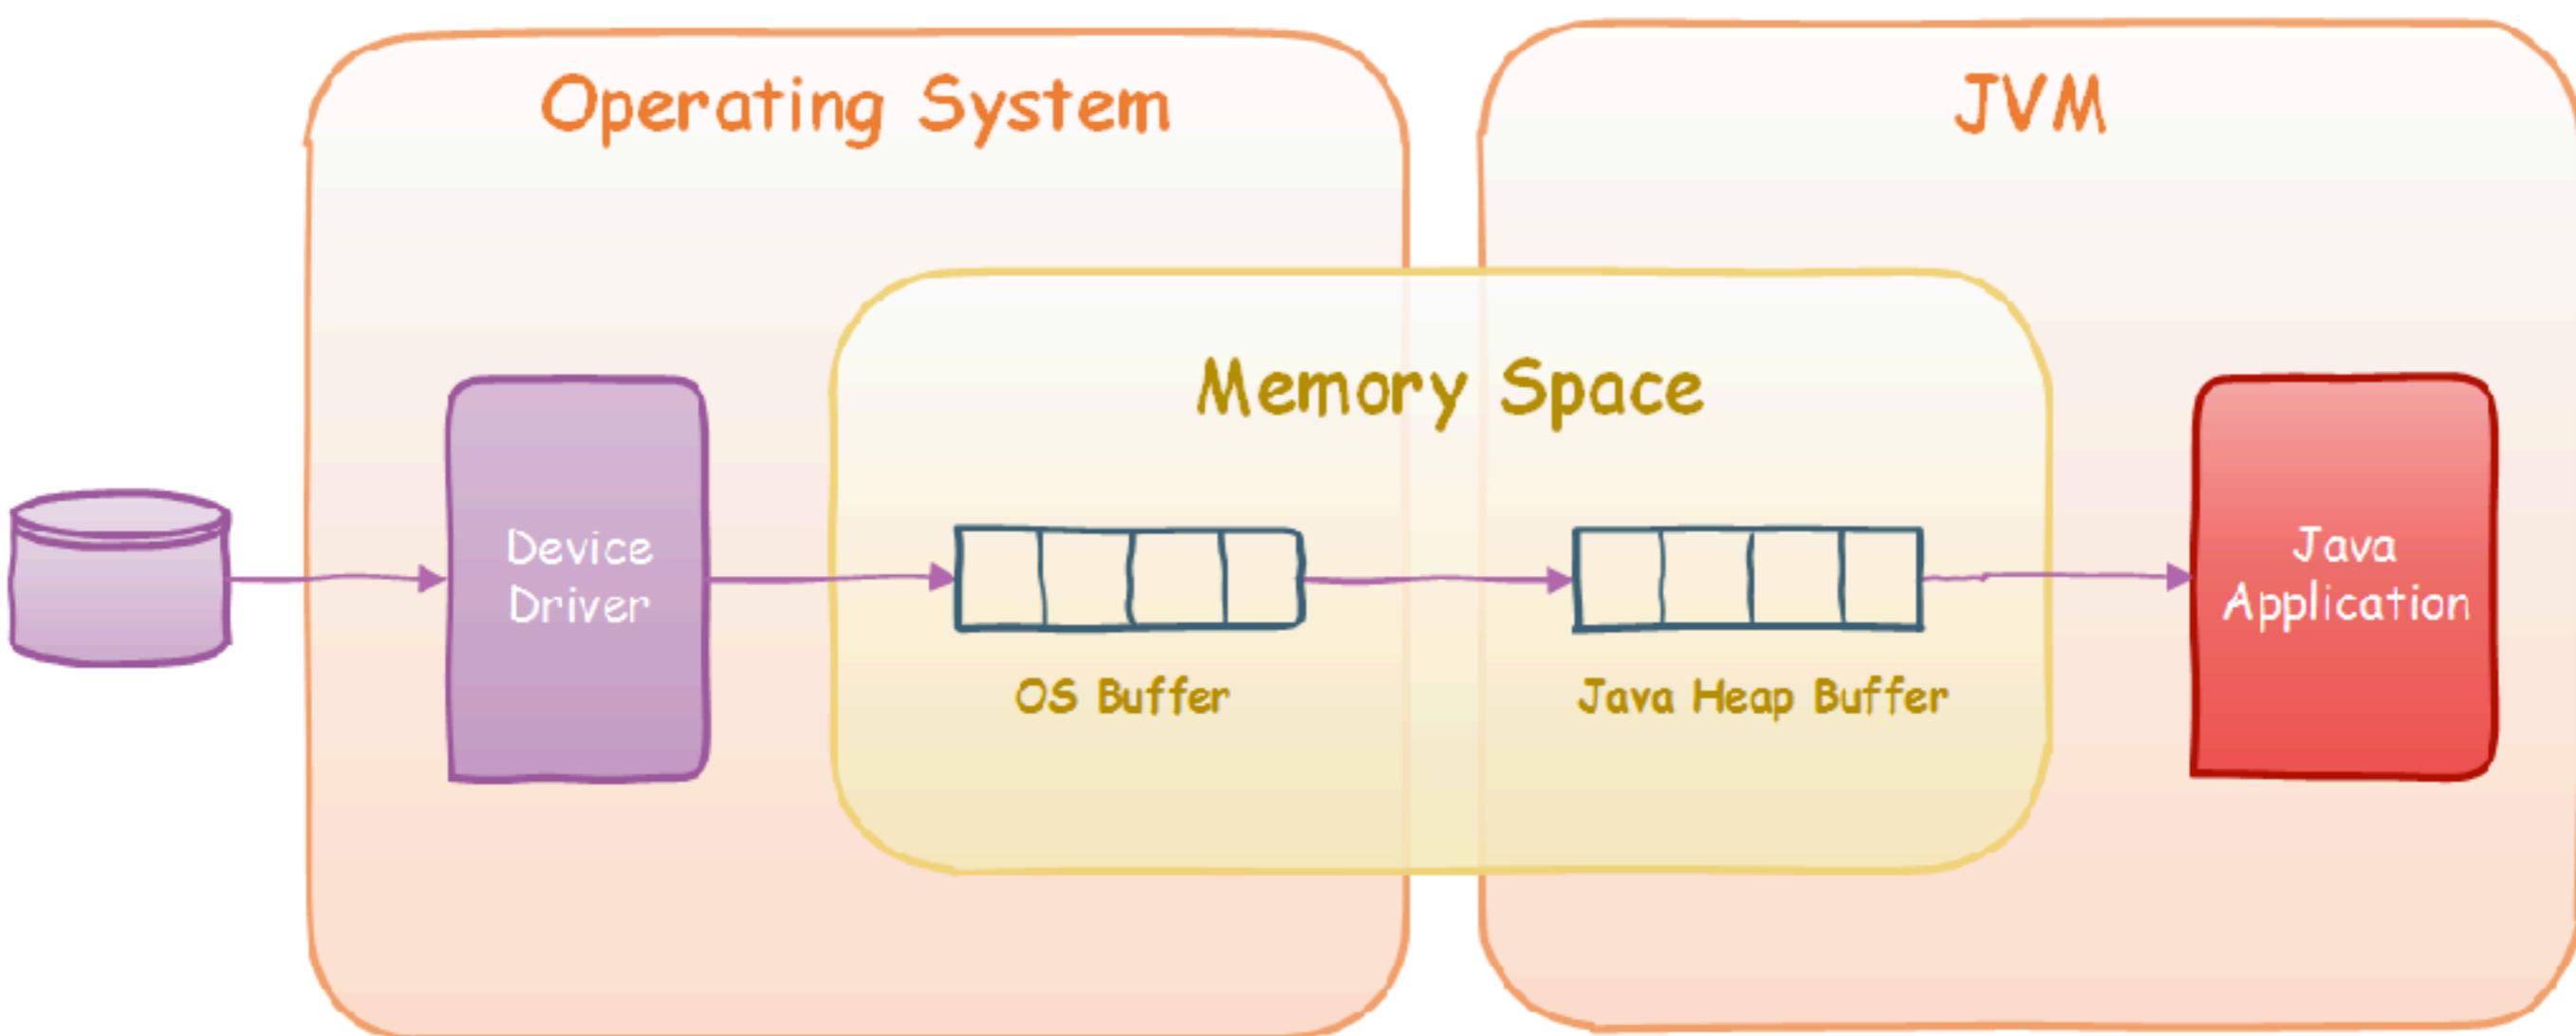
Filling Mode



Draining Mode

The buffer is in **filling mode**. After the producer has finished writing data, the buffer is then **flipped** to prepare it for **draining mode**. At this point, the buffer is ready for the consumer to read the data. Once done, the buffer is then **cleared** and ready for writing again.

Buffers - Direct and Non-Direct



```
ByteBuffer byteBuffer = ByteBuffer.allocateDirect(100);
```